



ID23-5311 Kubernetes Workshop

Wolfgang Ofner

Simon Birrer



Wolfgang Ofner

Freelance Cloud Architect, Toronto, Canada

Focus on Azure, Kubernetes, DevOps and .NET

<https://programmingwithwolfgang.com>

<https://www.linkedin.com/in/wolfgangofner>

https://twitter.com/wolfgang_ofner



Agenda

- Docker
 - Theory
 - Practice Time
- Kubernetes
 - Theory
- Lunch Break
- Kubernetes
 - Theory
 - Practice Time
- Advanced Tools and Learning Path





Challenges of modern Software

Deploy 100 times a day

Versioning

Dependencies

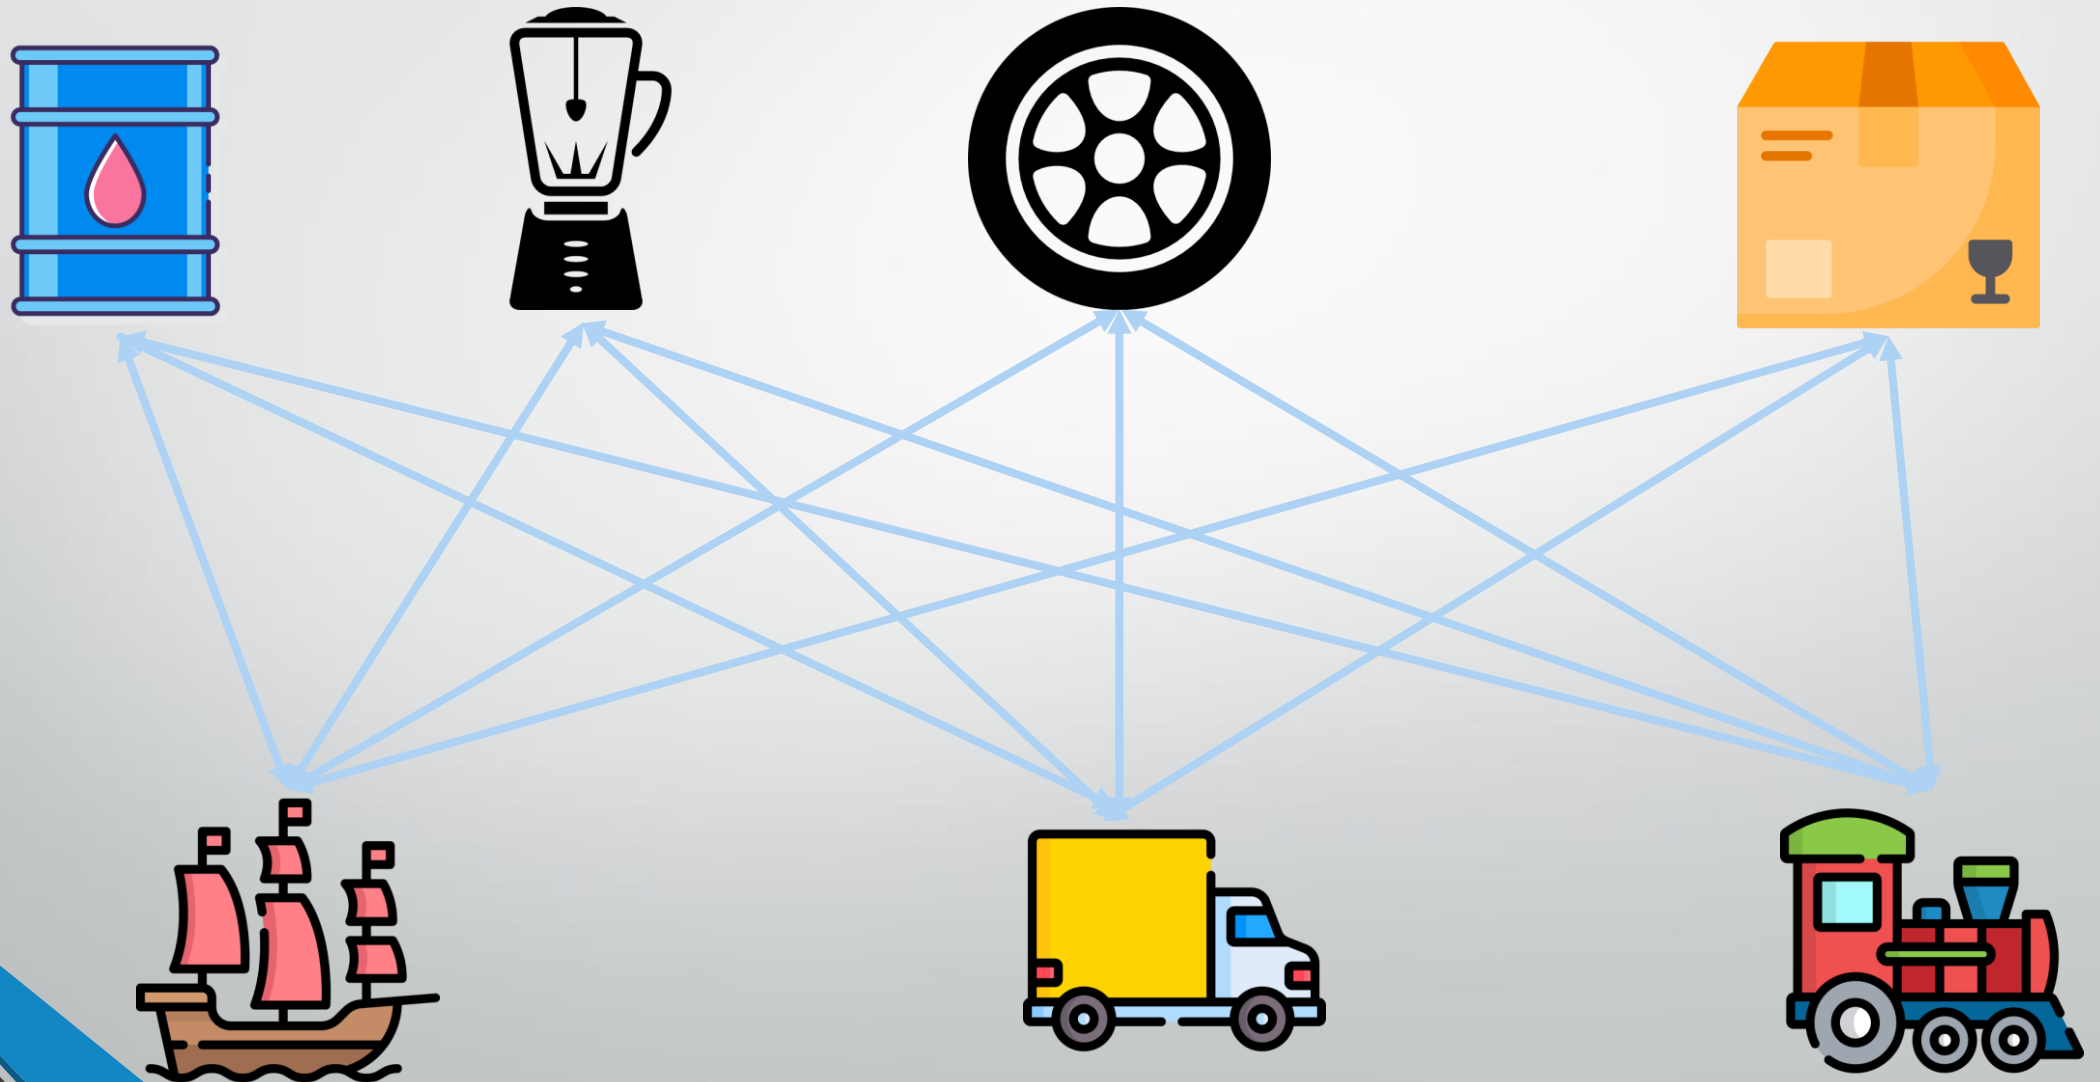
Easy to test

Fast to set up on target machine

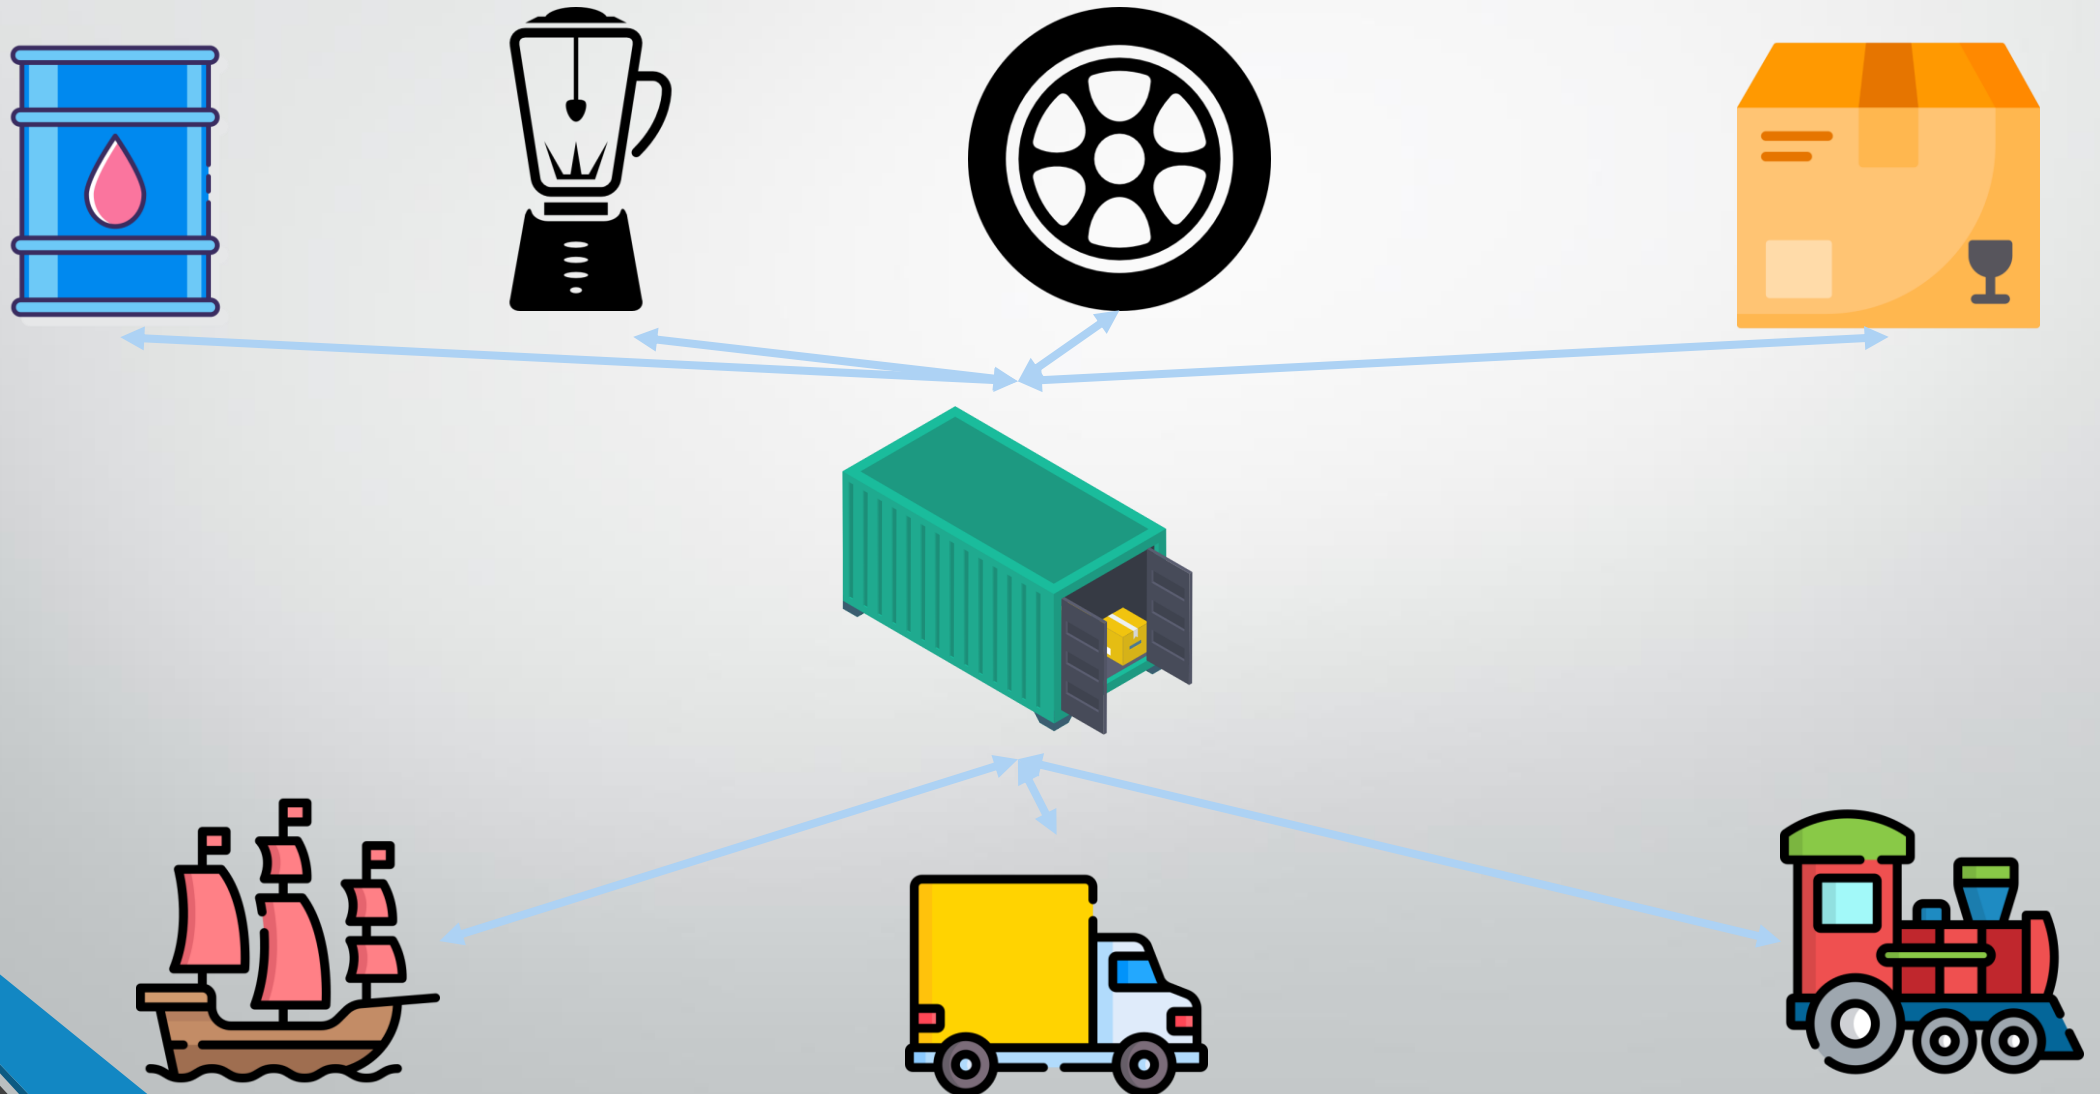
Monolithic software

Database deployments

Freighter transport before 1956



Invention of the Container (1956)







Container Solutions

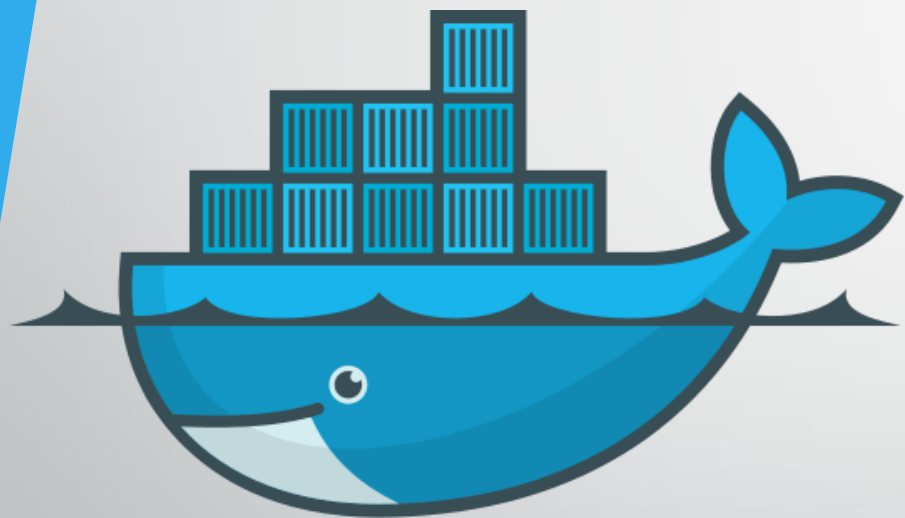
Docker

Podman

Lxc

Crio

Rancher Desktop



docker

Docker Containers

Dockerfile: blueprint

Container: Instance of this blueprint

Versioned artifact

Container image is always bit by bit identical when deployed

Docker Containers

Images for different build platforms, e.g. x86, x64, ARM

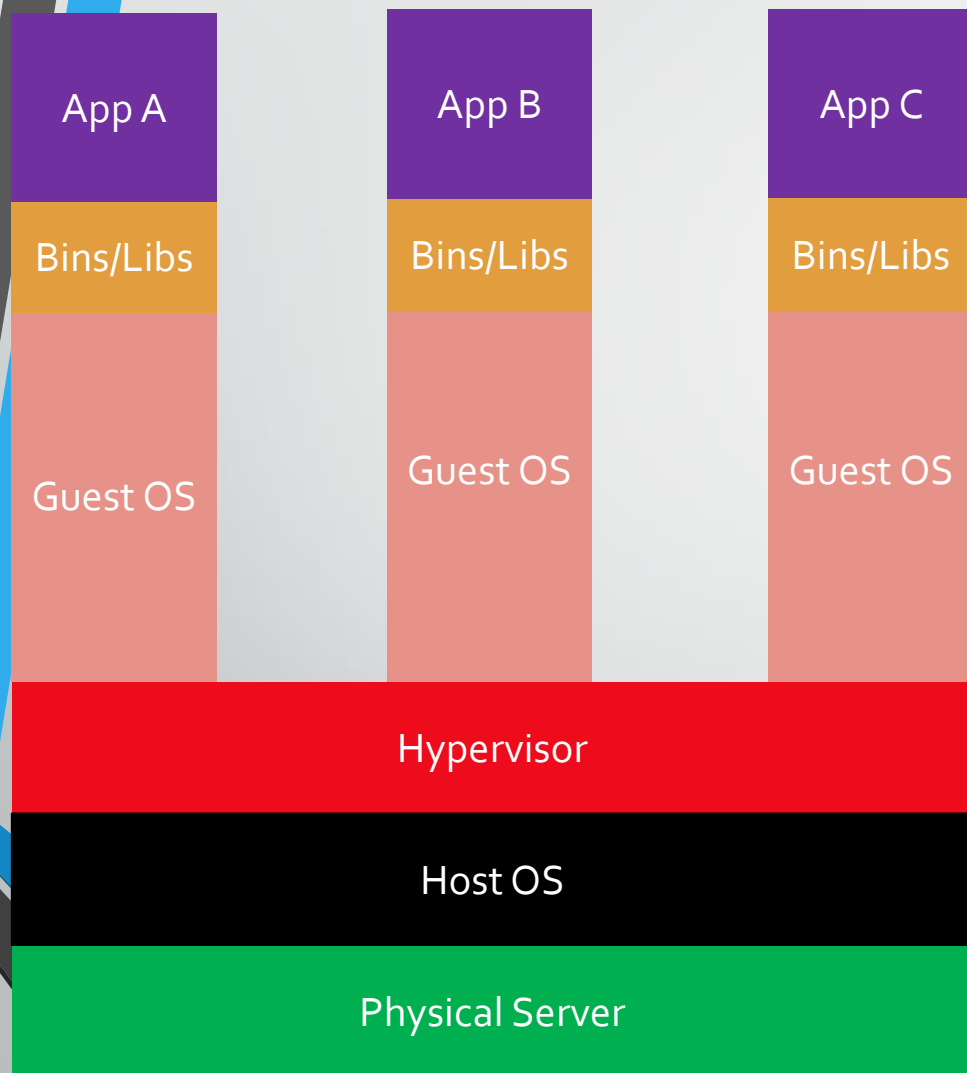
OCI (Open Container Initiative) compliant

Abstracts underlying infrastructure

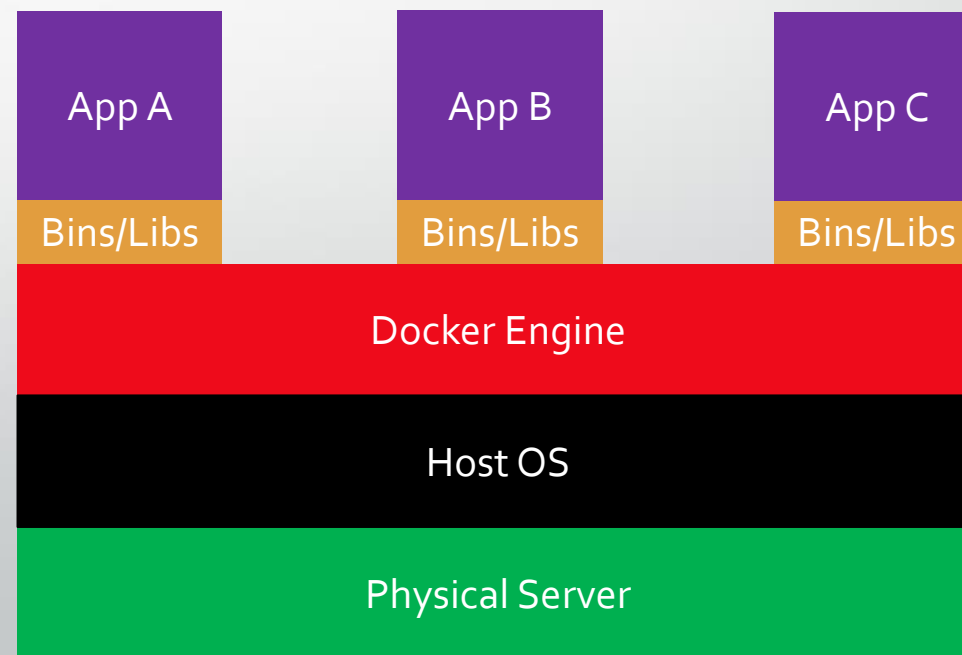
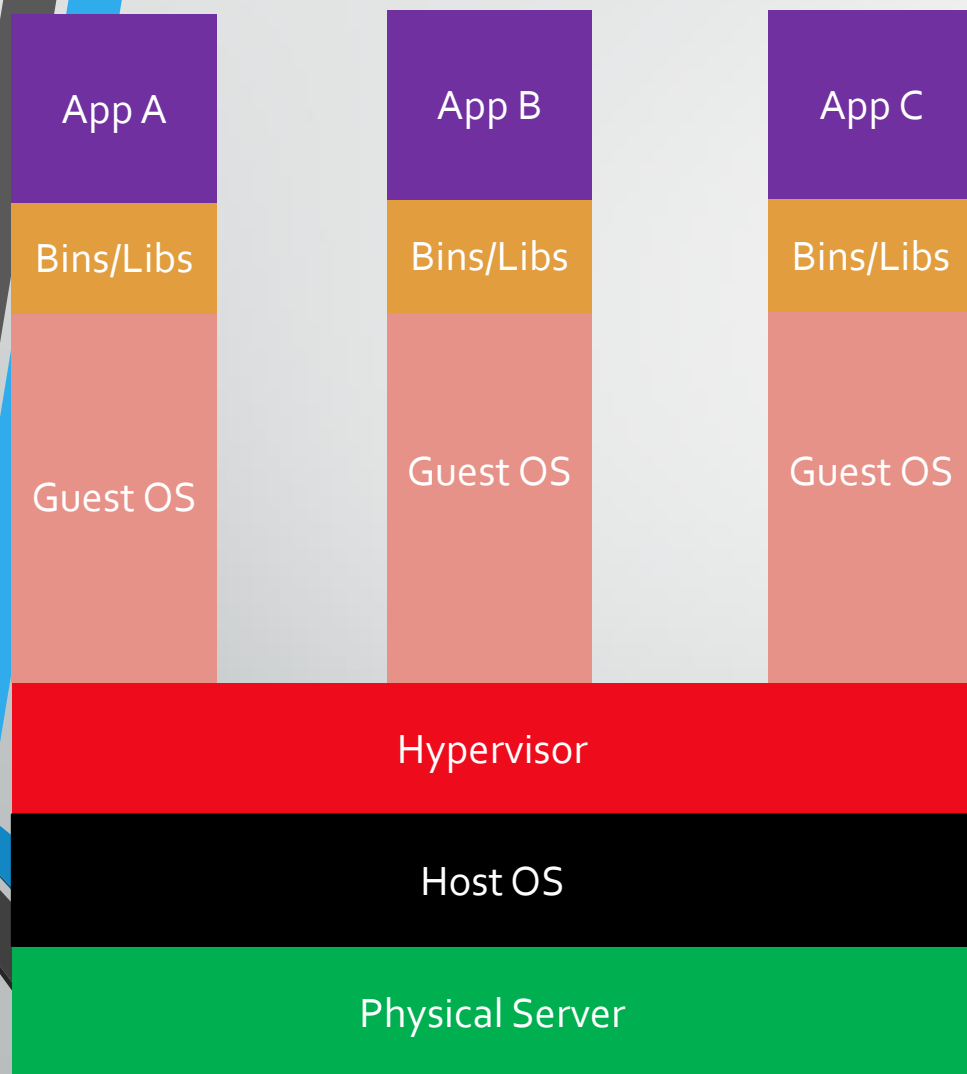
Fast start up times

Pet vs. Kettle

Virtual Machine vs. Container



Virtual Machine vs. Container



Dockerfile

Blueprint to build Docker Image

Can be based on existing images

Commands to update the base OS and install additional software

Build artifacts to include, such as a developed application

Command to run when the container is launched

Dockerfile

```
FROM mcr.microsoft.com/dotnet/aspnet:5.0 AS base
WORKDIR /app
EXPOSE 80
EXPOSE 443

FROM mcr.microsoft.com/dotnet/sdk:5.0 AS build
WORKDIR /src
COPY ["CustomerApi/CustomerApi.csproj", "CustomerApi/"]
RUN dotnet restore "CustomerApi/CustomerApi.csproj"
COPY . .
WORKDIR "/src/CustomerApi"
RUN dotnet build "CustomerApi.csproj" -c Release -o /app/build

FROM build AS publish
RUN dotnet publish "CustomerApi.csproj" -c Release -o /app/publish

FROM base AS final
WORKDIR /app
COPY --from=publish /app/publish .
ENTRYPOINT ["dotnet", "CustomerApi.dll"]
```

```
FROM mcr.microsoft.com/dotnet/aspnet:5.0 AS base
WORKDIR /app
EXPOSE 80
EXPOSE 443

FROM mcr.microsoft.com/dotnet/sdk:5.0 AS build
WORKDIR /src
COPY ["CustomerApi/CustomerApi.csproj", "CustomerApi/"]
COPY ["CustomerApi.Data/CustomerApi.Data.csproj", "CustomerApi.Data/"]
COPY ["CustomerApi.Domain/CustomerApi.Domain.csproj", "CustomerApi.Domain/"]
COPY ["CustomerApi.Service/CustomerApi.Service.csproj", "CustomerApi.Service/"]
COPY ["CustomerApi.Messaging.Send/CustomerApi.Messaging.Send.csproj", "CustomerApi.Messaging.Send/"]
COPY ["CustomerApi.Database.Build/CustomerApi.Database.Build.csproj", "CustomerApi.Database.Build/"]
COPY ["Tests/CustomerApi.Test/CustomerApi.Test.csproj", "Tests/CustomerApi.Test/"]
COPY ["Tests/CustomerApi.Service.Test/CustomerApi.Service.Test.csproj", "Tests/CustomerApi.Service.Test/"]
COPY ["Tests/CustomerApi.Data.Test/CustomerApi.Data.Test.csproj", "Tests/CustomerApi.Data.Test/"]
COPY ["CustomerApi/nuget.config", ""]
COPY ["*.props", "./*"]

ARG PAT=localhost
RUN sed -i "s|</configuration>|<packageSourceCredentials><MicroserviceDemoNugets><add key=\"Username\" value=\"PAT\" /><add key=\"ClearTextPassword\" value=\"${PAT}\" /></MicroserviceDemoNugets></p>|g"

RUN dotnet restore "CustomerApi/CustomerApi.csproj" --configfile "./nuget.config"
RUN dotnet restore "CustomerApi.Database.Build/CustomerApi.Database.Build.csproj" --configfile "./nuget.config"
RUN dotnet restore "Tests/CustomerApi.Test/CustomerApi.Test.csproj" --configfile "./nuget.config"
RUN dotnet restore "Tests/CustomerApi.Service.Test/CustomerApi.Service.Test.csproj" --configfile "./nuget.config"
RUN dotnet restore "Tests/CustomerApi.Data.Test/CustomerApi.Data.Test.csproj" --configfile "./nuget.config"
COPY . .

RUN dotnet build "CustomerApi/CustomerApi.csproj" -c Release -o /app/build --no-restore
RUN dotnet build "Tests/CustomerApi.Test/CustomerApi.Test.csproj" -c Release --no-restore
RUN dotnet build "Tests/CustomerApi.Service.Test/CustomerApi.Service.Test.csproj" -c Release --no-restore
RUN dotnet build "Tests/CustomerApi.Data.Test/CustomerApi.Data.Test.csproj" -c Release --no-restore

FROM build AS dacpac
ARG BuildId=localhost
LABEL dacpac=${BuildId}
WORKDIR /src
RUN dotnet build "CustomerApi.Database.Build/CustomerApi.Database.Build.csproj" -c Release -o /dacpacs --no-restore

FROM build AS test
ARG BuildId=localhost
LABEL test=${BuildId}
RUN dotnet test --no-build -c Release --results-directory /testresults --logger "trx;LogFileName=test_results.trx" /p:CollectCoverage=true /p:CoverletOutputFormat=json%2cCobertura /p:CoverletOutputPath=test_results2.trx
RUN dotnet test --no-build -c Release --results-directory /testresults --logger "trx;LogFileName=test_results2.trx" /p:CollectCoverage=true /p:CoverletOutputFormat=json%2cCobertura /p:CoverletOutputPath=test_results3.trx
RUN dotnet test --no-build -c Release --results-directory /testresults --logger "trx;LogFileName=test_results3.trx" /p:CollectCoverage=true /p:CoverletOutputFormat=json%2cCobertura /p:CoverletOutputPath=test_results3.trx

FROM build AS publish
RUN dotnet publish "CustomerApi/CustomerApi.csproj" --no-restore -c Release -o /app/publish

FROM base AS final
WORKDIR /app
COPY --from=publish /app/publish .
ENTRYPOINT ["dotnet", "CustomerApi.dll"]
```



```
FROM ubuntu:20.04
```

```
RUN DEBIAN_FRONTEND=noninteractive apt-get update
```

```
RUN DEBIAN_FRONTEND=noninteractive apt-get upgrade -y
```

```
RUN DEBIAN_FRONTEND=noninteractive apt-get install -y -qq --no-install-recommends \  
apt-transport-https \  
apt-utils \  
ca-certificates \  
curl \  
git \  
iputils-ping \  
jq \  
lsb-release \  
software-properties-common \  
wget
```

```
RUN curl -sL https://aka.ms/InstallAzureCLIDeb | bash
```

```
RUN wget https://packages.microsoft.com/config/ubuntu/20.04/packages-microsoft-prod.deb -O packages-microsoft-prod.deb
```

```
RUN dpkg -i packages-microsoft-prod.deb
```

```
RUN rm packages-microsoft-prod.deb
```

```
RUN echo 'deb http://download.opensuse.org/repositories/devel:/kubic:/libcontainers:/stable/xUbuntu\_20.04/ /' | tee /etc/apt/sources.list.d/kubernetes.list
```

```
RUN curl -fsSL https://download.opensuse.org/repositories/devel:kubic:libcontainers:stable/xUbuntu\_20.04/Release.key |
```

```
RUN apt-get update
```

```
RUN apt-get install -y dotnet-sdk-6.0
```

```
RUN apt-get install -y dotnet-sdk-7.0
```

```
RUN apt -y install podman fuse-overlayfs
```



Docker Demo



Onion File System

Every command is a
new layer

Layers can be cached

Faster builds

Onion File System

Every command is a
new layer

Layers can be cached

Faster builds

```
FROM mcr.microsoft.com/dotnet/aspnet:5.0 AS base
WORKDIR /app
EXPOSE 80
EXPOSE 443

FROM mcr.microsoft.com/dotnet/sdk:5.0 AS build
WORKDIR /src
COPY ["CustomerApi/CustomerApi.csproj", "CustomerApi/"]
RUN dotnet restore "CustomerApi/CustomerApi.csproj"
COPY . .
WORKDIR "/src/CustomerApi"
RUN dotnet build "CustomerApi.csproj" -c Release -o /app/build

FROM build AS publish
RUN dotnet publish "CustomerApi.csproj" -c Release -o /app/publish

FROM base AS final
WORKDIR /app
COPY --from=publish /app/publish .
ENTRYPOINT ["dotnet", "CustomerApi.dll"]
```

Onion File System

Every command is a
new layer

Layers can be cached

Faster builds

```
FROM mcr.microsoft.com/dotnet/aspnet:5.0 AS base
```

```
WORKDIR /app
```

```
EXPOSE 80
```

```
EXPOSE 443
```

```
FROM mcr.microsoft.com/dotnet/sdk:5.0 AS build
```

```
WORKDIR /src
```

```
COPY ["CustomerApi/CustomerApi.csproj", "CustomerApi/"]
```

```
RUN dotnet restore "CustomerApi/CustomerApi.csproj"
```

```
COPY . .
```

```
WORKDIR "/src/CustomerApi"
```

```
RUN dotnet build "CustomerApi.csproj" -c Release -o /app/build
```

```
FROM build AS publish
```

```
RUN dotnet publish "CustomerApi.csproj" -c Release -o /app/publish
```

```
FROM base AS final
```

```
WORKDIR /app
```

```
COPY --from=publish /app/publish .
```

```
ENTRYPOINT ["dotnet", "CustomerApi.dll"]
```

Onion File System

Every command is a
new layer

Layers can be cached

Faster builds

11 Layer

```
FROM mcr.microsoft.com/dotnet/aspnet:5.0 AS base
```

```
WORKDIR /app
```

```
EXPOSE 80
```

```
EXPOSE 443
```

```
FROM mcr.microsoft.com/dotnet/sdk:5.0 AS build
```

```
WORKDIR /src
```

```
COPY ["CustomerApi/CustomerApi.csproj", "CustomerApi/"]
```

```
RUN dotnet restore "CustomerApi/CustomerApi.csproj"
```

```
COPY . .
```

```
WORKDIR "/src/CustomerApi"
```

```
RUN dotnet build "CustomerApi.csproj" -c Release -o /app/build
```

```
FROM build AS publish
```

```
RUN dotnet publish "CustomerApi.csproj" -c Release -o /app/publish
```

```
FROM base AS final
```

```
WORKDIR /app
```

```
COPY --from=publish /app/publish .
```

```
ENTRYPOINT ["dotnet", "CustomerApi.dll"]
```



```
FROM mcr.microsoft.com/dotnet/aspnet:5.0 AS base
```

```
WORKDIR /app
```

```
EXPOSE 80
```

```
EXPOSE 443
```

```
FROM mcr.microsoft.com/dotnet/sdk:5.0 AS build
```

```
WORKDIR /src
```

```
COPY ["CustomerApi/CustomerApi.csproj", "CustomerApi/"]
```

```
COPY ["CustomerApi.Data/CustomerApi.Data.csproj", "CustomerApi.Data/"]
```

```
COPY ["CustomerApi.Domain/CustomerApi.Domain.csproj", "CustomerApi.Domain/"]
```

```
COPY ["CustomerApi.Service/CustomerApi.Service.csproj", "CustomerApi.Service/"]
```

```
COPY ["CustomerApi.Messaging.Send/CustomerApi.Messaging.Send.csproj", "CustomerApi.Messaging.Send/"]
```

```
COPY ["CustomerApi.Database.Build/CustomerApi.Database.Build.csproj", "CustomerApi.Database.Build/"]
```

```
COPY ["Tests/CustomerApi.Test/CustomerApi.Test.csproj", "Tests/CustomerApi.Test/"]
```

```
COPY ["Tests/CustomerApi.Service.Test/CustomerApi.Service.Test.csproj", "Tests/CustomerApi.Service.Test/"]
```

```
COPY ["Tests/CustomerApi.Data.Test/CustomerApi.Data.Test.csproj", "Tests/CustomerApi.Data.Test/"]
```

```
COPY ["CustomerApi/nuget.config", ""]
```

```
COPY ["*.props", "."]
```

```
ARG PAT=localhost
```

```
RUN sed -i "s|</configuration>|<packageSourceCredentials><MicroserviceDemoNugets><add key=\"Username\" value=\"PAT\" /><add key=\"ClearTextPassword\" value=\"${PAT}\" /></MicroserviceDemoNugets></p>|g"
```

```
RUN dotnet restore "CustomerApi/CustomerApi.csproj" --configfile "./nuget.config"
```

```
RUN dotnet restore "CustomerApi.Database.Build/CustomerApi.Database.Build.csproj" --configfile "./nuget.config"
```

```
RUN dotnet restore "Tests/CustomerApi.Test/CustomerApi.Test.csproj" --configfile "./nuget.config"
```

```
RUN dotnet restore "Tests/CustomerApi.Service.Test/CustomerApi.Service.Test.csproj" --configfile "./nuget.config"
```

```
RUN dotnet restore "Tests/CustomerApi.Data.Test/CustomerApi.Data.Test.csproj" --configfile "./nuget.config"
```

```
COPY . .
```

```
RUN dotnet build "CustomerApi/CustomerApi.csproj" -c Release -o /app/build --no-restore
```

```
RUN dotnet build "Tests/CustomerApi.Test/CustomerApi.Test.csproj" -c Release --no-restore
```

```
RUN dotnet build "Tests/CustomerApi.Service.Test/CustomerApi.Service.Test.csproj" -c Release --no-restore
```

```
RUN dotnet build "Tests/CustomerApi.Data.Test/CustomerApi.Data.Test.csproj" -c Release --no-restore
```

```
FROM build AS dacpac
```

```
ARG BuildId=localhost
```

```
LABEL dacpac=${BuildId}
```

```
WORKDIR /src
```

```
RUN dotnet build "CustomerApi.Database.Build/CustomerApi.Database.Build.csproj" -c Release -o /dacpacs --no-restore
```

```
FROM build AS test
```

```
ARG BuildId=localhost
```

```
LABEL test=${BuildId}
```

```
RUN dotnet test --no-build -c Release --results-directory /testresults --logger "trx;LogFileName=test_results.trx" /p:CollectCoverage=true /p:CoverletOutputFormat=json%2cCobertura /p:CoverletOutputPath=test_results.trx
```

```
RUN dotnet test --no-build -c Release --results-directory /testresults --logger "trx;LogFileName=test_results2.trx" /p:CollectCoverage=true /p:CoverletOutputFormat=json%2cCobertura /p:CoverletOutputPath=test_results2.trx
```

```
RUN dotnet test --no-build -c Release --results-directory /testresults --logger "trx;LogFileName=test_results3.trx" /p:CollectCoverage=true /p:CoverletOutputFormat=json%2cCobertura /p:CoverletOutputPath=test_results3.trx
```

```
FROM build AS publish
```

```
RUN dotnet publish "CustomerApi/CustomerApi.csproj" --no-restore -c Release -o /app/publish
```

```
FROM base AS final
```

```
WORKDIR /app
```

```
COPY --from=publish /app/publish .
```

```
ENTRYPOINT ["dotnet", "CustomerApi.dll"]
```

57 Layer

Inspect Layers

Docker history <IMAGE_ID>

Dive: <https://github.com/wagoodman/dive>

1: Terminal ▾

□ ×

```
[wagoodman@kiwi dive] ➤ master $ dive someproj:latest
```

Tags

Decide what version you run at any given time

“Latest” by default

Used for versioning

Tags

Decide what version you run at any given time

“Latest” by default

Used for versioning

Tag	:latest	
	:1	
	:1.0	:1.1
Digest	91efj6	u82lq



Container Registry

Repository to store container images

Docker Hub

Filters

Products

- ☐ Images
- ☐ Extensions
- ☐ Plugins

Trusted Content

- ☐ Docker Official Image
- ☐ Verified Publisher
- ☐ Sponsored OSS

Operating Systems

- ☐ Linux
- ☐ Windows

Architectures

- ☐ ARM
- ☐ ARM 64

1 - 15 of 15 results for wolfgangofner.

Best Match



wolfgangofner/microservicedemo • 100K+ • 0

By [wolfgangofner](#) • Updated 3 years ago

Linux x86-64



wolfgangofner/customerapi • 3.3K • 0

By [wolfgangofner](#) • Updated 17 days ago

Image for my NET 6 Microservice demo.

Linux x86-64



wolfgangofner/kubernetesdeploymentdemo • 2.1K • 0

By [wolfgangofner](#) • Updated 4 years ago

Linux x86-64



wolfgangofner/orderapi • 1.5K • 0

Container Registry

Repository to store container images

Docker Hub

Public vs. private registry

Azure Container Registry (ACR)

Additional functionalities like:

- Geo-replication
- Retention Policies
- Security scanning

Docker Compose

YAML file

Define container dependencies

Run all dependent containers

Docker Compose

YAML file

Define container dependencies

Run all dependent containers

Advantages

- Configure dependencies between containers
- Restart policy
- Easy to start

Disadvantages

- Monitoring
- Load Balancing
- Deployment
- SSL Certificate

Docker Compose

YAML file

Define container dependencies

Run all dependent containers

Advantages

- Configure dependencies between containers
- Restart policy
- Easy to start

Disadvantages

- Monitoring
- Load Balancing
- Deployment
- SSL Certificate

```
version: "3.9"
```

```
services:
```

```
  wordpress:
```

```
    image: wordpress
```

```
    restart: always
```

```
    ports:
```

```
      - 8080:80
```

```
    environment:
```

```
      WORDPRESS_DB_HOST: db
```

```
      WORDPRESS_DB_USER: exampleuser
```

```
      WORDPRESS_DB_PASSWORD: examplepass
```

```
      WORDPRESS_DB_NAME: exampledb
```

```
    volumes:
```

```
      - wordpress:/var/www/html
```

```
  db:
```

```
    image: mysql:5.7
```

```
    restart: always
```

```
    environment:
```

```
      MYSQL_DATABASE: exampledb
```

```
      MYSQL_USER: exampleuser
```

```
      MYSQL_PASSWORD: examplepass
```

```
      MYSQL_RANDOM_ROOT_PASSWORD: '1'
```

```
    volumes:
```

```
      - db:/var/lib/mysql
```

Docker Recap

Small images

Fast start up and deployment

Reusable and portable

Immutable → “Works on my machine”

Containers allow you to run your software even if your infrastructure provider does not support it

Docker Commands

List running containers

```
docker ps
```

List images

```
docker image ls
```

Download an image from a registry

```
docker pull wolgangofer/customerapi
```

Build an image from a Dockerfile

```
docker build . [-f CustomerApi/Dockerfile]
```

Tag an image

```
docker tag customerapi wolgangofer/customerapi
```

Push an image to a registry

```
docker push wolgangofer/customerapi
```

Start a container

```
docker run -p 32789:80 -p 32788:443 wolgangofer/customerapi
```



Docker Compose Demo



Exercise

Docker Exercise

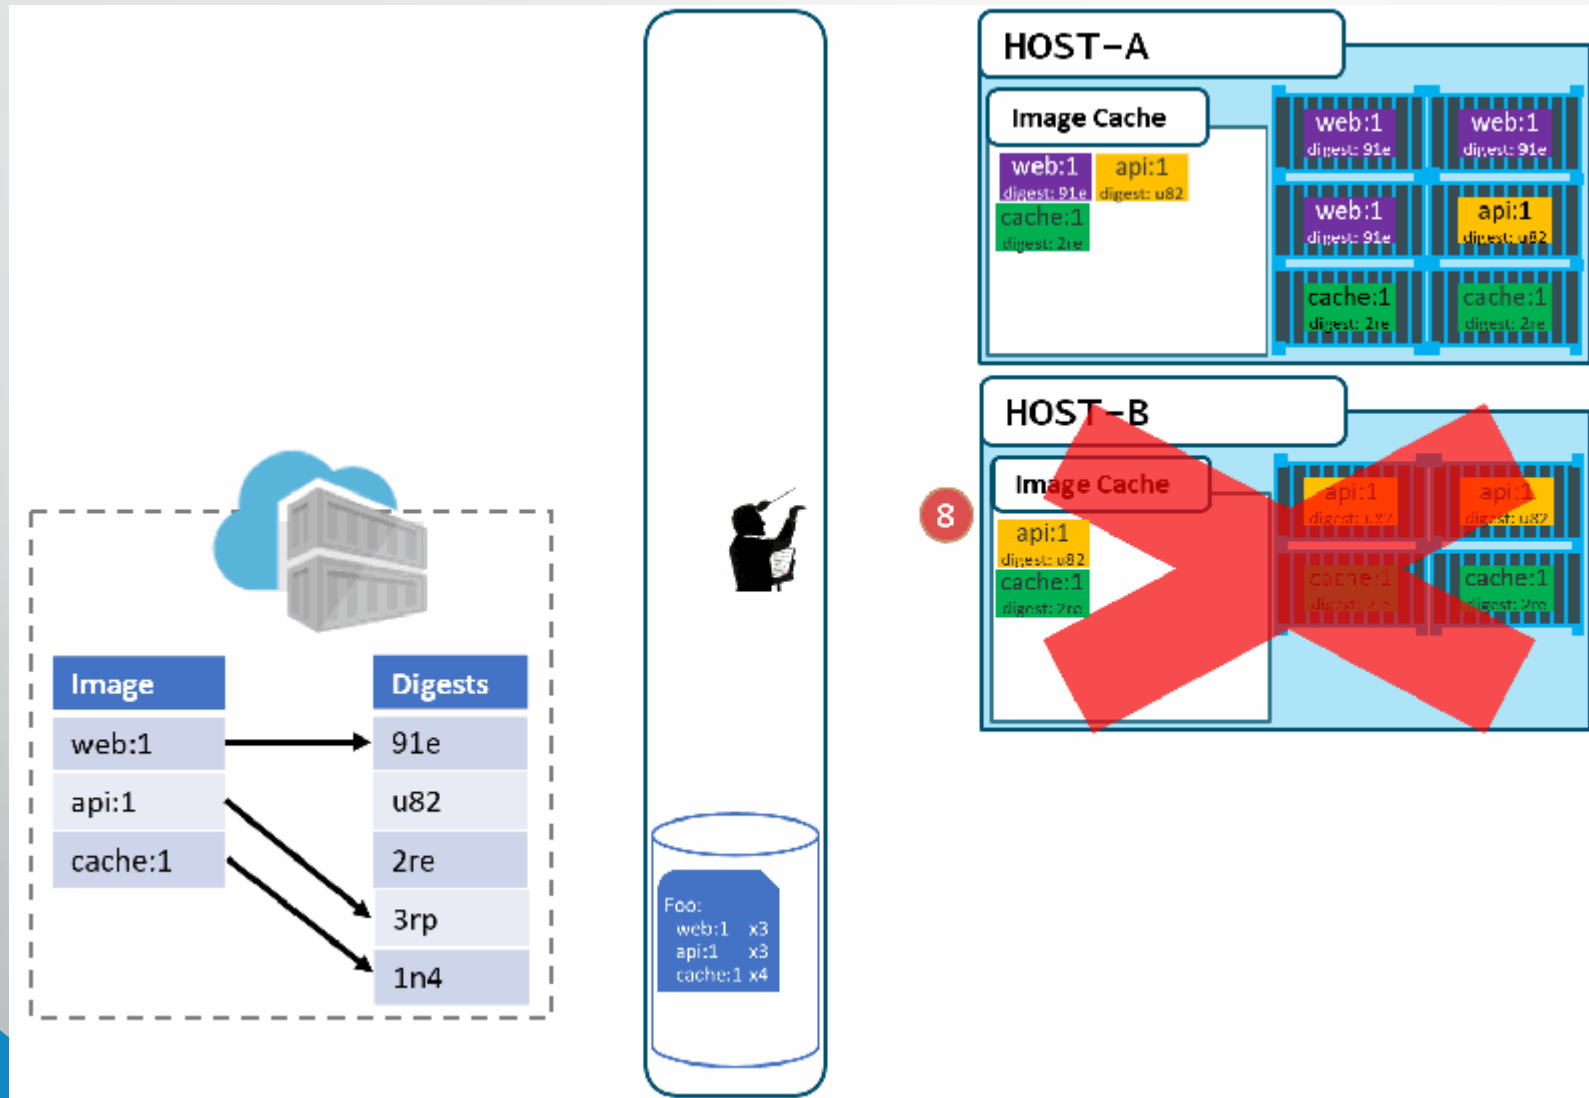
- Run an image from Dockerhub
- Create a new application and build it in a Dockerfile
- Upload your image to Dockerhub
- Build and run a docker-compose file
- Use different docker commands to interact with images



containers don't solve all problems



Container Orchestrator





Container Orchestrator

Multi-Node Management

Resource Management

Load balancing

Monitoring and Self Healing

Zero Downtime Deployments

Manage SSL certificates



Container Orchestrator

Multi-Node Management

Resource Management

Load balancing

Monitoring and Self Healing

Zero Downtime Deployments

Manage SSL certificates

Kubernetes

Docker Swarm

Marathon



Container Orchestrator

Multi-Node Management

Resource Management

Load balancing

Monitoring and Self Healing

Zero Downtime Deployments

Manage SSL certificates

Kubernetes

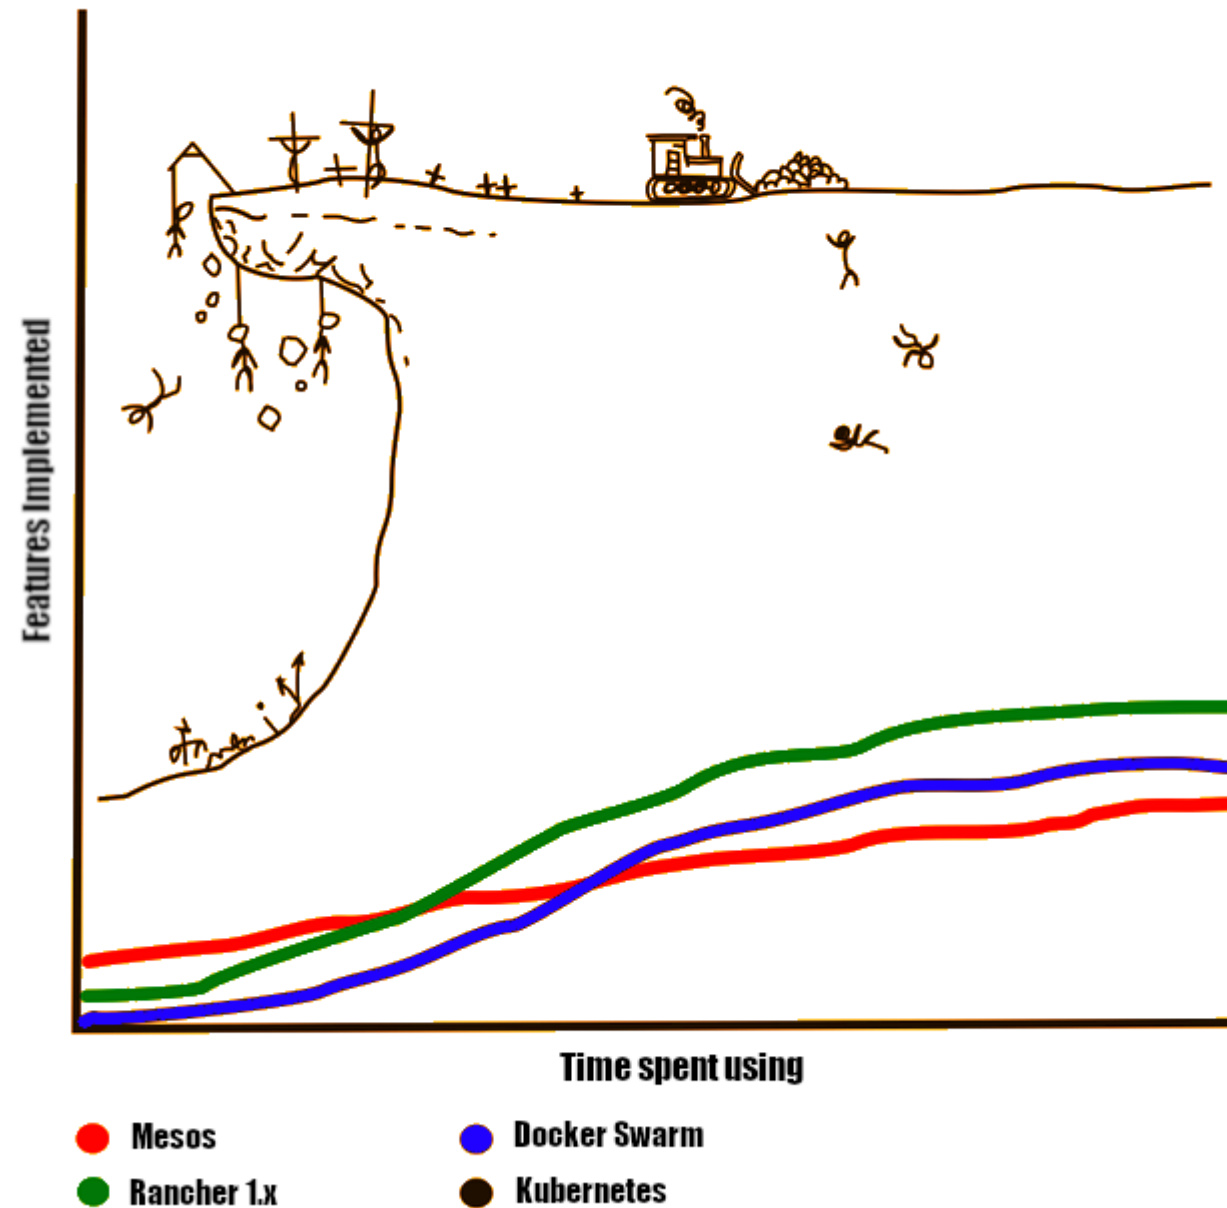
Docker Swarm


Marathon



kubernetes

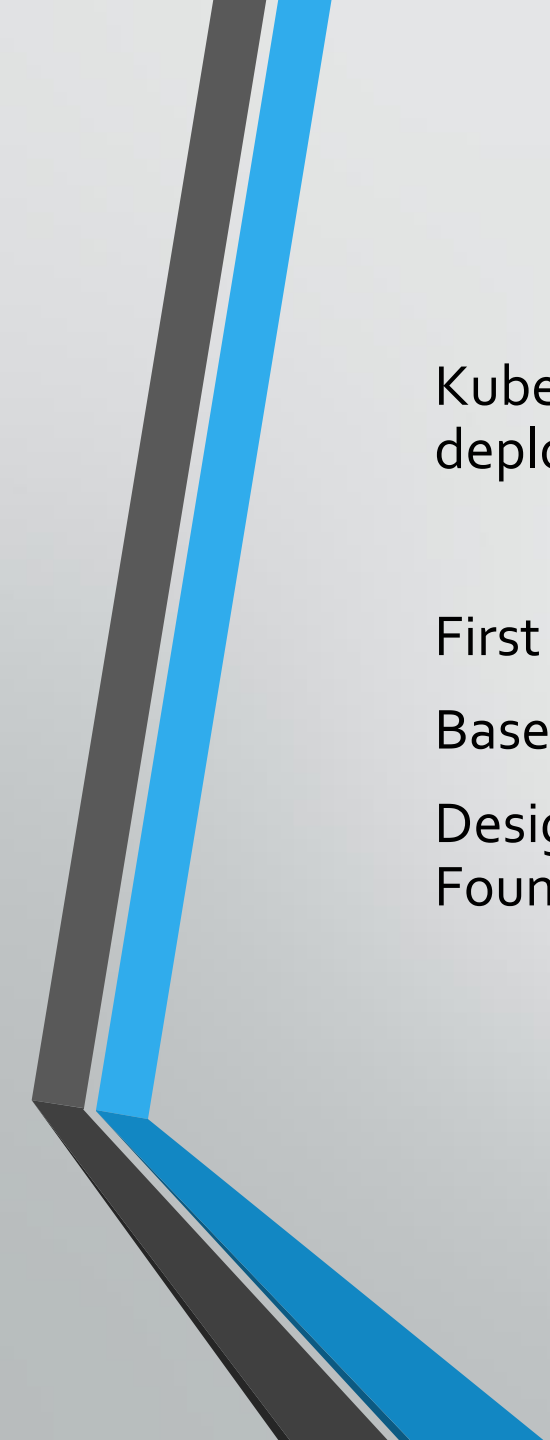
Learning curves of some Container Orchestration Engines





Kubernetes

Kubernetes is an open-source system for automating computer application deployment, scaling, and **management of container applications**



Kubernetes

Kubernetes is an open-source system for automating computer application deployment, scaling, and **management of container applications**

First Release on 10 July 2015

Based on Google's Borg

Designed by Google and is now maintained by the Cloud Native Computing Foundation

Kubernetes

Kubernetes is an open-source system for automating computer application deployment, scaling, and **management of container applications**

First Release on 10 July 2015

Based on Google's Borg

Designed by Google and is now maintained by the Cloud Native Computing Foundation

Written in Go

Open-source

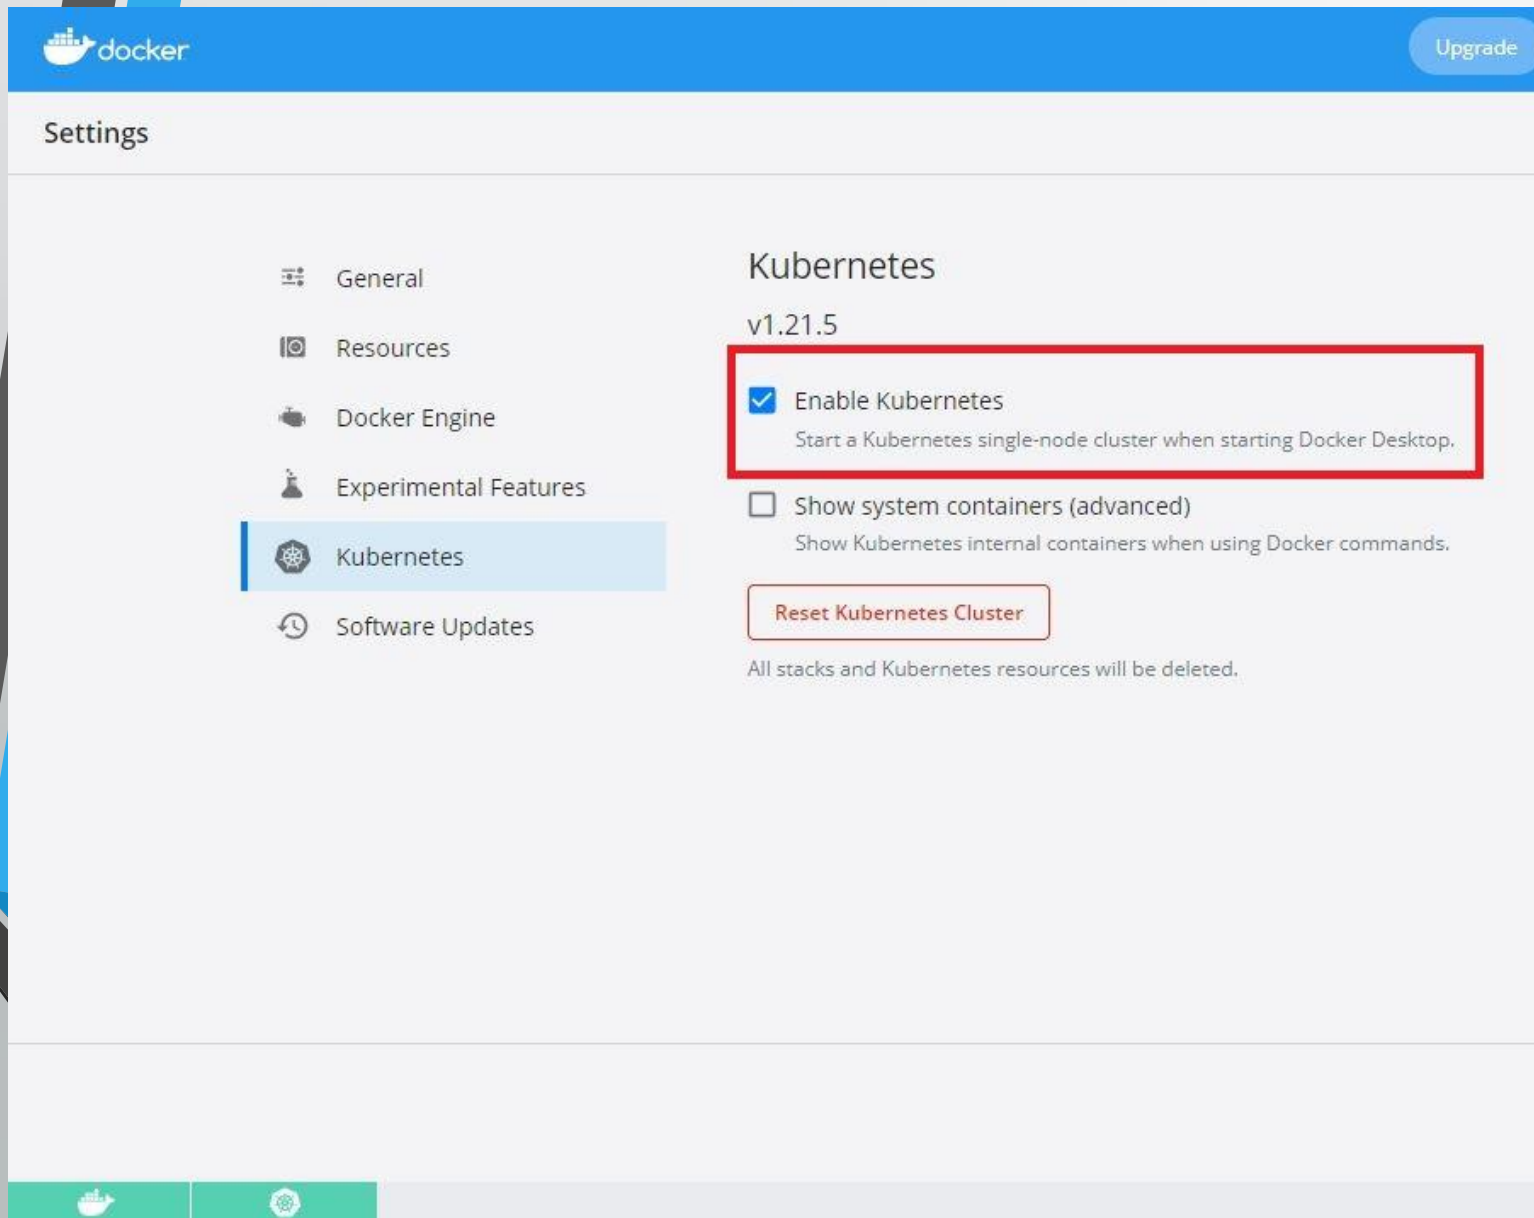
"K8s" → K-8 character-s

Kubernetes

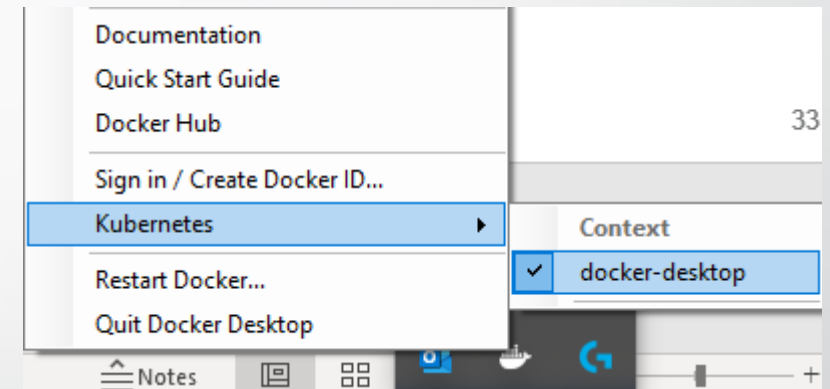
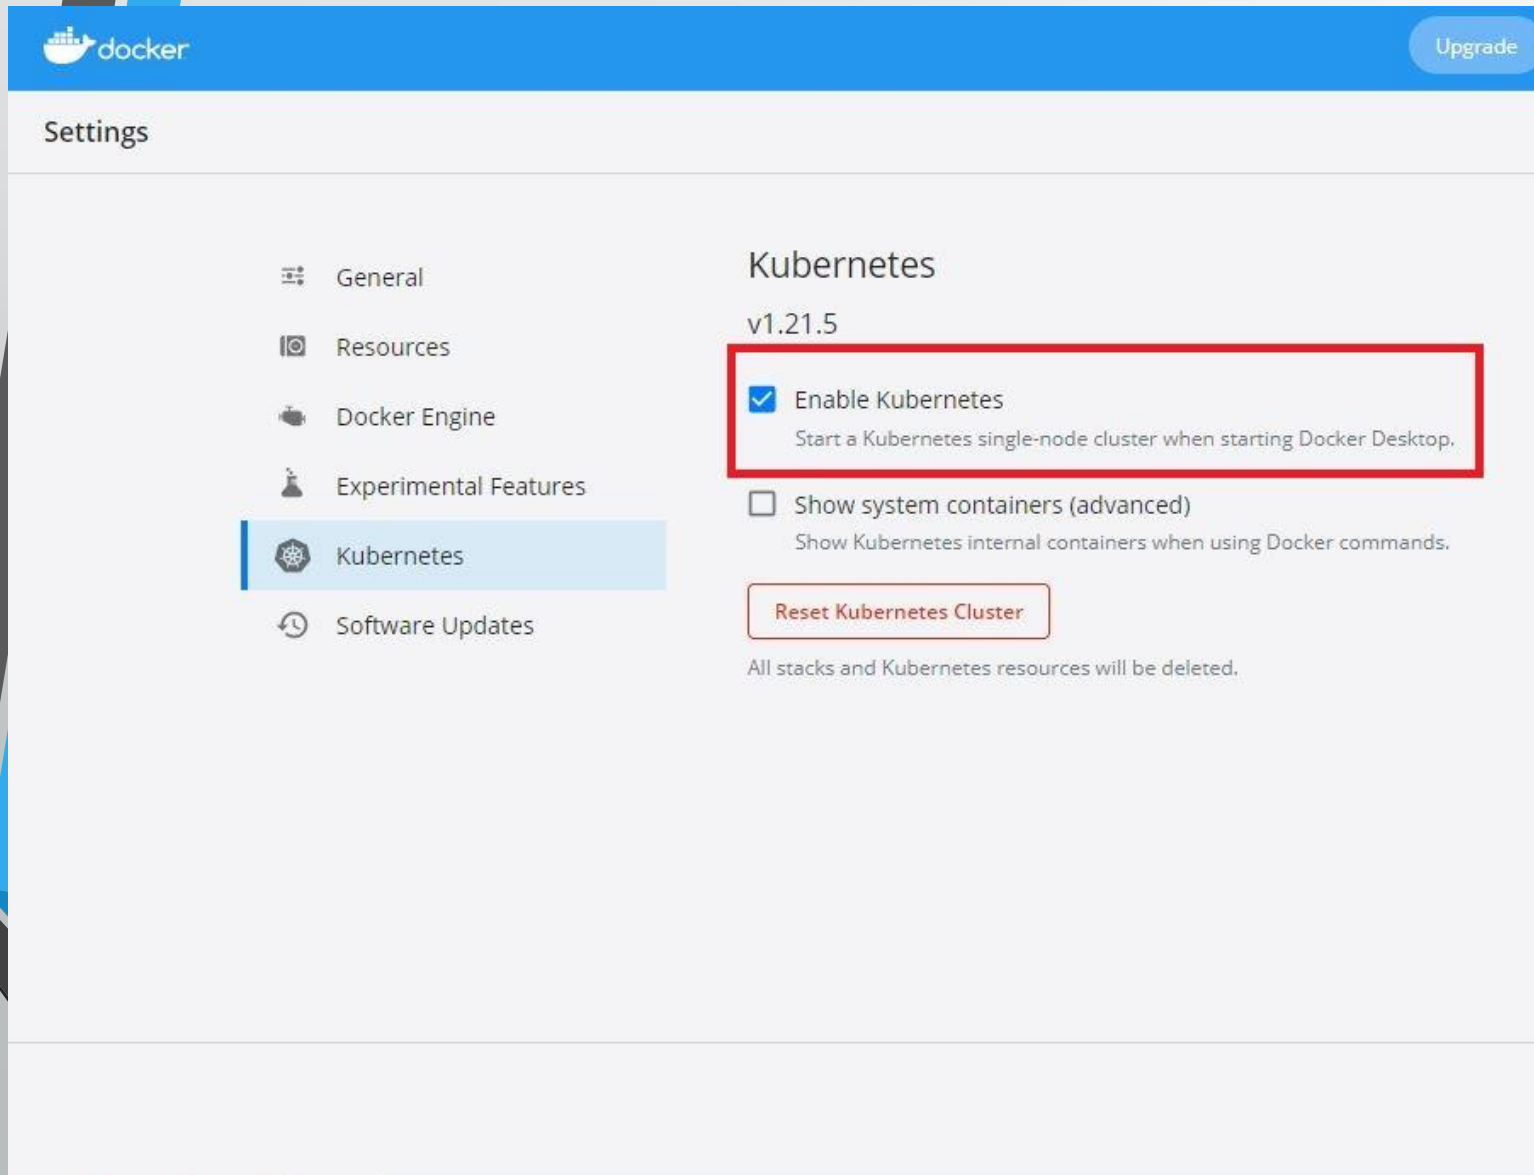
Multiple distribution

- K3s
- MicroK8s
- RedHat OpenShift
- Mirantis
- Azure Kubernetes Service
- Amazon Elastic Kubernetes
- Google Kubernetes Engine

K8s locally



K8s locally



K8s locally

K3s:

- Lightweight
- Great for Edge and IoT
- Easy to install

K8s locally

K3s:

- Lightweight
- Great for Edge and IoT
- Easy to install

```
curl -sfL https://get.k3s.io | sh -
```

Kubernetes Features

Self-healing

Service discovery and load balancing

Secret and configuration management

Horizontal scaling

Zero downtime deployments

Batch execution

Namespaces

Easily extensible

Configuration in JSON or YAML



Self-healing Demo

Kubernetes Components

Master Node (Control Plane)

- kube-apiserver
- etcd
- kube-scheduler
- kube-control-manager
- Master Node is managed by cloud vendor

Kubernetes Components

Master Node (Control Plane)

- kube-apiserver
- etcd
- kube-scheduler
- kube-control-manager
- Master Node is managed by cloud vendor

Worker Node

- kubelet
- kube-proxy
- Container runtime

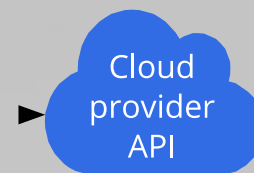
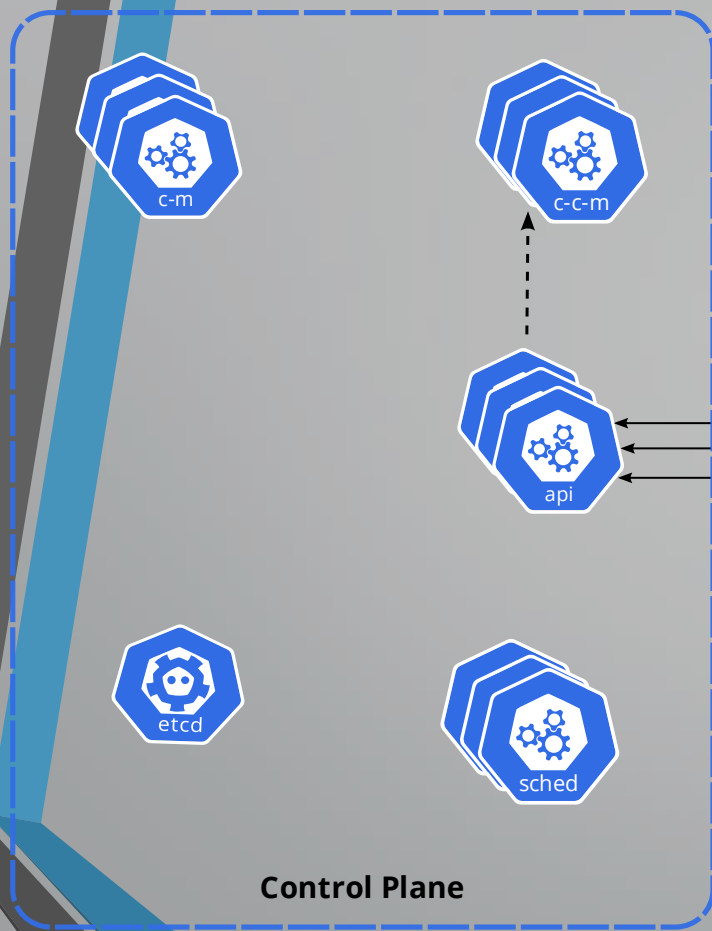
Kubernetes Components

Addons

- DNS
- Networking
- Storage
- Dashboard
- ...

Kubernetes Components

Kubernetes cluster



API server



Cloud controller manager
(optional)

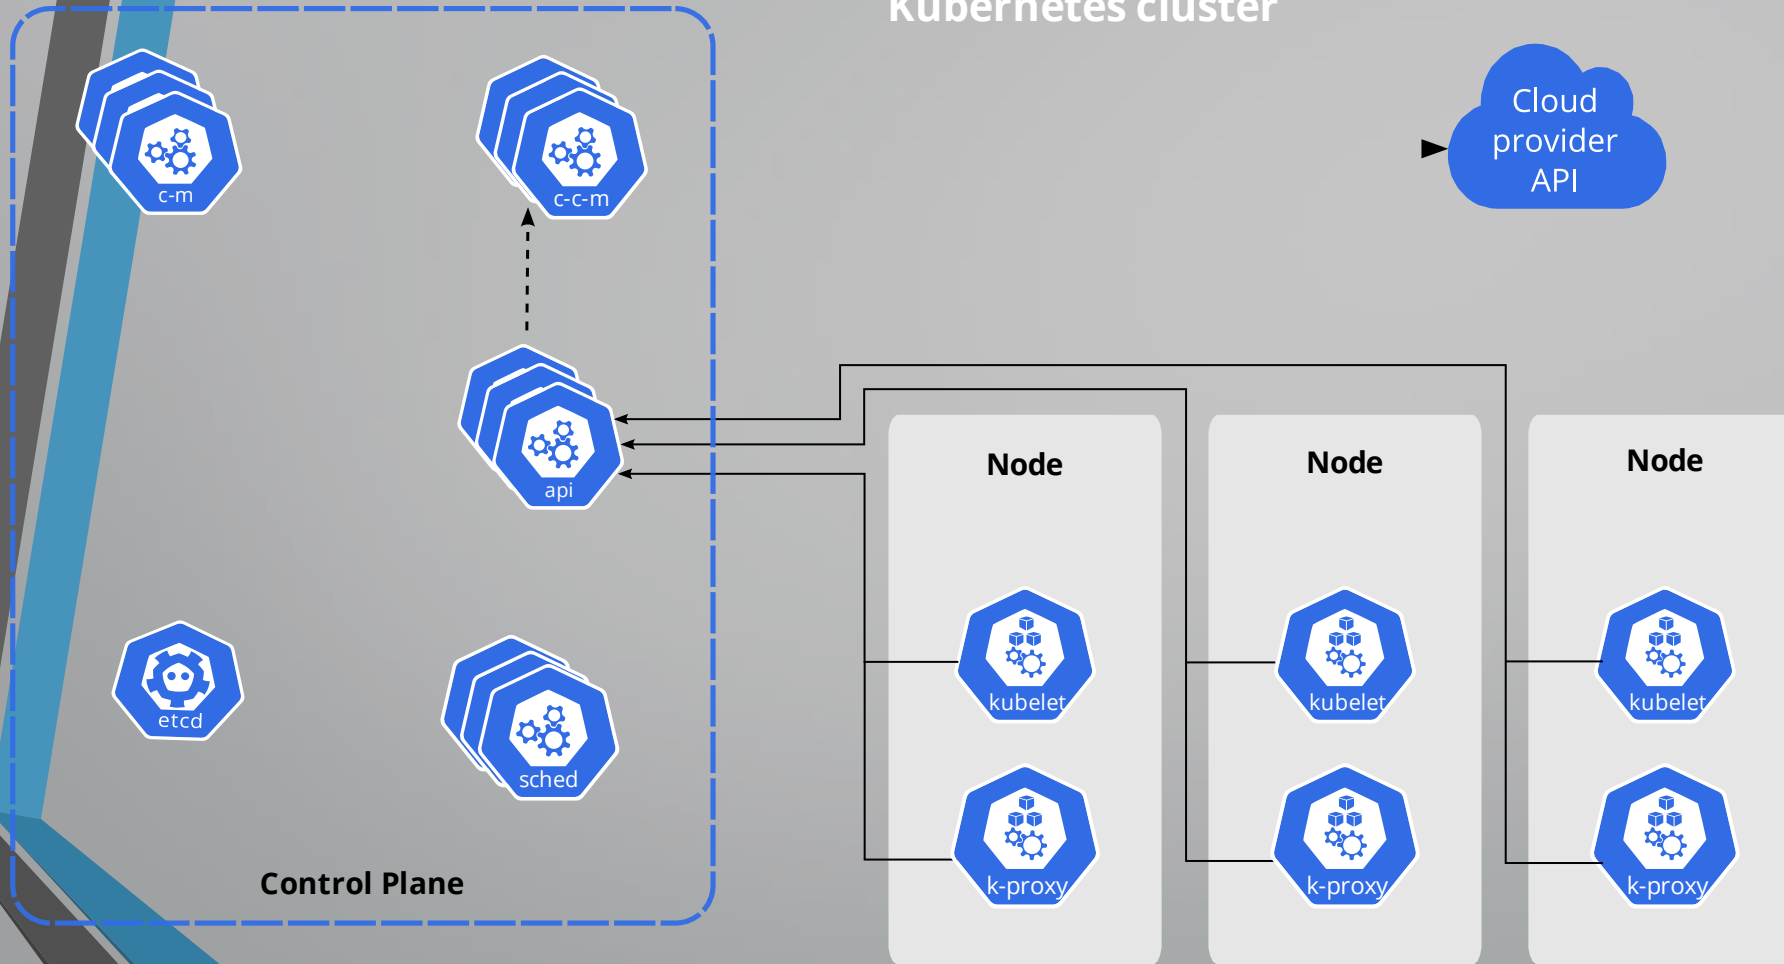


Controller manager



Kubernetes Components

Kubernetes cluster



API server



Cloud controller manager
(optional)



Controller manager



etcd
(persistence store)



kubelet



kube-proxy

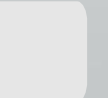


Scheduler



Control plane

Node



Pod

A pod is the smallest unit in K8s

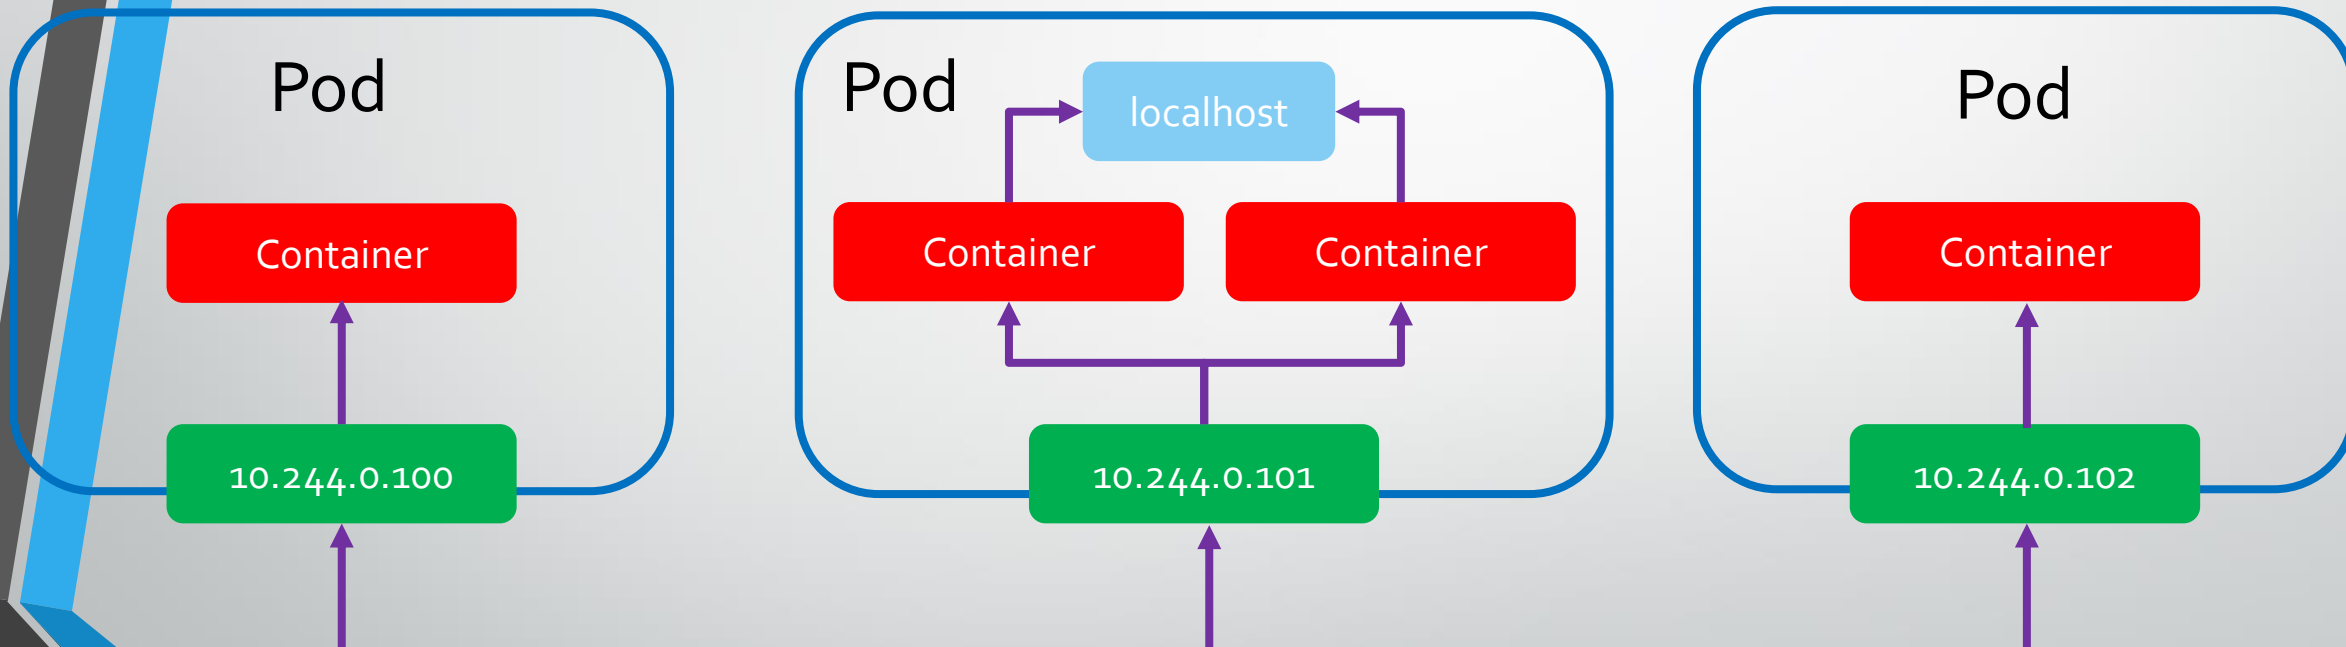
Pods wrap one or more containers

Provides a way to set environment variables and mount storage

Containers inside a pod can communicate via localhost

Multiple containers should only be combined in a pod if they are interdependent

Pods and Containers





Kubernetes Networking

The default network IP range for pods is 10.244.0.0/16

Kubernetes Networking

The default network IP range for pods is 10.244.0.0/16

Network CIDR range can be configured with addon:

- Cilium
- Flannel
- Calico



Namespaces

Used to create virtual cluster inside a physical cluster

Namespaces

Used to create virtual cluster inside a physical cluster

- Isolation
- Resource Segregation
- Multiple Environments
- Resource Quotas
- Default Namespace

Kubernetes Configuration

Declarative Model and Desired State

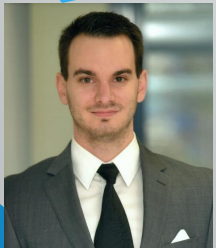
- Tell Kubernetes what you want
- Kubernetes will figure out a way to get to the desired state
- Etcd holds the current status of any K8s component

Kubernetes Configuration

Declarative Model and Desired State

- Tell Kubernetes what you want
- Kubernetes will figure out a way to get to the desired state
- Etcd holds the current status of any K8s component

Hey Kubernetes, run 3 pods of
wolfgangofner/customerapi



Kubernetes Configuration

Declarative Model and Desired State

- Tell Kubernetes what you want
- Kubernetes will figure out a way to get to the desired state
- Etcd holds the current status of any K8s component

Hey Kubernetes, run 3 pods of
wolfgangofner/customerapi



Let me check if I am already
running your pods



Kubernetes Configuration

Declarative Model and Desired State

- Tell Kubernetes what you want
- Kubernetes will figure out a way to get to the desired state
- Etcd holds the current status of any K8s component

Hey Kubernetes, run 3 pods of
wolfgangofner/customerapi



Let me check if I am already
running your pods



Currently there is one pod of
wolfgangofner/customerapi
running

Kubernetes Configuration

Declarative Model and Desired State

- Tell Kubernetes what you want
- Kubernetes will figure out a way to get to the desired state
- Etcd holds the current status of any K8s component

Hey Kubernetes, run 3 pods of
wolfgangofner/customerapi



Let me check if I am already
running your pods



Currently there is one pod of
wolfgangofner/customerapi
running

Starting two more pods of
wolfgangofner/customerapi

Kubernetes Configuration

Declarative Model and Desired State

- Tell Kubernetes what you want
- Kubernetes will figure out a way to get to the desired state
- Etcd holds the current status of any K8s component

Configuration Handling

- YAML or JSON files
- Kubernetes CLI called kubectl
- Kubectl communicates with the Kubernetes API

Kubernetes Configuration

Declarative Model and Desired State

- Tell Kubernetes what you want
- Kubernetes will figure out a way to get to the desired state
- Etcd holds the current status of any K8s component

Configuration Handling

- YAML or JSON files
- Kubernetes CLI called kubectl
- Kubectl communicates with the Kubernetes API

kubectl



Kube Control

Kube Cuddle

```
apiVersion: v1
kind: Service
metadata:
  name: kubernetesdemo-service
spec:
  type: LoadBalancer
  selector:
    app: kubernetesdemo
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubernetesdemo-deployment
  labels:
    app: kubernetesdemo
spec:
  replicas: 1
  selector:
    matchLabels:
      app: kubernetesdemo
  template:
    metadata:
      labels:
        app: kubernetesdemo
    spec:
      containers:
        - name: kubernetesdemo
          image: wolfgangofner/kubernetesdeploymentdemo:start
          ports:
            - containerPort: 80
```

Labels and Annotations

Labels

- Key value pairs that are bound to objects like deployments or pods with a maximum of 63 character
- app:MyAppName
- Used to filter or select objects
- Can be changed or deleted at any times

Labels and Annotations

Labels

- Key value pairs that are bound to objects like deployments or pods with a maximum of 63 character
- app:MyAppName
- Used to filter or select objects
- Can be changed or deleted at any times

Annotations

- Also key value pairs but without the character limitation
- Can not be used for filtering or selecting objects

Labels and Annotations

Labels

- Key value pairs that are bound to objects like deployments or pods with a maximum of 63 character
- app:MyAppName
- Used to filter or select objects
- Can be changed or deleted at any times

Annotations

- Also key value pairs but without the character limitation
- Can not be used for filtering or selecting objects

```
metadata:  
  creationTimestamp: "2021-10-17T11:58:22Z"  
  labels:  
    component: apiserver  
    provider: kubernetes
```

Metadata

Age 51m

Labels

component:apiserver

provider:kubernetes



Services

Pods come and go

IP addresses will change

Services

Pods come and go

IP addresses will change

Service stay for the entire lifetime of the application

Persistent entry point

Fixed IP address

Load Balancing



Load Balancer Demo

Services

Pods come and go

IP addresses will change

Service stay for the entire lifetime of the application

Persistent entry point

Fixed IP address

Load Balancing

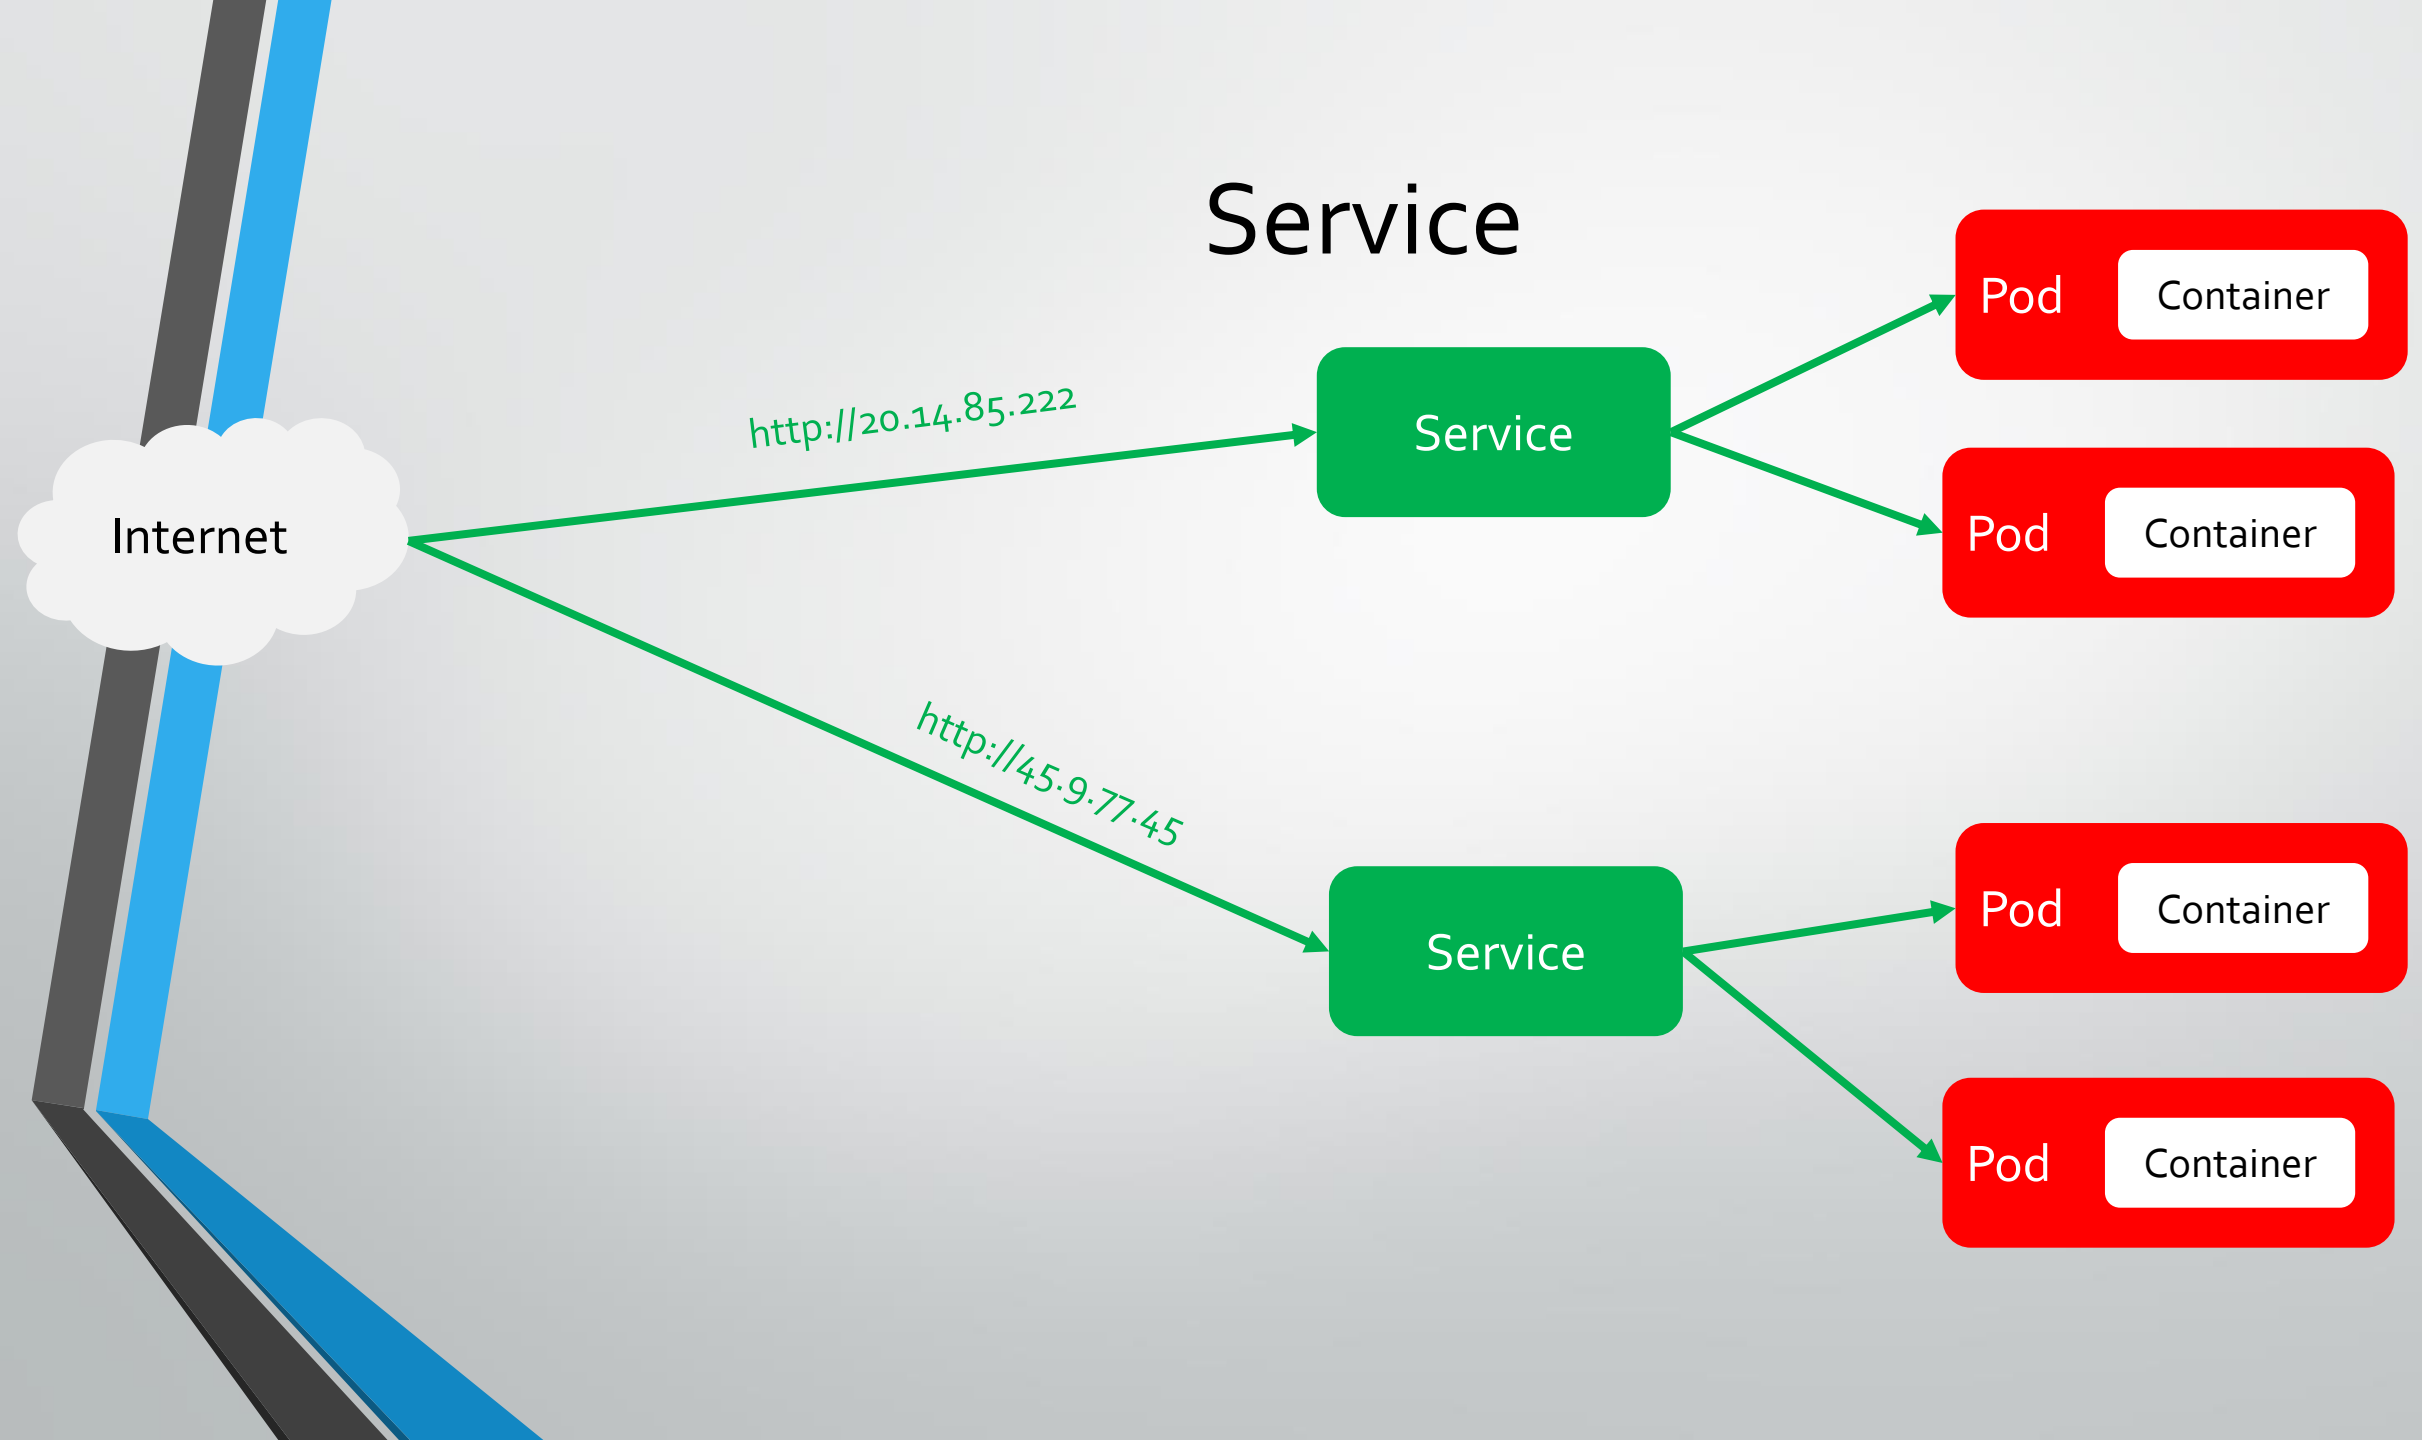
Pods and Services are matched using Labels

Services

Four types of Services

- LoadBalancer: exposes the Service using an external load balancer
- ClusterIP: makes the Services accessible only from within the cluster
- NodePort: exposes the Service at each node's IP at a static port
- ExternalName: maps the service to an existing DNS FQDN

Service



Swagger UI

Not secure | 20.103.233.94/index.html

swagger

Select a spec

My API V1

kubernetesdemo-deployment-599d8f48c-gl4t9^{v1}

/swagger/v1/swagger.json

A collection of Web APIs

Values

GET	/api/Values
POST	/api/Values
GET	/api/Values/{id}
PUT	/api/Values/{id}
DELETE	/api/Values/{id}

Kubectl Commands

Get resource

kubectl get pods/service/deployment

Delete resource

kubectl delete pod/service/deployment

Display information about resource

kubectl describe pod/node/service resource-name

Add/update new resource

kubectl apply -f myfile.yaml [namespace=my-namespace]

Set current namespace

kubectl config set-context --current --namespace=my-namespace

Kubernetes Cheat Sheet: <https://kubernetes.io/docs/reference/kubectl/cheatsheet>



Exercise

Exercise

Get Kubernetes up and running

Create a new namespace

```
kubectl create ns my-namespace
```

Apply provided YAML file to your namespace

```
kubectl apply --f filepath --namespace=my-namespace
```

```
kubectl get all --n my-namespace
```

Set Replicas to 3 and check what happens

Replace image in Deployment





Lunch Break



Recap

Questions about what we have learned so far?

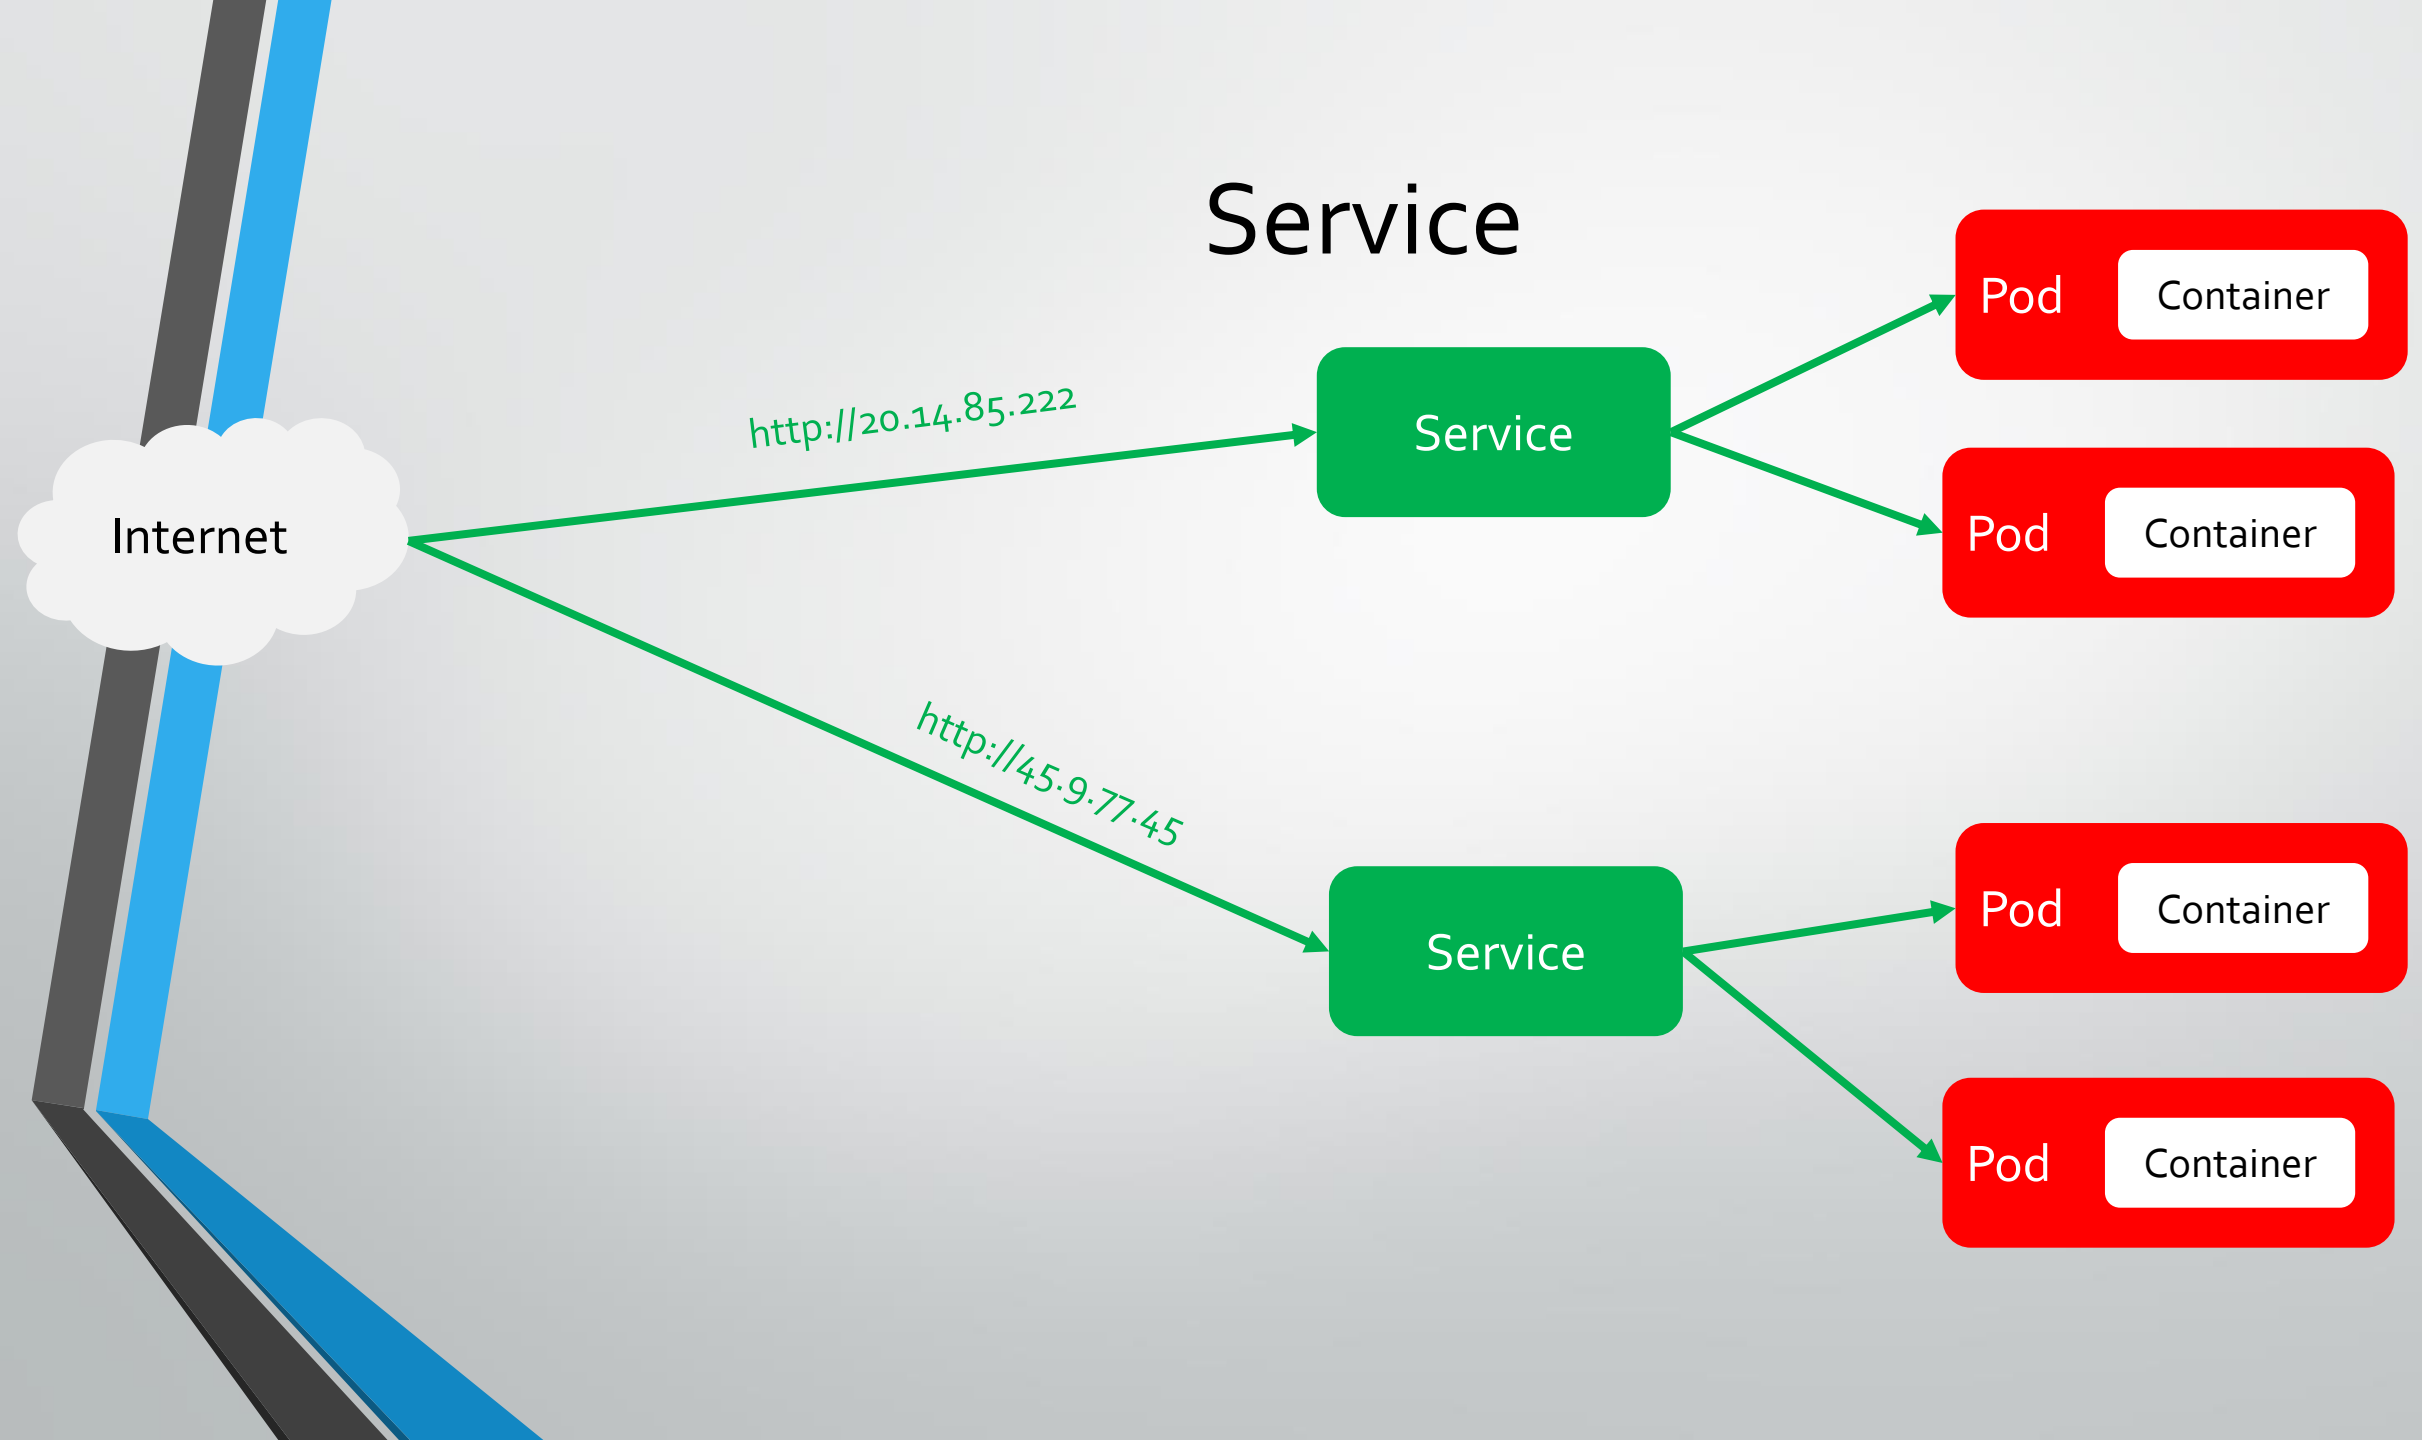
Recap

Containers are immutable objects

Pods run containers in Kubernetes

Services are an entry point and route requests to pods

Service



Swagger UI

Not secure | 20.103.233.94/index.html

swagger

Select a spec

My API V1

kubernetesdemo-deployment-599d8f48c-gl4t9^{v1}

/swagger/v1/swagger.json

A collection of Web APIs

Values

GET

/api/Values

POST

/api/Values

GET

/api/Values/{id}

PUT

/api/Values/{id}

DELETE

/api/Values/{id}

Recap

Containers are immutable objects

Pods run containers in Kubernetes

Services are an entry point and route requests to pods

Needed for production environment:

- Secret Management
- Health Checks
- SSL Certificates
- Deployments
- Resource Management

Secrets

Base64 encoded

Automatically decrypted when attached to pod

Can be used in config file or environment variable

Secrets

Base64 encoded

Automatically decrypted when attached to pod

Can be used in config file or environment variable

Config and Storage > Secrets > kedademoapi-tls

kedademoapi-tls

Summary

Metadata

Resource Viewer

YAML

```
1  ---
2  apiVersion: v1
3  data:
4    tls.crt: LS0tLS1CRUdJTiBDRVJUSUZJJQ0FURS0tLS0tCk1JSUZXXVENDQkVHZ0F3SUJBZ01T
5    tls.key: LS0tLS1CRUdJTiBSU0EgUFJJVkFURSBLRVktLS0tLQpNSU1Fb2dJJQkFBS0NBUEVB
6  kind: Secret
```

Cert-Issuer

Kubernetes resource

Handles certificate requests

Supported sources:

- Let's Encrypt
- HashiCorp Vault
- Venafi
- private PKI (Public Key Infrastructure)

Cert-Issuer

Kubernetes resource

Handles certificate requests

Supported sources:

- Let's Encrypt
- HashiCorp Vault
- Venafi
- private PKI (Public Key Infrastructure)

```
apiVersion: cert-manager.io/v1
kind: ClusterIssuer
metadata:
  name: letsencrypt
spec:
  acme:
    server: https://acme-v02.api.letsencrypt.org/directory
    email: <Your Email>
    privateKeySecretRef:
      name: letsencrypt
    solvers:
      - http01:
          ingress:
            class: nginx
            podTemplate:
              spec:
                nodeSelector:
                  "kubernetes.io/os": linux
```

Cert-Manager

Manages obtaining and renewing of certificates

Can use variety of CAs like Let's Encrypt, HashiCorp Vault, and Venafi

Updates certificates at a configured time before expiry

Uses Cert Issuer to issue certificates

Cert-Manager

Issuers

letsencrypt

venafi-tpp

hashicorp-vault

Cert-Manager

Certificates

example.com
Issuer: letsencrypt

foo.bar.com
Issuer: hashicorp-vault

Kubernetes
Secrets

Signed keypair

Signed keypair

kedademoapi-tls

[Summary](#)[Metadata](#)[Resource Viewer](#)[YAML](#)

```
1  ---
2  apiVersion: v1
3  data:
4    tls.crt: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUZXXVENDQkVHZ0F3SUJBZ01TQSc
5    tls.key: LS0tLS1CRUdJTiBSU0EgUFJJVkFURSBLRVktLS0tLQpNSU1Fb2dJQkFBS0NBUEUvBNG
6  kind: Secret
7  metadata:
8    annotations:
9      cert-manager.io/alt-names: test.kedademo.programmingwithwolfgang.com
10     cert-manager.io/certificate-name: kedademoapi-tls
11     cert-manager.io/common-name: test.kedademo.programmingwithwolfgang.com
12     cert-manager.io/ip-sans: ""
13     cert-manager.io/issuer-group: cert-manager.io
14     cert-manager.io/issuer-kind: ClusterIssuer
15     cert-manager.io/issuer-name: letsencrypt
16     cert-manager.io/uri-sans: ""
17  creationTimestamp: "2021-10-17T12:07:46Z"
```



Your information (for example, passwords or credit card numbers) is private when it is sent to this site.

[Learn more](#)

Cu

/swagger

A simp

Wolfgang

Send e

Customer

GET

/v1/Customers Action to see all existing customers.

POST

/v1/Customers Action to create a new customer in the database.

PUT

/v1/Customers Action to update an existing customer



Ingress Controller

Entry point into the cluster

Processes HTTPS traffic

Reverse proxy redirects requests to Application

Ingress Controller

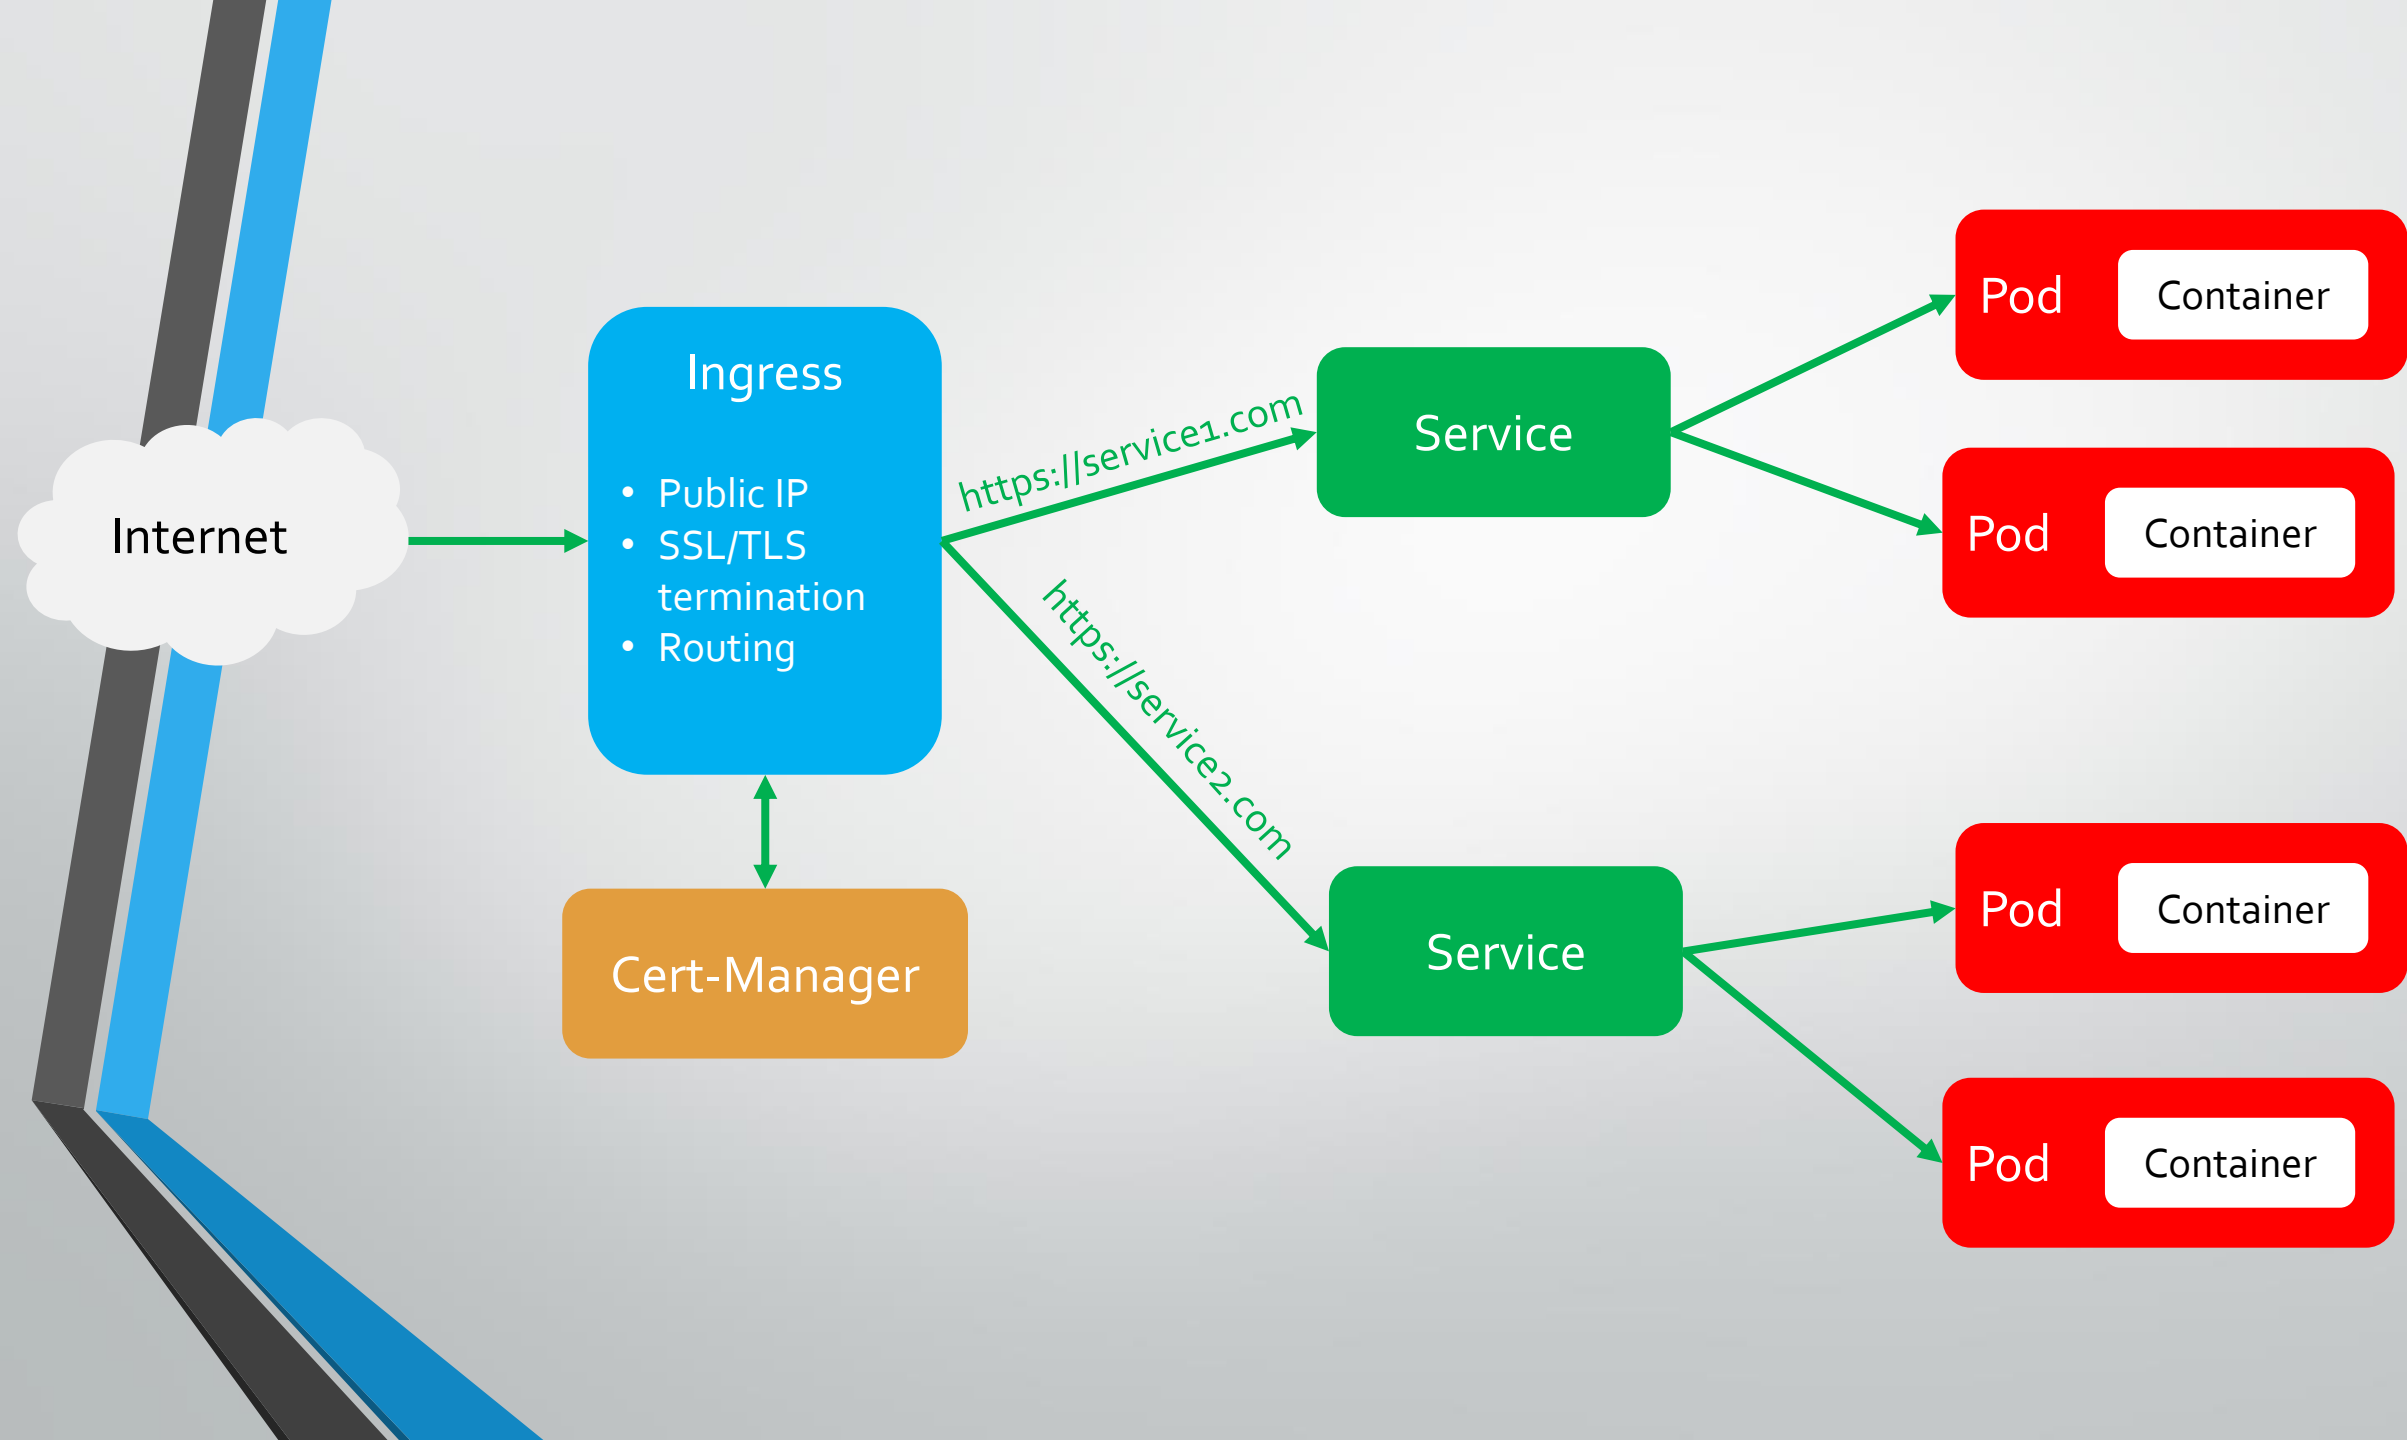
Entry point into the cluster

Processes HTTPS traffic

Reverse proxy redirects requests to Application

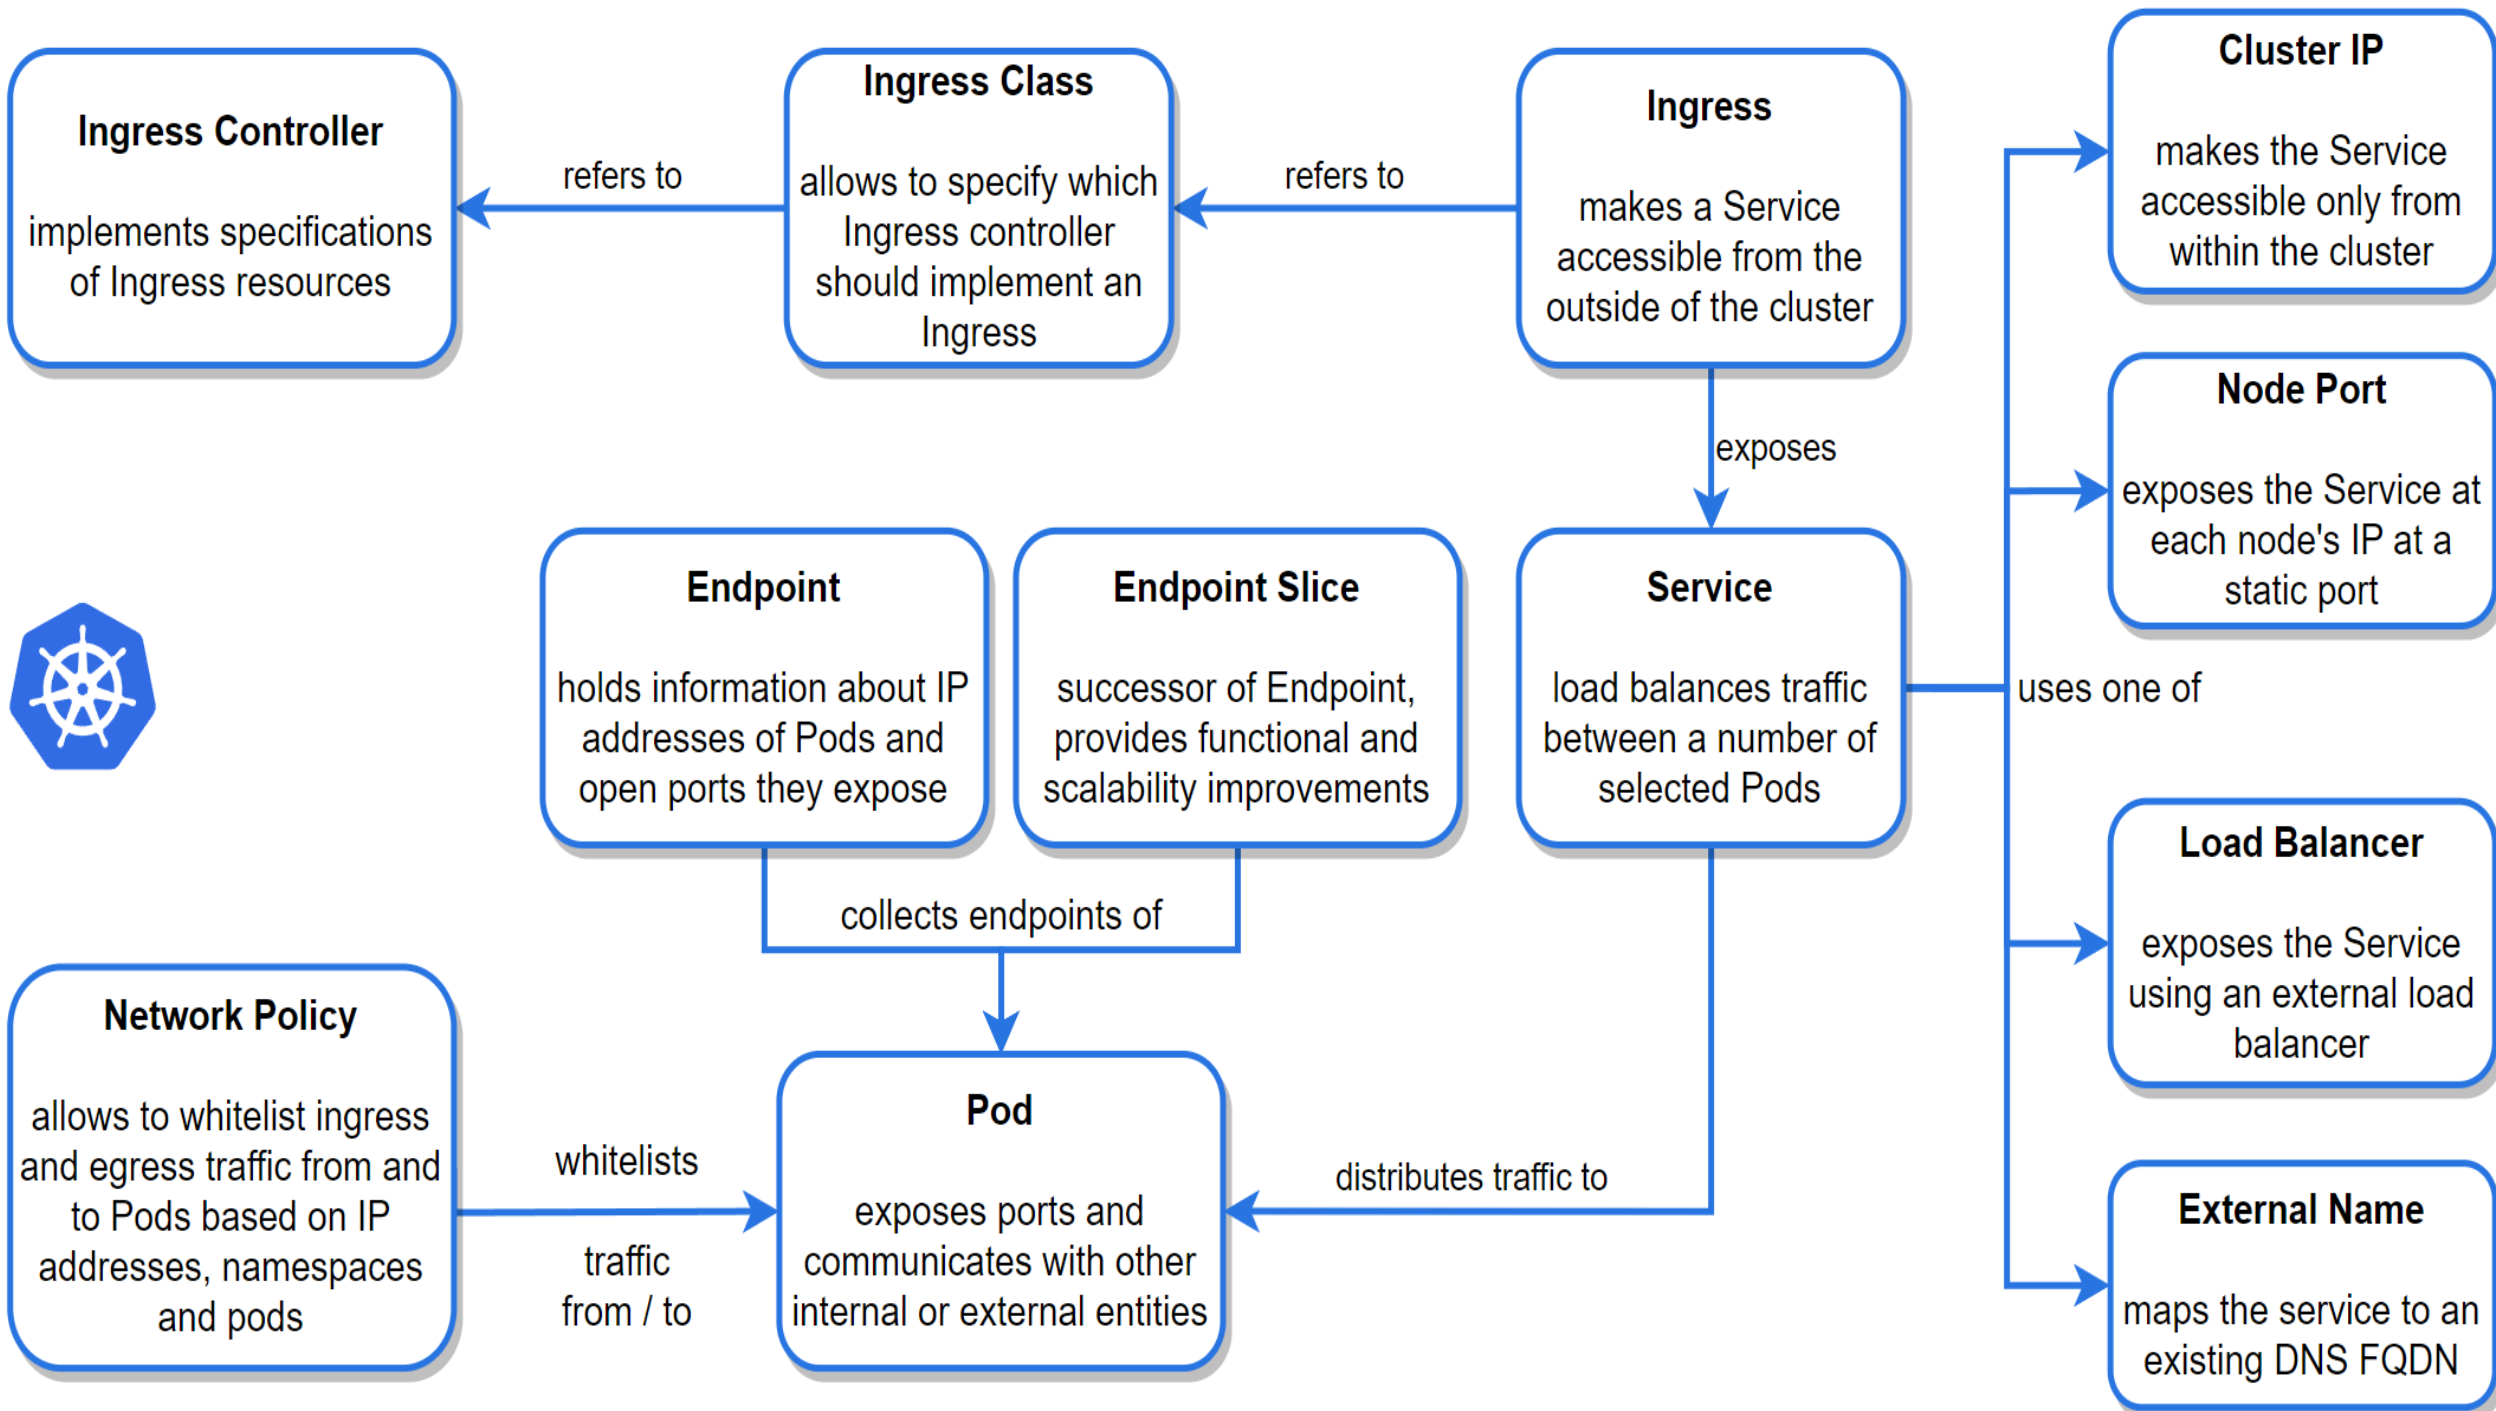
Popular Ingress Controller:

- NGINX
- Traefik
- HAProxy





Pull Request Demo



Pod Deployment

Pods are not directly deployed

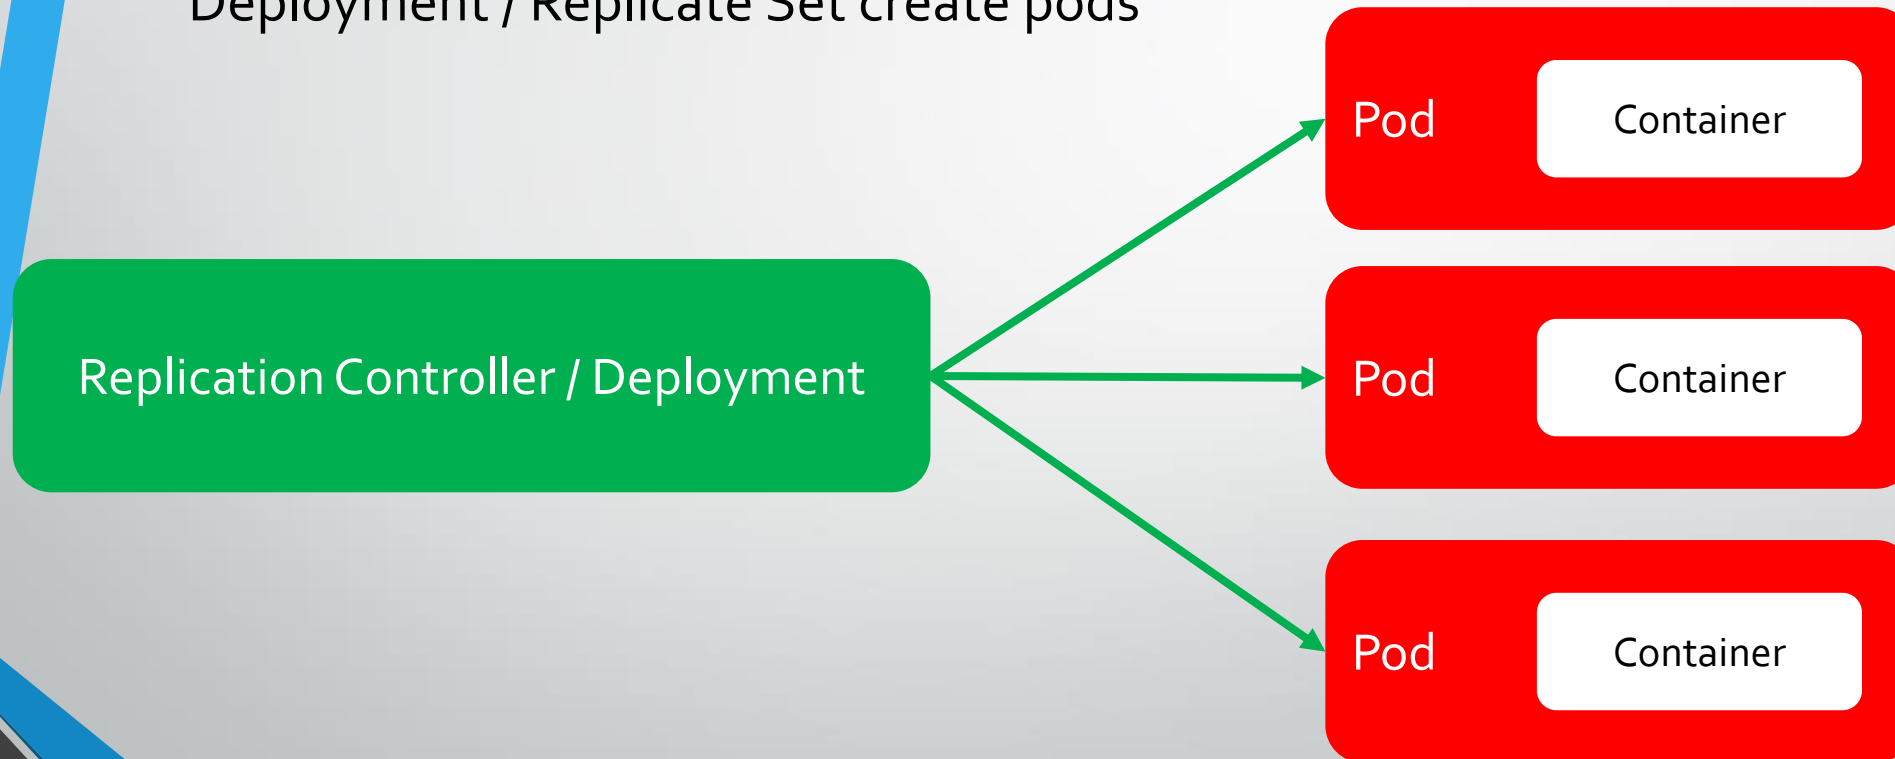
Deployment / Replicate Set create pods



Pod Deployment

Pods are not directly deployed

Deployment / Replicate Set create pods



Pod Deployment

Pods are not directly deployed

Deployment / Replicate Set create pods

Deployments manage ReplicaSets

Manages stateless applications



DaemonSet, CronJob, StatefulSet

Alternatively to Deployments, pods can be run using DaemonSets, CronJobs, and StatefulSets

DaemonSet, CronJob, StatefulSet

Alternatively to Deployments, pods can be run using DaemonSets, CronJobs, and StatefulSets

CronJobs can be scheduled to start pods

StatefulSets manage stateful applications

DaemonSet, CronJob, StatefulSet

DaemonSets run pods on every node in the cluster

- Logging
- Monitoring
- Backup
- Reports
- Automated testing

Horizontal Pod Autoscaler (HPA)

Queries resource utilization, e.g. CPU and RAM usage

Instructs ReplicationSet to scale out or scale in

Configures minimum and maximum number of pods

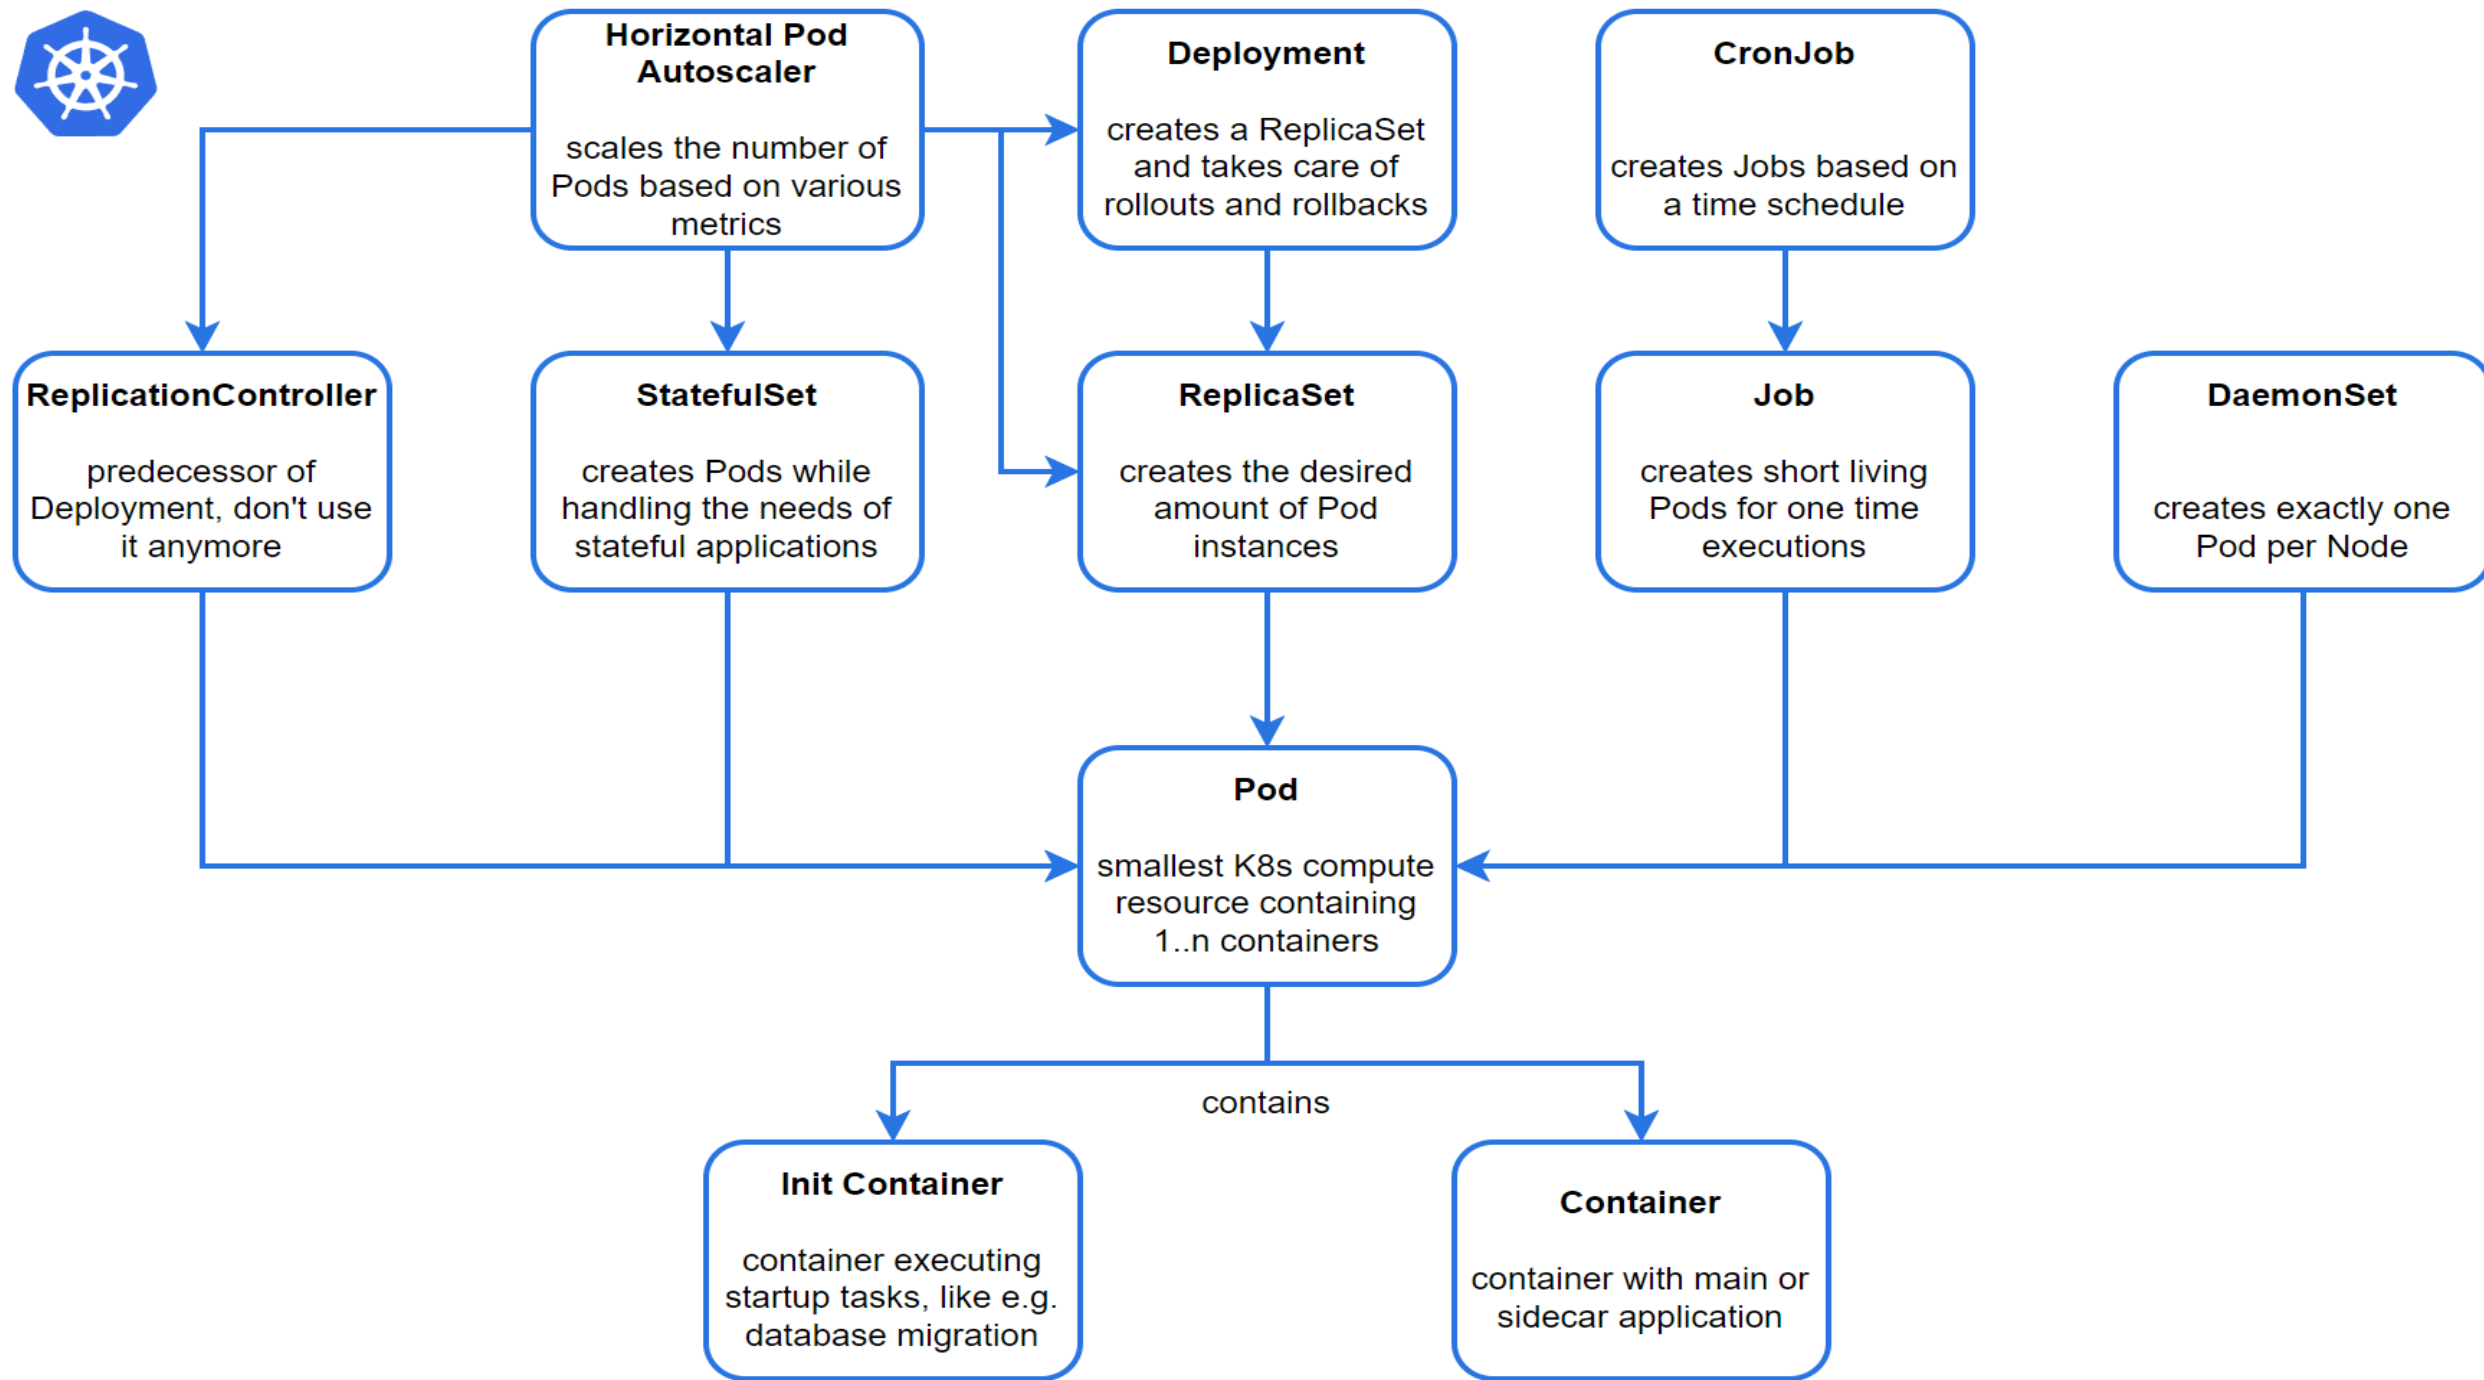
Horizontal Pod Autoscaler (HPA)

Queries resource utilization, e.g. CPU and RAM usage

Instructs ReplicationSet to scale out or scale in

Configures minimum and maximum number of pods

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: customerapi
spec:
  maxReplicas: 10
  minReplicas: 1
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: customerapi
  targetCPUUtilizationPercentage: 50
```





Auto Scaling Demo

Docker Health Checks

Docker provides a simple health check

--interval=DURATION

--timeout=DURATION

--start-period=DURATION

--retries should be set to N

Docker Health Checks

Problems with simple health checks:

- Application startup may be longer than expected
- Different startup and health check
- Specific port for checks

Liveness Probe

Checks if pod is alive

Sends HTTP request to check pod

Alive if answer \geq HTTP 200 & $<$ HTTP 400

- Pod will be restarted if dead

Liveness Probe

Checks if pod is alive

Sends HTTP request to check pod

Alive if answer \geq HTTP 200 & $<$ HTTP 400

- Pod will be restarted if dead

Configuration part of the Deployment

```
livenessProbe:  
  httpGet:  
    path: /health  
    port: http  
  initialDelaySeconds: 15
```


Readiness Probe

Checks if pod is ready to receive traffic

Sends HTTP request to check pod

Alive if answer \geq HTTP 200 & $<$ HTTP 400

- Traffic will be routed to the pod when ready

Readiness Probe

Checks if pod is ready to receive traffic

Sends HTTP request to check pod

Alive if answer \geq HTTP 200 & $<$ HTTP 400

- Traffic will be routed to the pod when ready

Configuration part of the Deployment

```
readinessProbe:  
  httpGet:  
    path: /health  
    port: http  
    initialDelaySeconds: 15
```

Resource Requests & Resource Limits

1000 Millicores = 1 Core

Memory is defined in bytes

Mebibyte = ~1MB

Configured in Deployment

Resource Requests & Resource Limits

Resource Requests

- Describe how many free resources a node has to have
- CPU and/or RAM

Resource Requests & Resource Limits

Resource Requests

- Describe how many free resources a node has to have
- CPU and/or RAM

Resource Limits

- Maximum resources a pod is allowed to use
- Pods gets throttled when it uses too many resources
- CPU and/or RAM

Resource Requests & Resource Limits

Resource Requests

- Describe how many free resources a node has to have
- CPU and/or RAM

```
resources:  
  limits:  
    cpu: 0.3  
    memory: 128Mi  
  requests:  
    cpu: 100m  
    memory: 64Mi
```

Resource Limits

- Maximum resources a pod is allowed to use
- Pods gets throttled when it uses too many resources
- CPU and/or RAM



Resource Quotas

Limit resource usage within a namespace

Configured in the ResourceQuota object

Resource Quotas

Limit resource usage within a namespace

Configured in the ResourceQuota object

Restrict the following resources:

- Number of pods
- CPU and RAM per request
- Total CPU and RAM



10 Min Break



Deployment Strategies

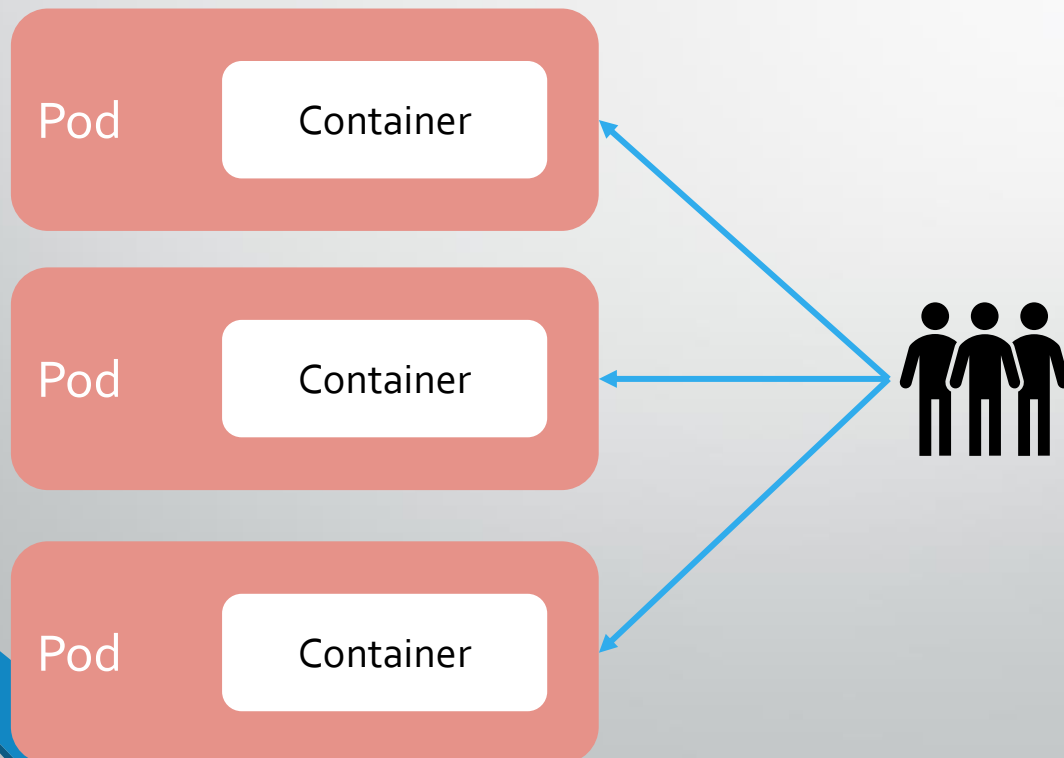
Blue Green deployment

Start all new pods and then switch

Deployment Strategies

Blue Green deployment

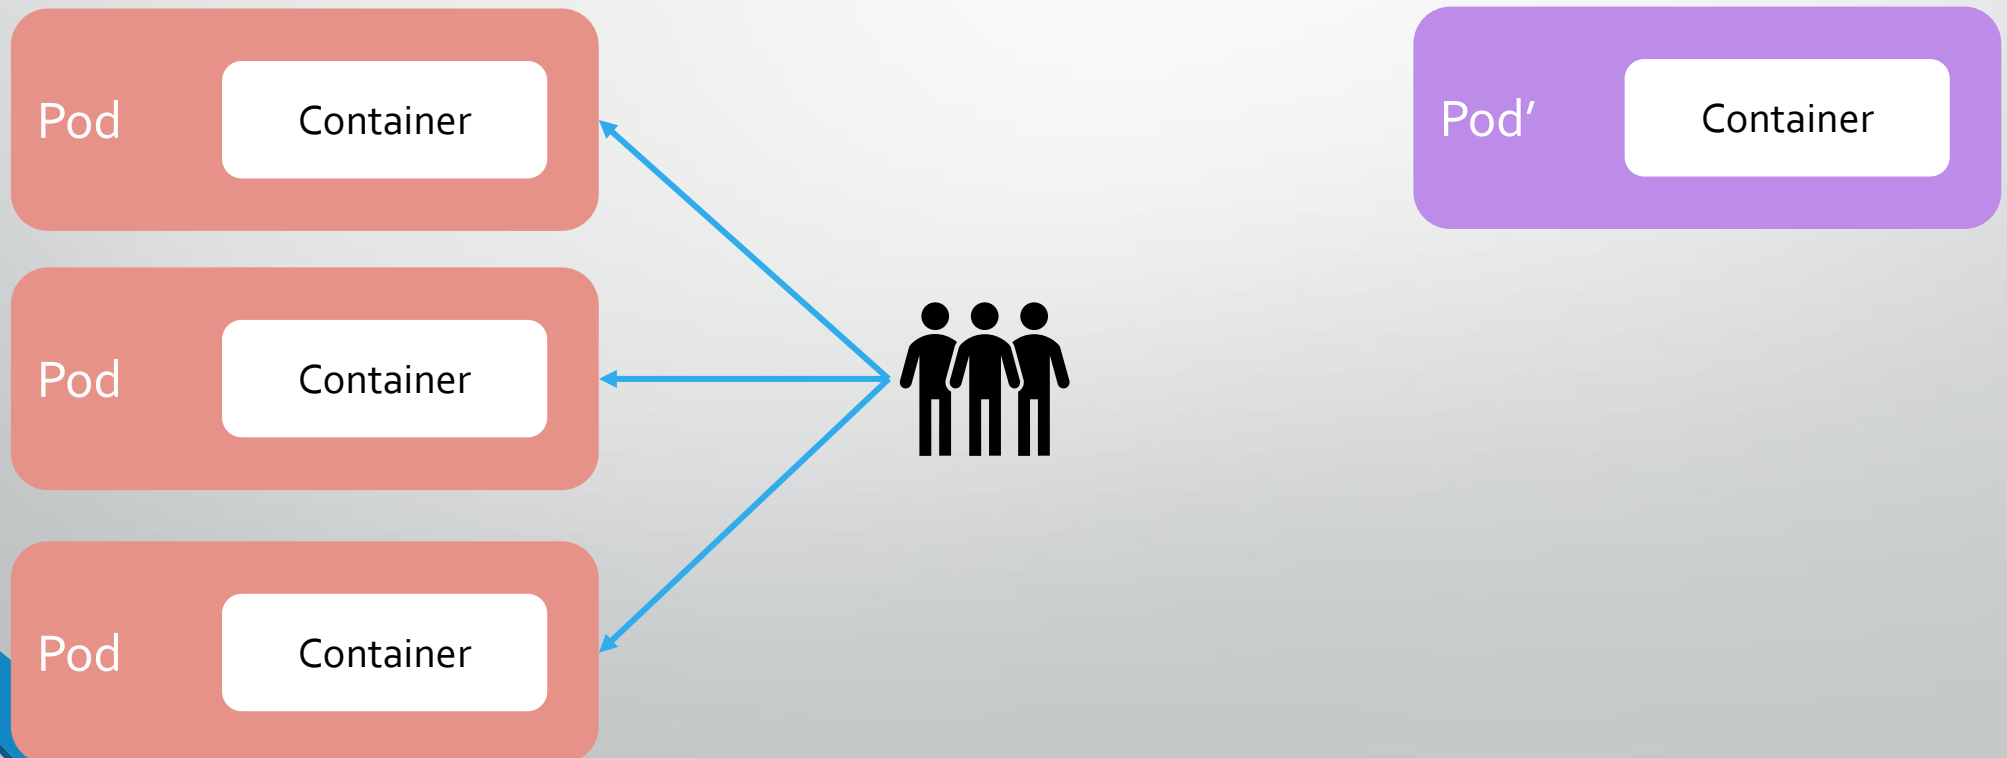
Start all new pods and then switch



Deployment Strategies

Blue Green deployment

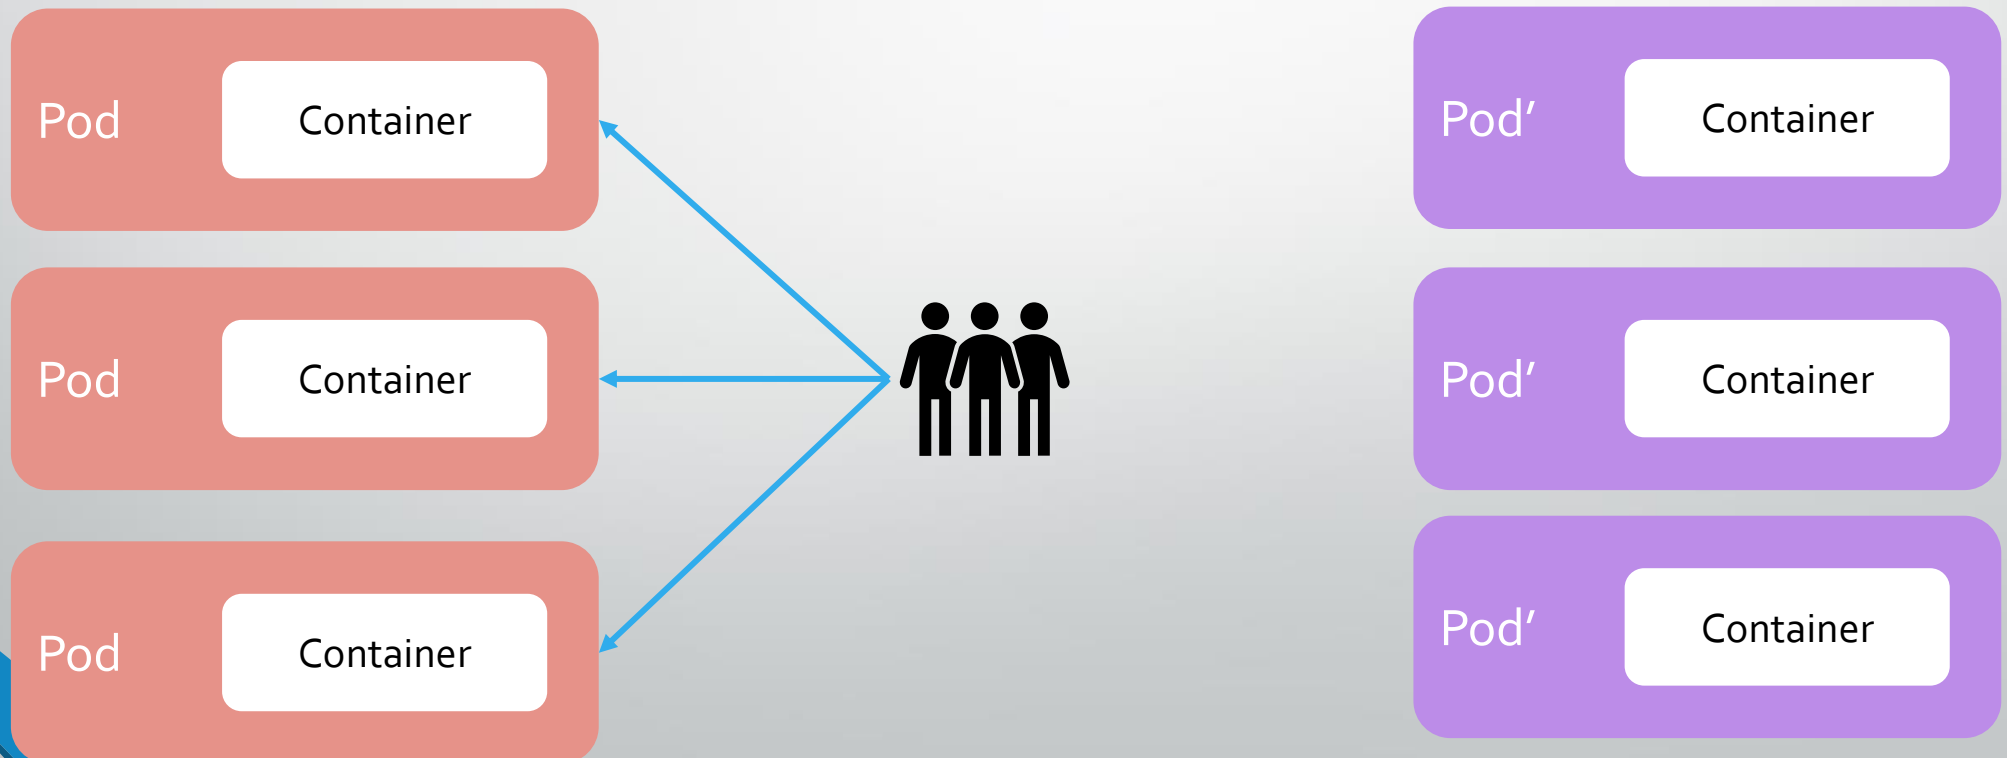
Start all new pods and then switch



Deployment Strategies

Blue Green deployment

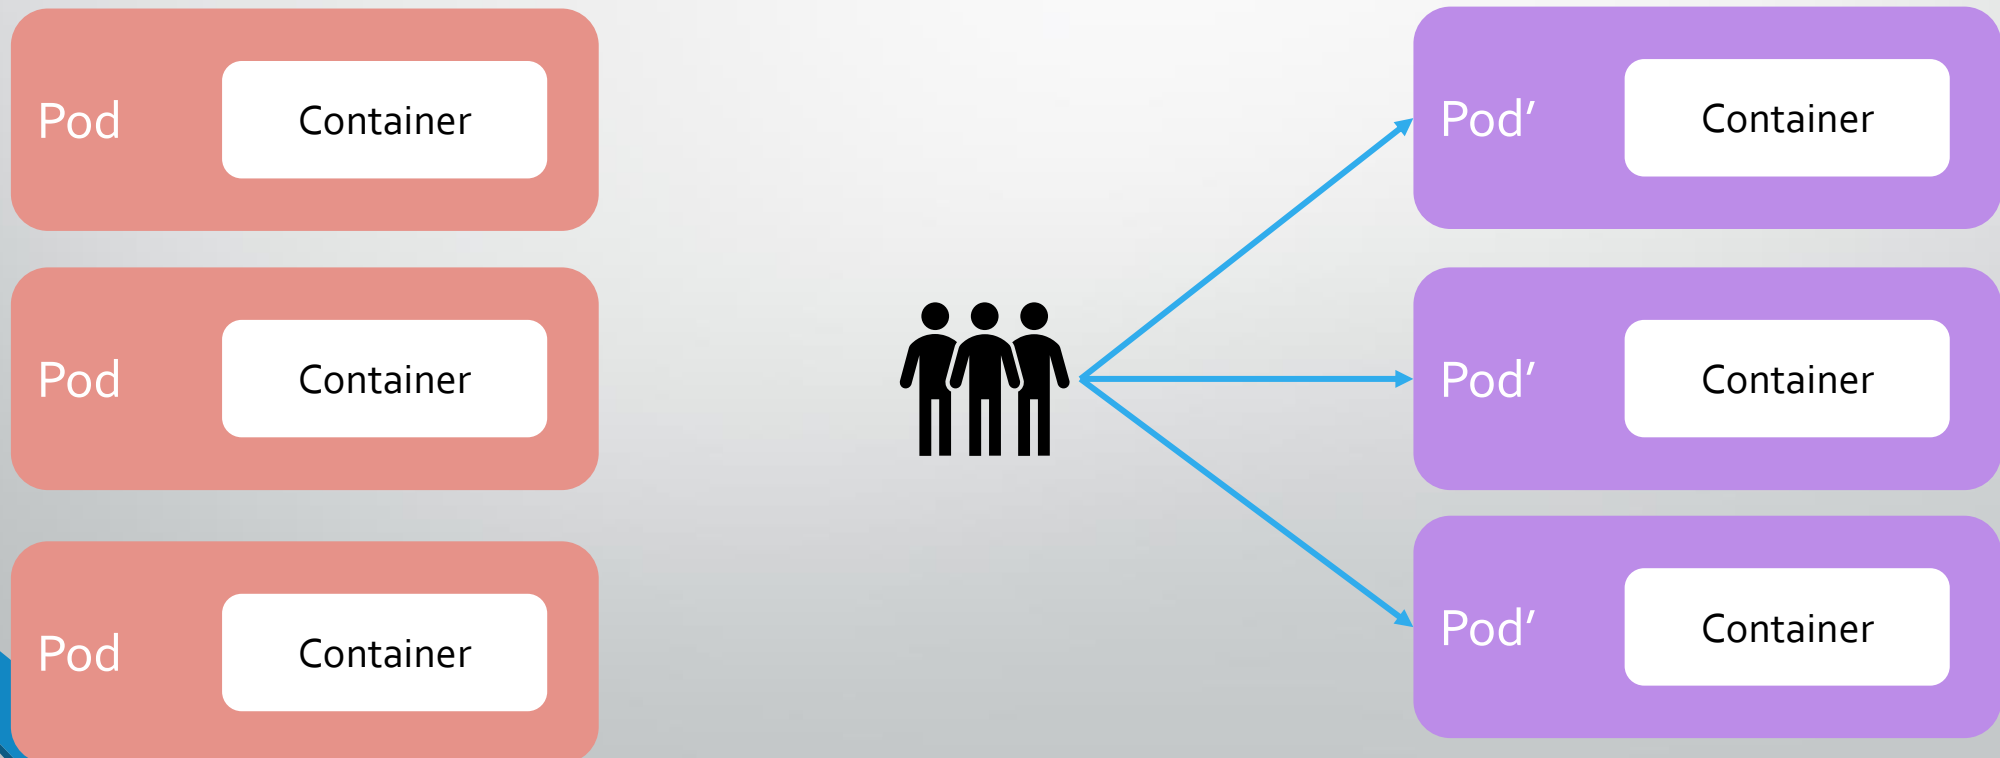
Start all new pods and then switch



Deployment Strategies

Blue Green deployment

Start all new pods and then switch





Deployment Strategies

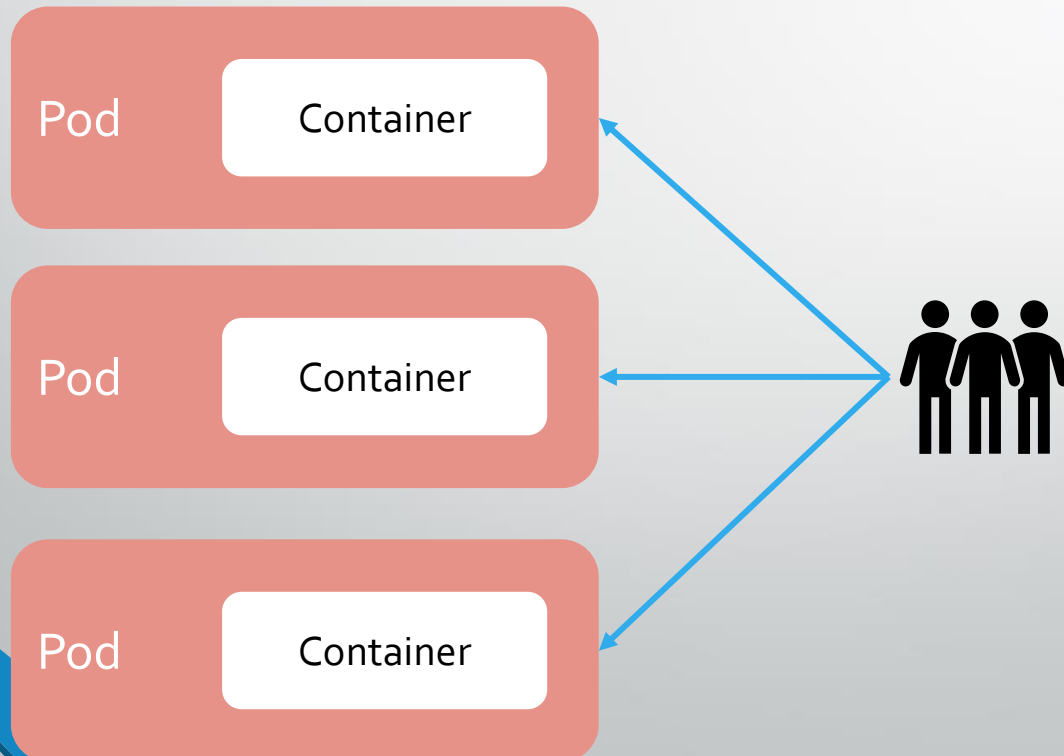
Rolling deployment

Replace old pods with new ones until all are replaced

Deployment Strategies

Rolling deployment

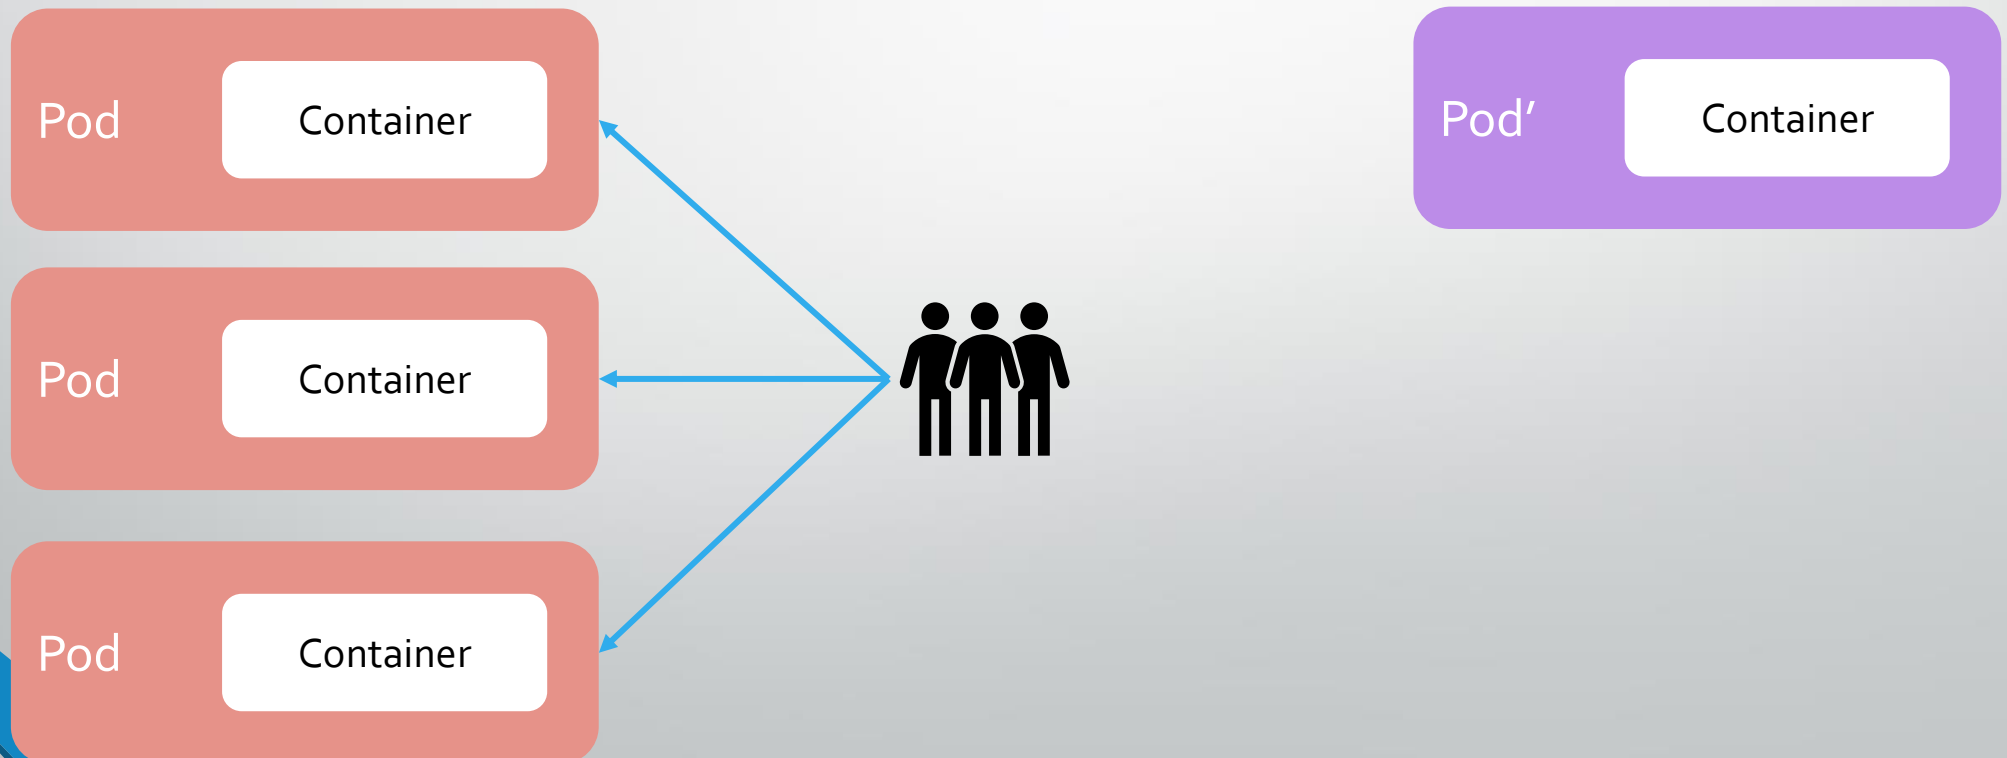
Replace old pods with new ones until all are replaced



Deployment Strategies

Rolling deployment

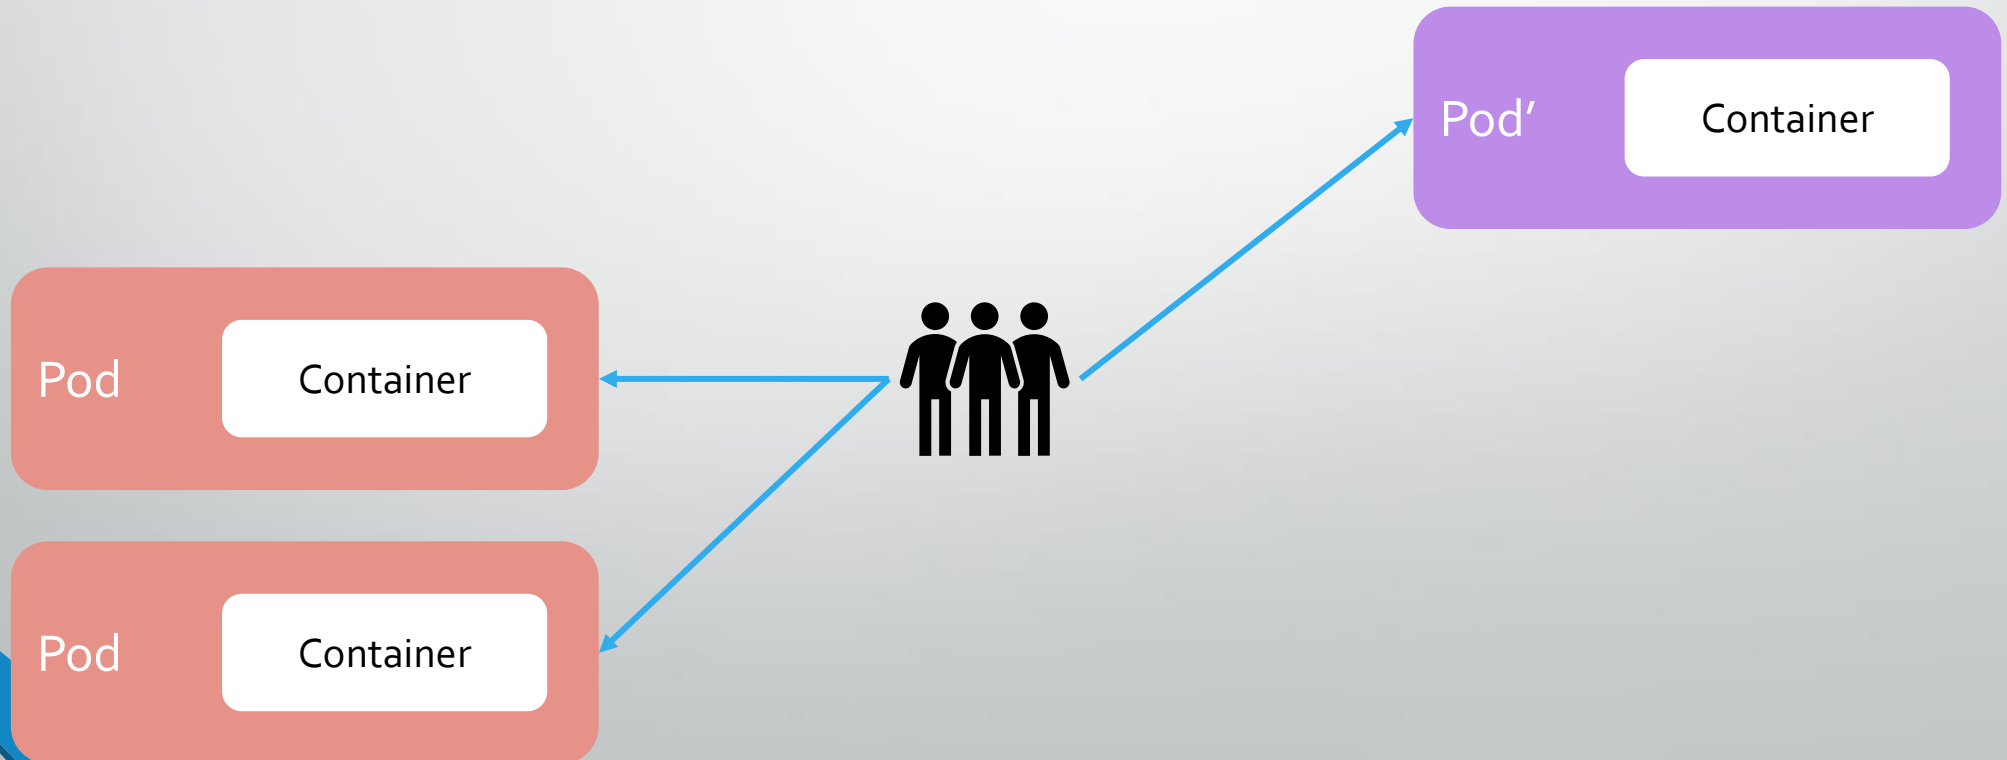
Replace old pods with new ones until all are replaced



Deployment Strategies

Rolling deployment

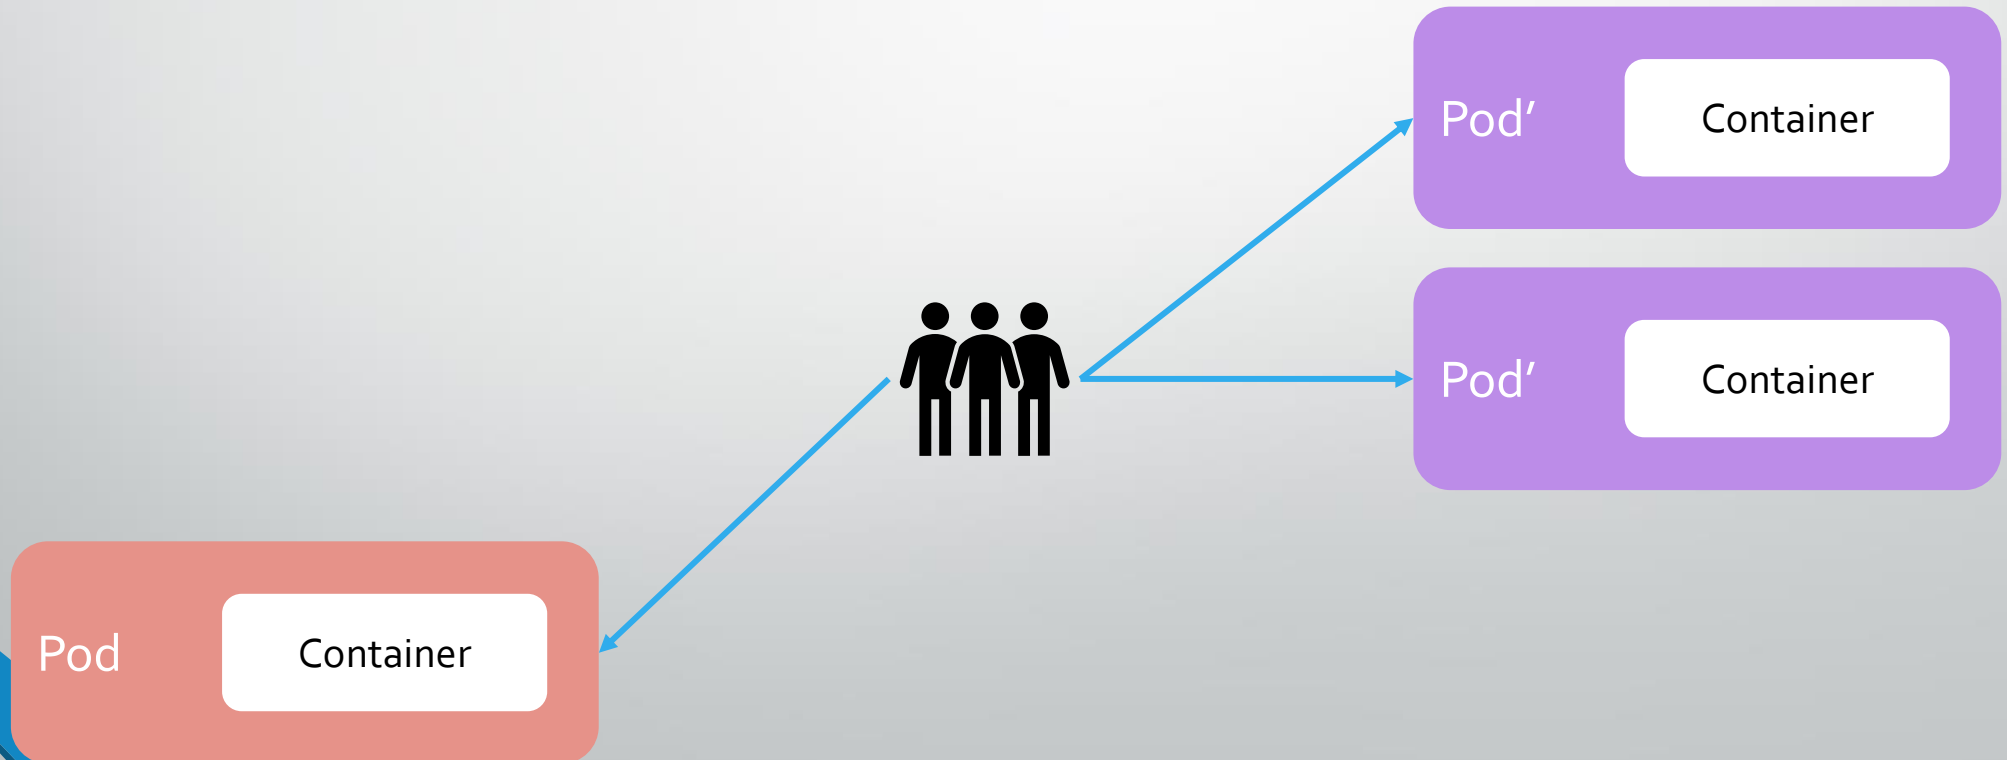
Replace old pods with new ones until all are replaced



Deployment Strategies

Rolling deployment

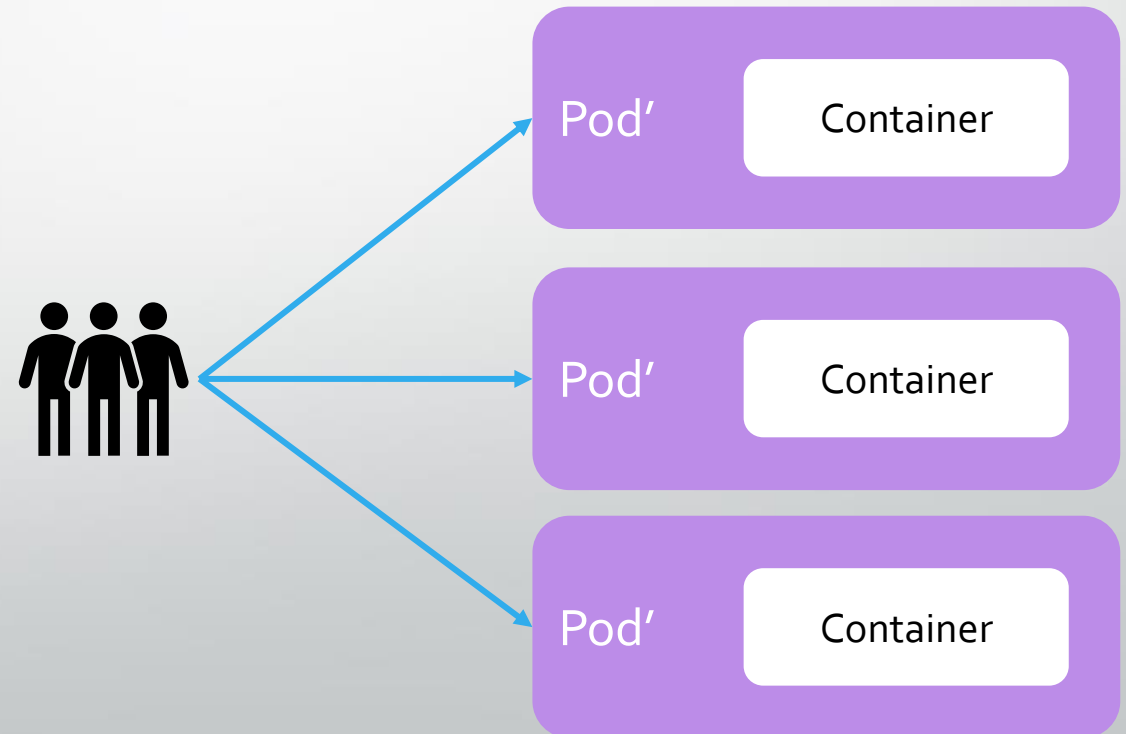
Replace old pods with new ones until all are replaced



Deployment Strategies

Rolling deployment

Replace old pods with new ones until all are replaced



Deployment Strategies

Rolling deployment

- Default deployment mode
- Application must support two versions

Deployment Strategies

Canary deployment

- Route only a small percentage of traffic to the new version
- Reduce the blast radius of a bad deployment
- Increase percentage of traffic gradually
- A/B testing



Zero Downtime Deployment Demo

Deployment Strategies

maxSurge:

- Max number of pods that can be created at a time
- Absolut number or percentage
- Default: 25%

Deployment Strategies

maxSurge:

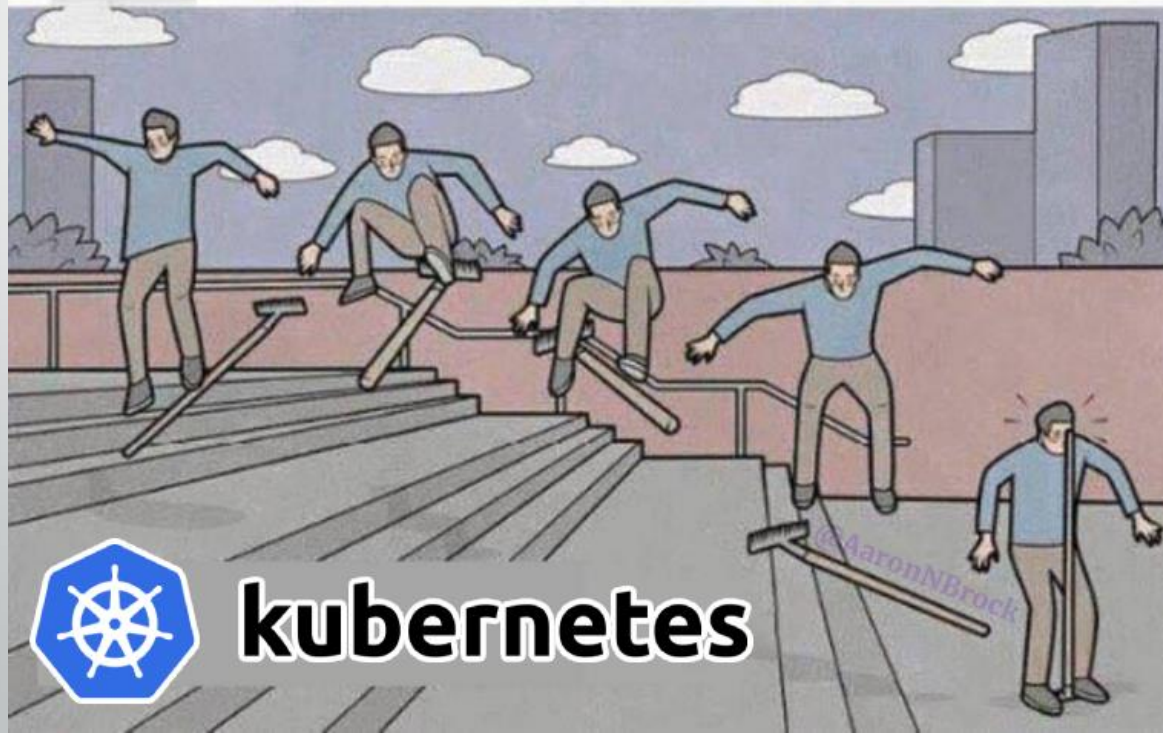
- Max number of pods that can be created at a time
- Absolut number or percentage
- Default: 25%

maxUnavailable:

- Max number of pods that can be available during the deployment
- Absolut number or percentage
- Default: 25%



virtual machines



kubernetes

Considerations when using K8s

Cloud-native architecture

Microservices

.NET Full Frameworks vs .NET Core

DevOps process and culture

Deploy fast and often

Fast paced development and deployment

Cloud hosted vs. on-premises

When not to use Kubernetes

Skills and experience of the team

Application that will be barely change

Big monolithic applications

Quick results

Very simple applications

Cloud hosted vs. on-premises



Exercise

Deploy your own Image from Dockerhub

Change ports in the Service

Deploy a new version of your container

Change Service Type to ClusterIP

Implement an HPA and test autoscaling

Readiness / Liveness Probe

Create an AKS cluster

Kubectl Commands

Get resource

kubectl get pods/service/deployment

Delete resource

kubectl delete pod/service/deployment

Display information about resource

kubectl describe pod/node/service resource-name


Add/update new resource

kubectl apply -f myfile.yaml [namespace=my-namespace]

Set current namespace

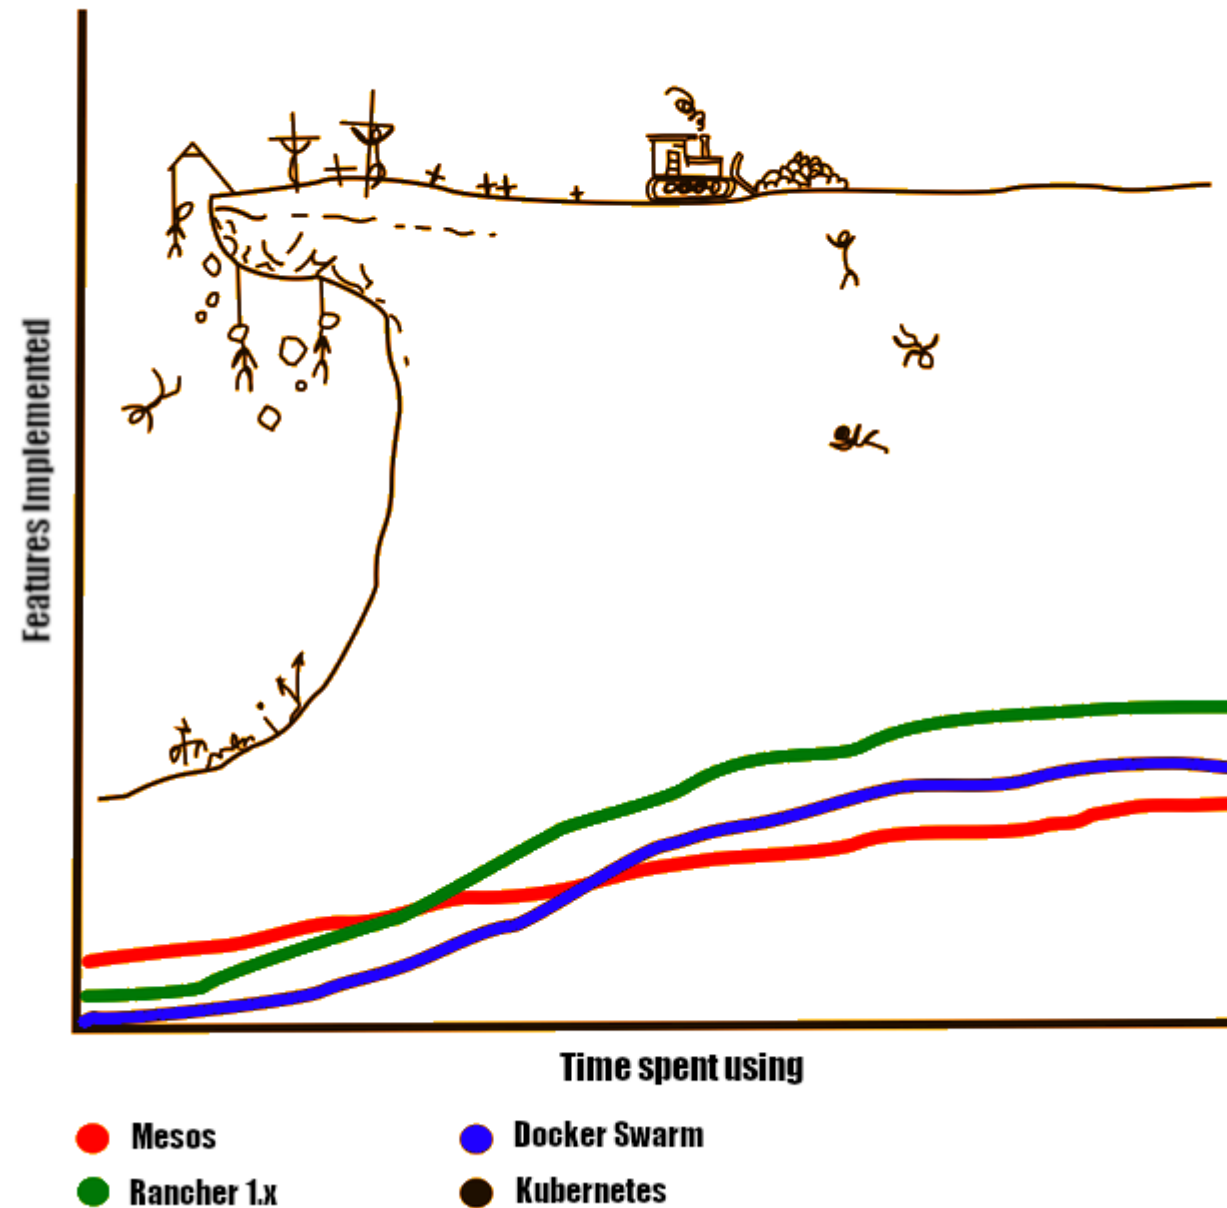
kubectl config set-context --current --namespace=my-namespace

Kubernetes Cheat Sheet: <https://kubernetes.io/docs/reference/kubectl/cheatsheet>



Advanced Tools and Learning Path

Learning curves of some Container Orchestration Engines



Certification

Kubernetes Fundamentals Training (35 hours):

<https://training.linuxfoundation.org/training/kubernetes-fundamentals>

Certified Kubernetes Administrator (CKA):

<https://training.linuxfoundation.org/certification/certified-kubernetes-administrator-cka>

Certified Kubernetes Application Developer (CKAD):

<https://training.linuxfoundation.org/certification/certified-kubernetes-application-developer-ckad>

Helm

Kubernetes Package Manager

Use Helm charts to define application

Share Helm charts

Open-source

Public Helm charts can be found on <https://artifacthub.io>

Documentation: <https://helm.sh>

Kustomize

Kubernetes configuration management tool

Scalable for various deployment sizes

Efficient configuration management

Enables customization without directly editing manifest files

Documentation: <https://kustomize.io>

GitOps

Automation using Git and CI/CD pipelines

Configuration as code

History of config changes + Pull Request reviews

ArgoCD: <https://argo-cd.readthedocs.io>

Flux: <https://fluxcd.io>

KEDA

Kubernetes Event-driven Autoscaling

Enables autoscaling workloads based on external events

- Azure Service Bus
- Apache Kafka
- Redis Streams
- MongoDB
- Azure DevOps Pipeline

KEDA

Optimize resource usage

64 built in scalers

Can be used without any changes in applications

Open-source

Documentation: <https://keda.sh>

Cilium

Networking, security and observability

Layer 3 to layer 7 security policies

mTLS between pods

Visibility into network and application layer interactions

Open-source

Documentation: <https://cilium.io>

Service Mesh

Manage communication between applications and pods

Service discovery

Failure handling

Observability: metrics, logging, tracing

Security: encryption, authentication, authorization

Service Mesh

Open-source Services:

Istio: <https://istio.io>

Linkerd: <https://linkerd.io>

Consul: <https://developer.hashicorp.com/consul>

Secret Management

Secrets are not encrypted by default

Secrets should be encrypted to keep them secret

- Database connection strings
- Passwords
- Certificates

Secret Management

Tools:

- Azure Key Vault Provider for Secrets Store CSI Driver: <https://learn.microsoft.com/en-us/azure/aks/csi-secrets-store-driver>
- HashiCorp Vault: <https://www.vaultproject.io>
- Sealed Secrets: <https://github.com/bitnami-labs/sealed-secrets>



Monitoring

Metrics

Logs

Traces

Alerting

Monitoring

Tools:

- Grafana: <https://grafana.com>
- Prometheus: <https://prometheus.io>
- Elastic Stack (ELK Stack):
<https://www.elastic.co/elastic-stack>
- Jaeger: <https://www.jaegertracing.io>
- Zipkin: <https://zipkin.io>

Azure Arc

Manage Kubernetes with Azure

Run Azure services on your Kubernetes cluster on-premises

- Azure App Service
- Azure Machine Learning
- Azure Managed SQL Database

Authentication with Azure

Access cluster which is hidden behind a firewall

<https://learn.microsoft.com/en-us/azure/azure-arc/kubernetes/overview>

Kubernetes Operations

Resource utilization

Tagging

Namespace limits

Pod security policies

Networking policies

RBAC policies

Logging

Checklist: <https://learnk8s.io/production-best-practices>

More Resources

Kubernetes: <https://kubernetes.io>

Docker: <https://www.docker.com>

K3s: <https://k3s.io>

Networking: <https://www.cni.dev/docs>

Graphics: <https://shipit.dev/posts/kubernetes-overview-diagrams.html>

Code:

<https://code.komaxgroup.com/global/training/innovation-days/2023/id23-5311-kubernetes-workshop>

Exercise for the Fast Ones

Helm

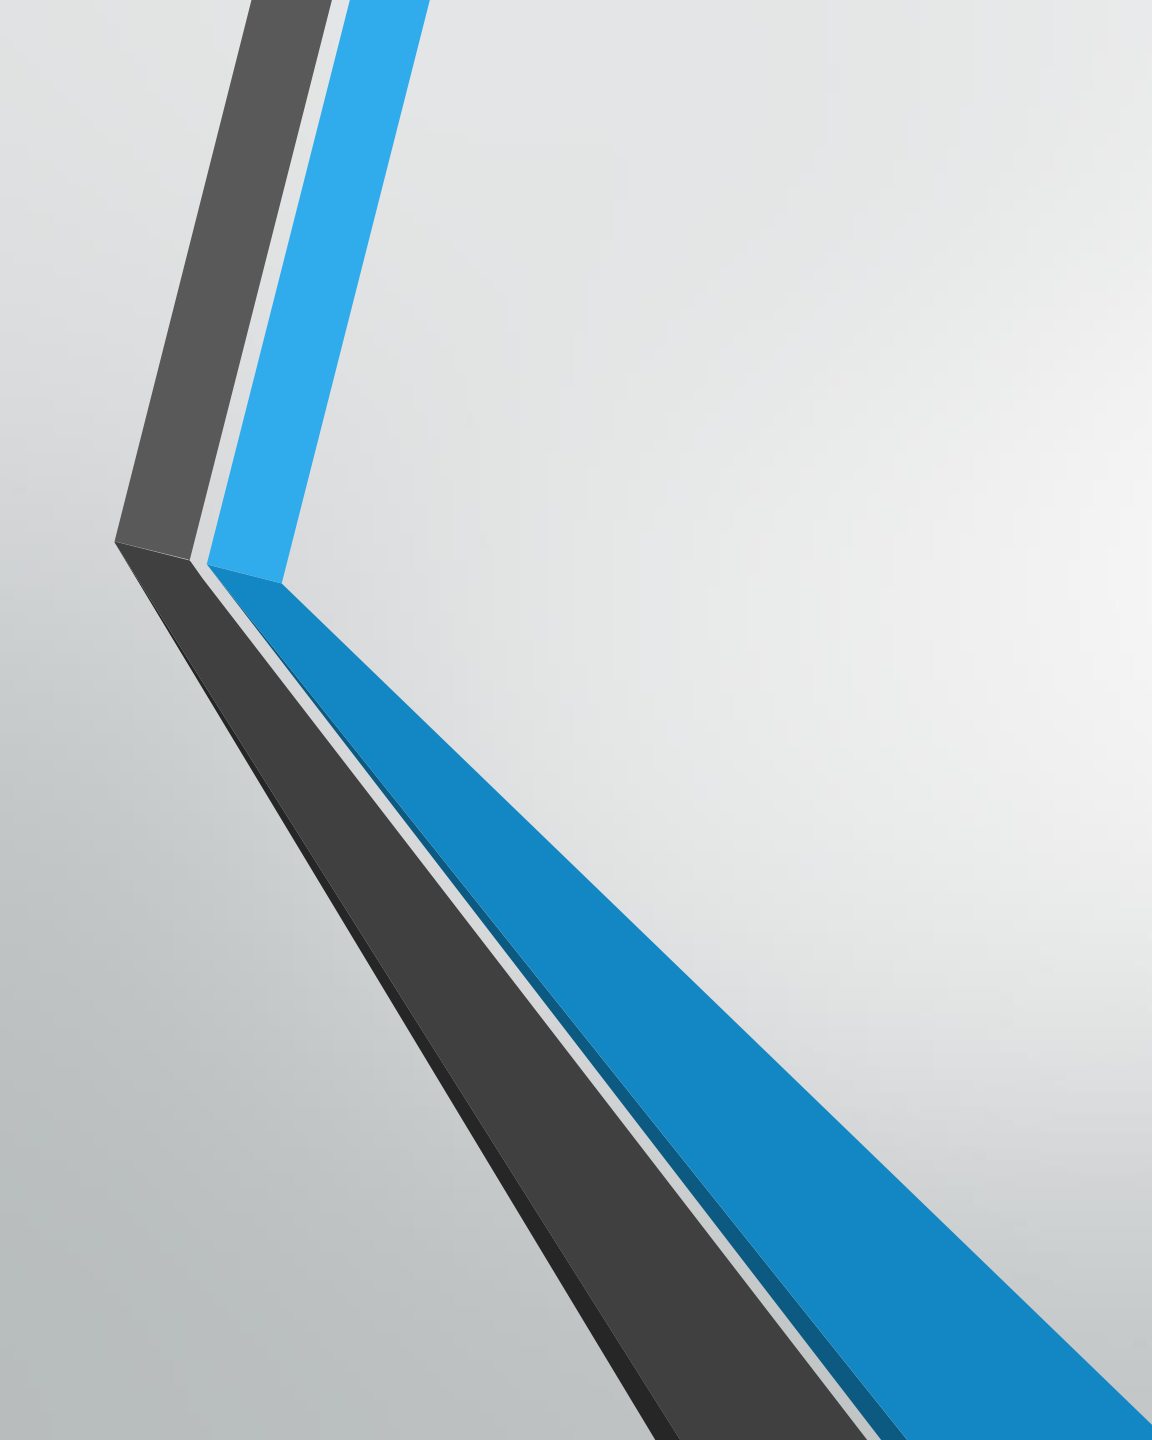
- Install Helm, create a Helm chart, check content, deploy helm chart, download public Helm chart

Keda Demo

- Azure Service Bus, Scale to zero, Azure DevOps Agent Scaling

Use a dashboard to see what's going on in your cluster

Install Istio and try out the demo



Q&A