


# Scaling to Success: Leveraging KEDA and Kubernetes for Optimal Azure DevOps Pipeline Performance

**Wolfgang Ofner**

**Freelance Cloud Architect, ProgrammingWithWolfgang**





# Wolfgang Ofner

Freelance Cloud Architect, Perth, Australia

Focus on Azure, Kubernetes, DevOps and .NET

<https://programmingwithwolfgang.com>

<https://www.linkedin.com/in/wolfgangofner>

[https://twitter.com/wolfgang\\_ofner](https://twitter.com/wolfgang_ofner)




# Agenda

- Architecture in SW projects
- Introduction to KEDA
- Scaling Azure DevOps Agents in Kubernetes
- KEDA Conclusion


# Simplified Architecture History

- Server – Client Architecture
- Only few clients
- No redundancy
- No high availability



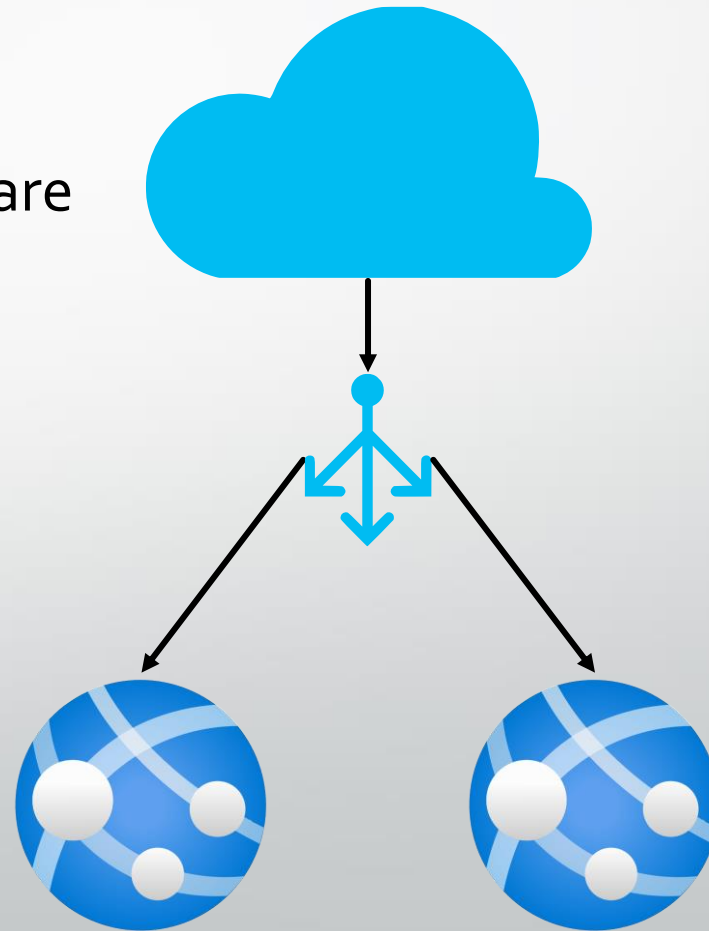
# Simplified Architecture History

- Static load balancing
- New VMs need to be added by hand
- Expensive on-premises hardware

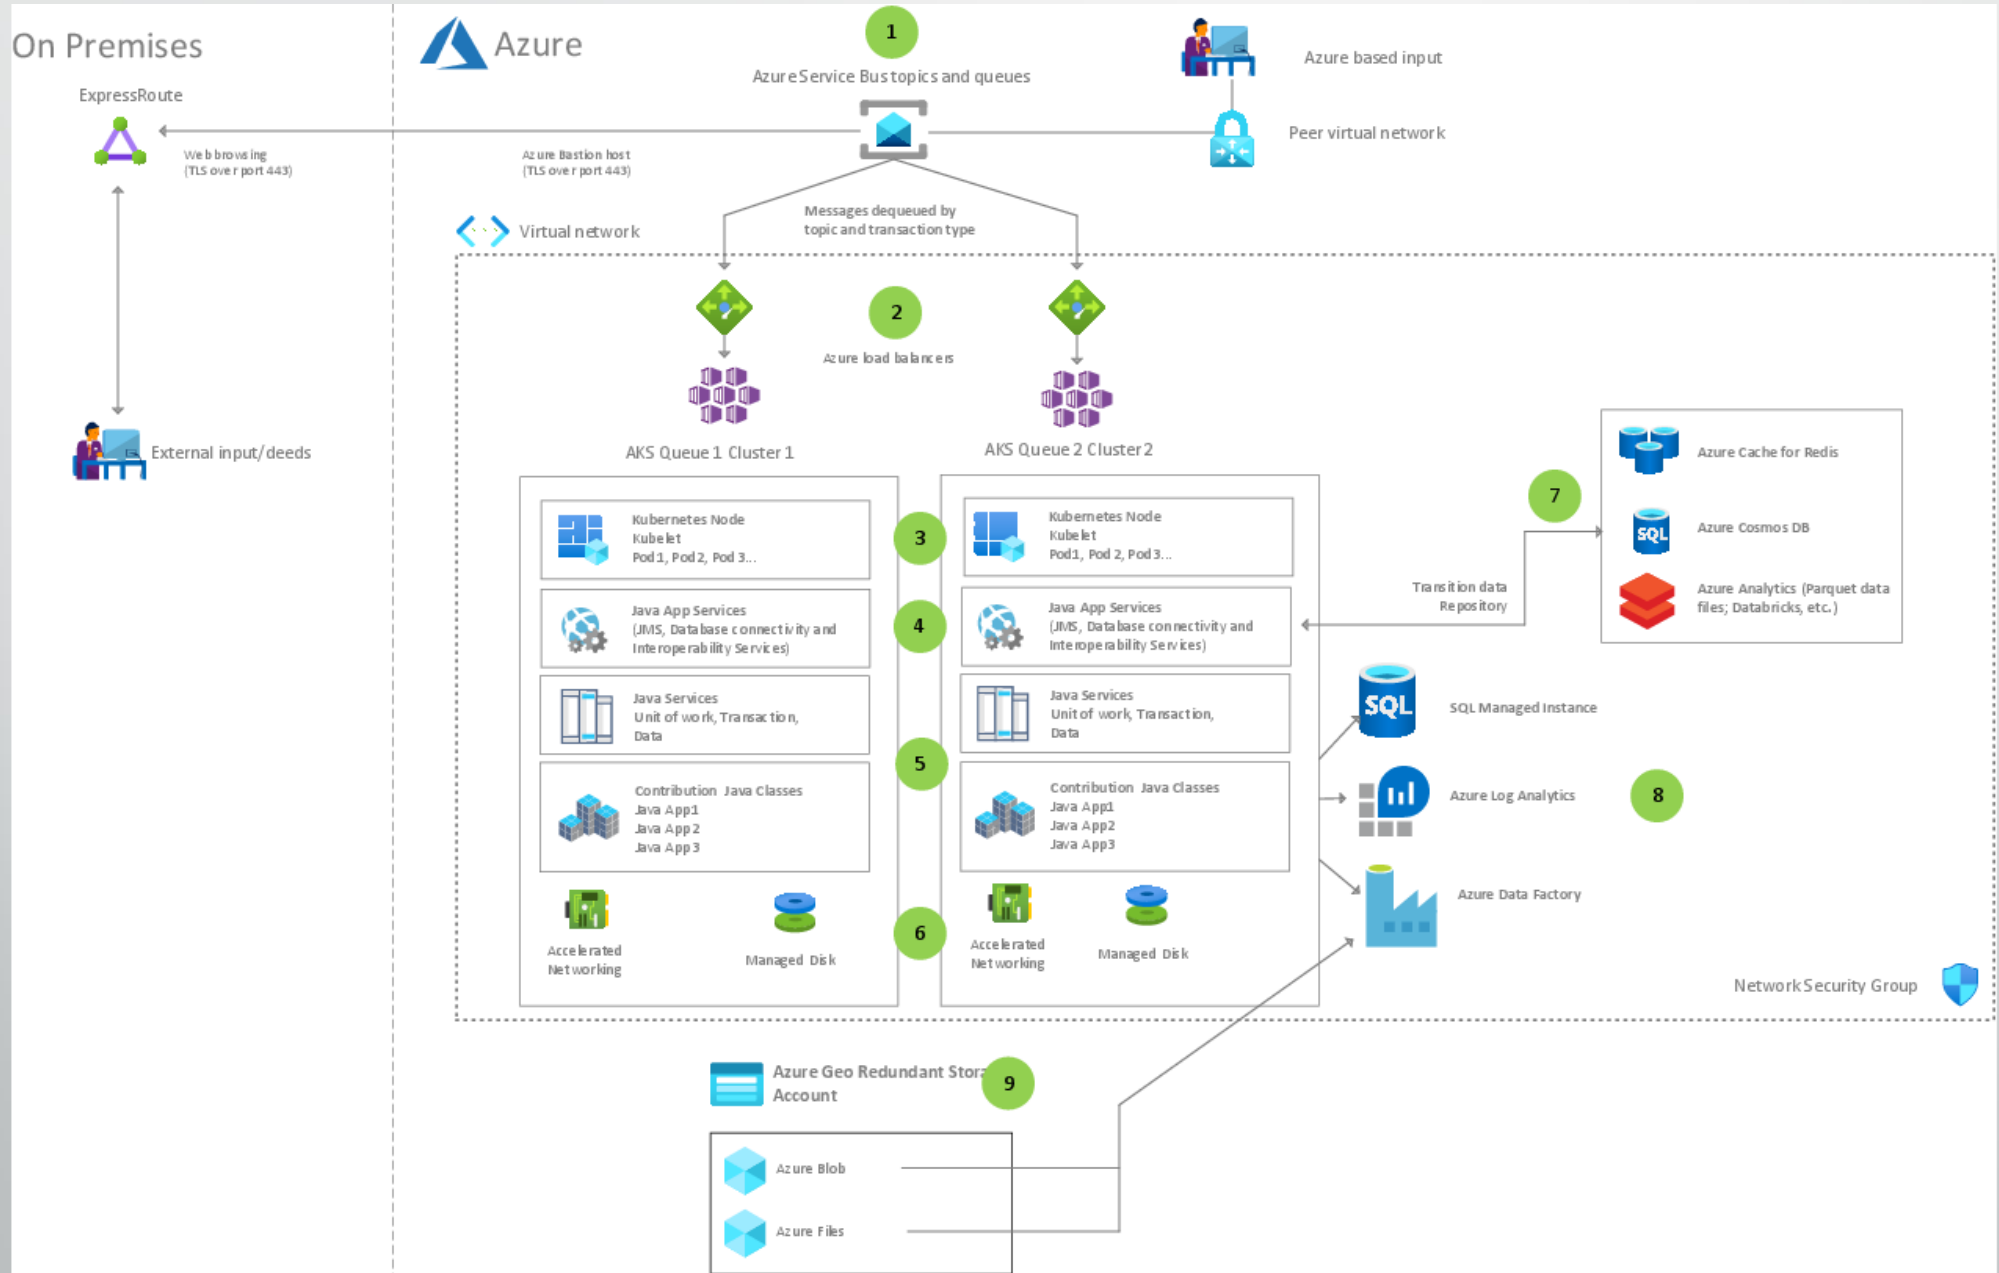


# Simplified Architecture History

- Automatically adding additional hardware
- Pay only what you need
- Mostly CPU or RAM based scaling



# Modern Architecture



# Kubernetes

- Horizontal Pod Autoscaler (HPA)
  - Scaling according to CPU and/or RAM
- Architectures get more and more complex
- Dependencies on external components
- Applications have to react to events
  - Database
  - Service Bus
  - Streams



# Horizontal Pod Autoscaler

- Scales Deployments or StatefulSets
- Adds or removes pods
- Scaling based on CPU or RAM usage
- Scaling on custom metrics
  - Query custom metrics from Kubernetes API
  - Prometheus
  - requests per second

# Horizontal Pod Autoscaler Configuration

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: customerapi
  namespace: customerapi-test
spec
  maxReplicas: 10
  minReplicas: 1
  averageCpuUtilization: 50
  scaleTargetRef
    apiVersion: apps/v1
    kind: Deployment
    name: customerapi
  behavior:
    scaleDown:
      policies:
        - type: Pods
          value: 4
          periodSeconds: 60
        - type: Percent
          value: 10
          periodSeconds: 60
      selectPolicy: Min
    scaleUp:
      policies:
        - type: Pods
          value: 5
          periodSeconds: 60
        - type: Percent
          value: 12
          periodSeconds: 60
      selectPolicy: Max
```

# Limitation of the HPA

- Black Friday
- Thousands of orders are stored in a queue
- Scaling using CPU or RAM is not sufficient
- No option for scaling in this scenario

# KEDA – Kubernetes Event-driven Autoscaling

- Kubernetes Event-driven Autoscaling
- Open source
- CNCF Project
- Maintained by
  - Docplanner Tech
  - Microsoft
  - Red Hat

# KEDA

- ~61 built-in Scaler
  - Apache Kafka
  - Azure Blob Storage
  - Azure Monitor
  - Azure Service Bus
  - Elastic Search
  - MongoDB
  - Prometheus
  - Redis Streams

# KEDA Use Cases

- Scale according to external events
- Scale to Zero
  - Bring serverless to your datacenter
  - Recreate Azure Functions architecture
  - Better resource usage

# KEDA Installation

- Installation via Helm charts
- Namespace: keda

# KEDA Installation

```
kubectl create namespace keda
```

```
helm repo add kedacore https://kedacore.github.io/charts
```

```
helm repo update
```

```
helm install keda kedacore/keda --namespace keda
```



# KEDA Resources

```
PS C:\Users\Wolfgang> kubectl get all -n keda
```

NAME	READY	STATUS	RESTARTS	AGE
pod/keda-operator-5748df494c-mxz9p	1/1	Running	0	124m
pod/keda-operator-metrics-apiserver-cb649dd48-jjhpc	1/1	Running	0	124m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/keda-operator-metrics-apiserver	ClusterIP	10.0.241.182	<none>	443/TCP,80/TCP	124m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/keda-operator	1/1	1	1	124m
deployment.apps/keda-operator-metrics-apiserver	1/1	1	1	124m

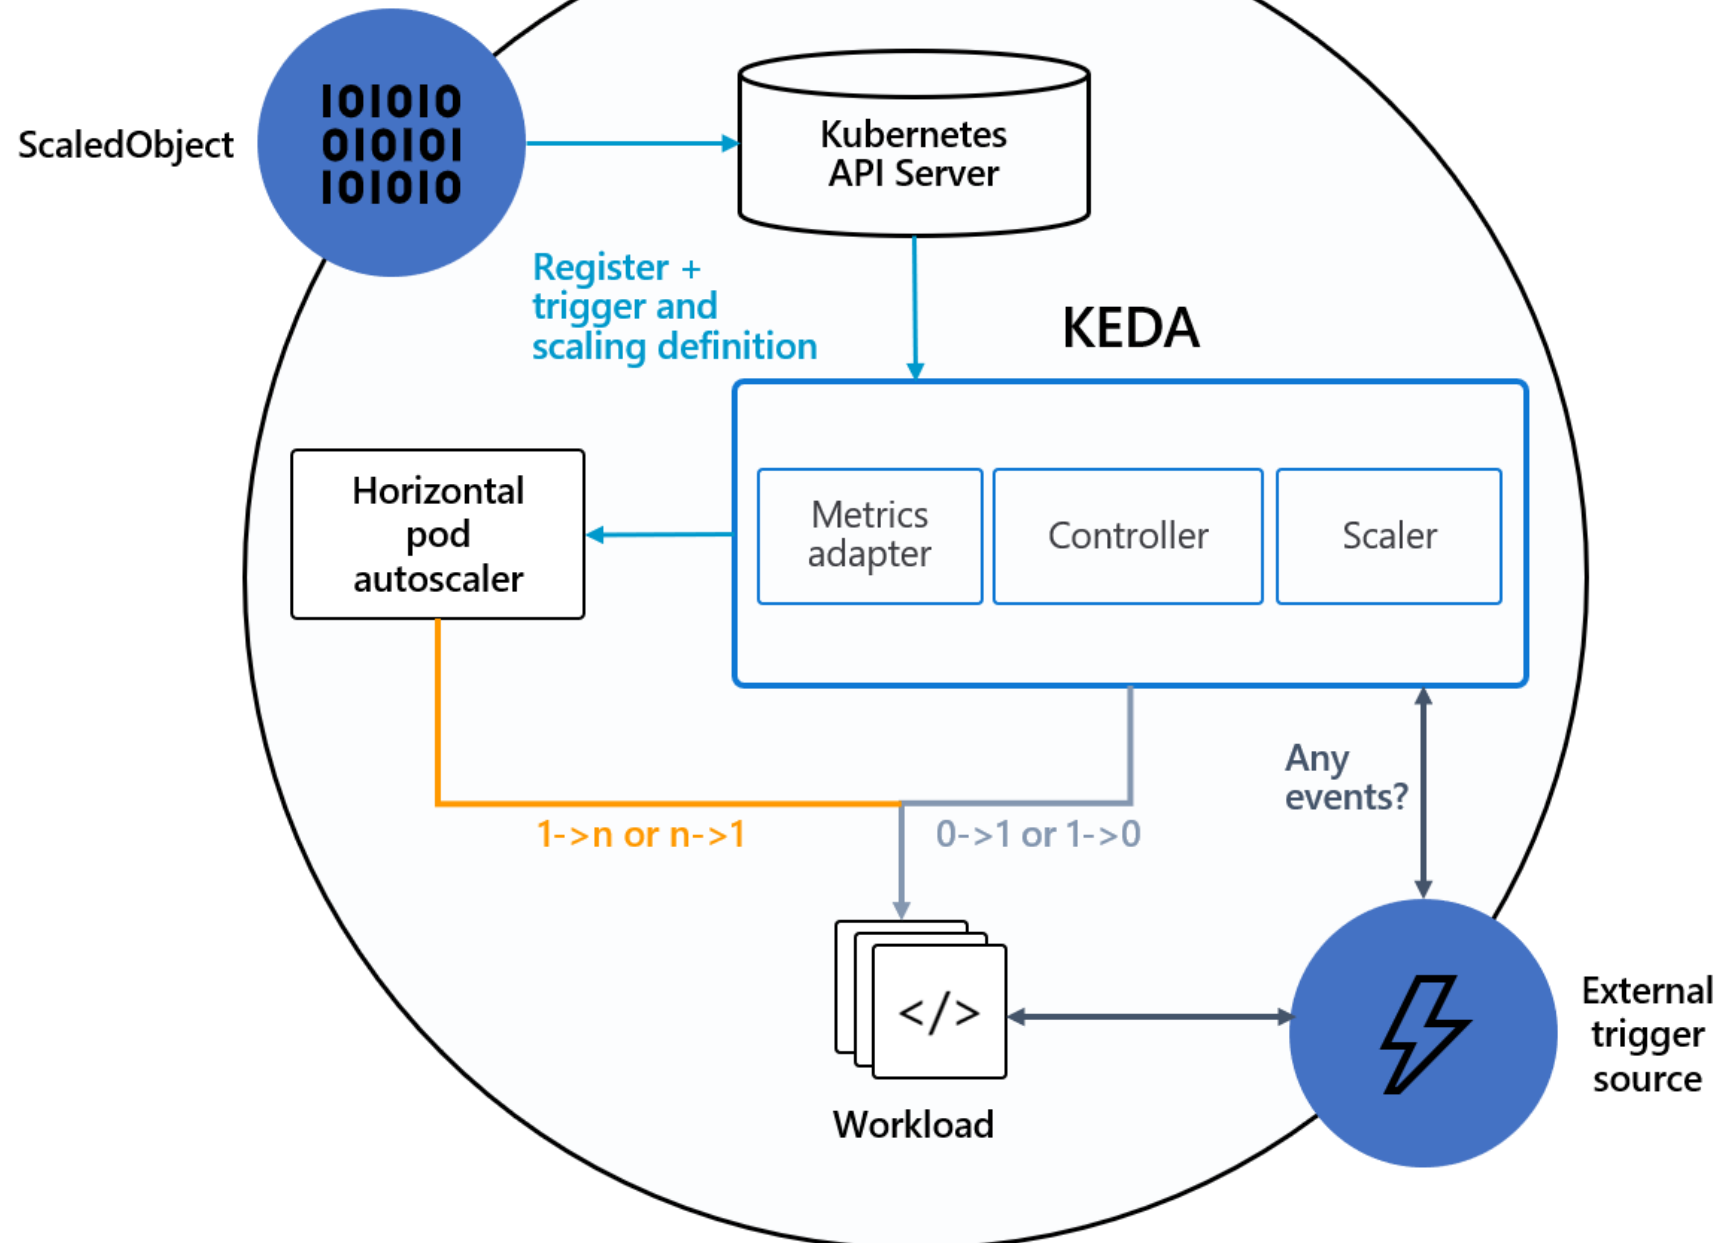
  

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/keda-operator-5748df494c	1	1	1	124m
replicaset.apps/keda-operator-metrics-apiserver-cb649dd48	1	1	1	124m

# KEDA Architecture

- 2 components for KEDA
  - Agent or Operator
  - Metrics Server
- Uses HPA for scaling
- Seamless integration into existing architecture

# Kubernetes cluster



# KEDA Architecture

- 2 components for KEDA
  - Agent
  - Metrics Server
- Uses HPA for scaling
- Seamless integration into existing architecture
- 2 custom K8s resources for scaler
  - ScaledObject
  - TriggerAuthentication

# ScaledObject

```
apiVersion: keda.sh/v1alpha1
kind: ScaledObject
metadata:
  name: kedademoapi-scaler
spec:
  scaleTargetRef:
    name: kedademoapi
  minReplicaCount: 0
  maxReplicaCount: 10
  pollingInterval: 30
  cooldownPeriod: 30
  triggers:
    - type: azure-servicebus
      metadata:
        queueName: KedaDemo
        queueLength: '5'
        authenticationRef:
          name: trigger-
            authentication-kedademoapi
```

# TriggerAuthentication

```
apiVersion: keda.sh/v1alpha1
kind: TriggerAuthentication
metadata:
  name: trigger-authentication-kedademoapi
spec:
  secretTargetRef:
    - parameter: connection
      name: kedademoapi-connectionstrings
      key: AzureServiceBus__ConnectionString
```

# Kubernetes Secret

```
PS C:\Users\Wolfgang> kubectl get secrets
```

NAME	TYPE	DATA	AGE
default-token-88lzb	kubernetes.io/service-account-token	3	26h
kedademoapi-connectionstrings	Opaque	1	26h
kedademoapi-tls	kubernetes.io/tls	2	26h
sh.helm.release.v1.kedademoapi-kedademoapi-test.v1	helm.sh/release.v1	1	26h
sh.helm.release.v1.kedademoapi-kedademoapi-test.v2	helm.sh/release.v1	1	22h

```
PS C:\Users\Wolfgang> kubectl describe secret kedademoapi-connectionstrings
```

```
Name:          kedademoapi-connectionstrings
Namespace:     kedademoapi-test
Labels:        app.kubernetes.io/managed-by=Helm
Annotations:   meta.helm.sh/release-name: kedademoapi-kedademoapi-test
               meta.helm.sh/release-namespace: kedademoapi-test
```

```
Type:  Opaque
```

```
Data
====
AzureServiceBus__ConnectionString: 165 bytes
```

# Kubernetes Secret

Namespace Overview ▾ > Config and Storage ▾ > Secrets ▾ > kedademoapi-connectionstrings

## kedademoapi-connectionstrings

Summary


Metadata

Resource Viewer

YAML

```
1 ---
2 apiVersion: v1
3 data:
4   AzureServiceBus__ConnectionString: RW5kcG9pbnQ9c2I6Ly93b2xmZ2FuZ2t1ZGFkZW1vLnNlcr
5 kind: Secret
```





# Azure DevOps Agent with KEDA

# Demo: Scaling ADO Agent with KEDA

- Azure DevOps preparation
- Build Docker image
- Test locally
- Deploy to Kubernetes
- Apply KEDA scaling

# Azure DevOps Limitations

- ADO Pipelines support scale to zero but need at least one agent registered
- ADO Pipelines can not queue a job with an empty agent pool
- Licensing limits parallel jobs

# KEDA ADO Scaling Limitations

- Cancelling a pipeline does not stop running pods
- KEDA does not remove completed pods
- Azure DevOps does not remove offline agents from the agent pool


# KEDA in Production

- Microsoft uses KEDA for Azure Services
  - Azure Container Apps
  - Azure App Services with Azure Arc
- KEDA 1.0.0 → 17. Nov 2019
- Currently 2.10
- Over 6k GitHub stars

# Resources

- Demo Application
  - <https://github.com/WolfgangOfner/MicroserviceDemo/tree/master/KedaDemoApi>
  - <https://github.com/WolfgangOfner/Ado-Agent-Keda>
- KEDA
  - <https://keda.sh>
- KEDA GitHub
  - <https://github.com/kedacore/keda>
- KEDA Architecture Screenshot
  - <https://keda.sh/docs/2.6/concepts/#architecture>





# Contact

Level Up your Kubernetes Scaling with KEDA

Wolfgang Ofner

<https://programmingwithwolfgang.com>


<https://www.linkedin.com/in/wolfgangofner>

[https://twitter.com/wolfgang\\_ofner](https://twitter.com/wolfgang_ofner)



**Thank you for watching!**

Remember to rate the presentation and  
leave your questions in the section below.



[www.WarszawskieDniInformatyki.pl](http://www.WarszawskieDniInformatyki.pl)



31 March - 1 April 2023



Politechnika Warszawska + online

**MAIN ORGANIZER:** **AcademicPartners**  
FUNDACJA

**ORGANIZING COMMITTEE:** dozens of organizations from the IT / data science sector (full list on the event website)